

R Companion for

Sampling

Design and Analysis

Third Edition

R Companion for
Sampling
Design and Analysis
Third Edition

Yan Lu and Sharon L. Lohr



CRC Press

Taylor & Francis Group

Boca Raton London New York

CRC Press is an imprint of the
Taylor & Francis Group, an **informa** business
A CHAPMAN & HALL BOOK

First edition published 2022
by CRC Press
6000 Broken Sound Parkway NW, Suite 300, Boca Raton, FL 33487-2742

and by CRC Press
2 Park Square, Milton Park, Abingdon, Oxon, OX14 4RN

© 2022 Yan Lu and Sharon L. Lohr

CRC Press is an imprint of Taylor & Francis Group, LLC

Reasonable efforts have been made to publish reliable data and information, but the author and publisher cannot assume responsibility for the validity of all materials or the consequences of their use. The authors and publishers have attempted to trace the copyright holders of all material reproduced in this publication and apologize to copyright holders if permission to publish in this form has not been obtained. If any copyright material has not been acknowledged please write and let us know so we may rectify in any future reprint.

Except as permitted under U.S. Copyright Law, no part of this book may be reprinted, reproduced, transmitted, or utilized in any form by any electronic, mechanical, or other means, now known or hereafter invented, including photocopying, microfilming, and recording, or in any information storage or retrieval system, without written permission from the publishers.

For permission to photocopy or use material electronically from this work, access www.copyright.com or contact the Copyright Clearance Center, Inc. (CCC), 222 Rosewood Drive, Danvers, MA 01923, 978-750-8400. For works that are not available on CCC please contact mpkbookspermissions@tandf.co.uk

Trademark notice: Product or corporate names may be trademarks or registered trademarks and are used only for identification and explanation without intent to infringe. SAS® and all other SAS Institute Inc. product or service names are registered trademarks or trademarks of SAS Institute Inc. in the USA and other countries. ® indicates USA registration.

Library of Congress Cataloging-in-Publication Data

Names: Lu, Yan, author. | Lohr, Sharon L., 1960- author.

Title: R companion for sampling : design and analysis / Yan Lu and Sharon L. Lohr.

Description: First edition. | Boca Raton : CRC Press, 2022. | Includes bibliographical references and index. | Summary: "The R Companion for Sampling: Design and Analysis, designed to be read alongside Sampling: Design and Analysis, Third Edition by Sharon L. Lohr (SDA; 2022, CRC Press), shows how to use functions in base R and contributed packages to perform calculations for the examples in SDA. No prior experience with R is needed. Chapter 1 tells you how to obtain R and RStudio, introduces basic features of the R statistical software environment, and helps you get started with analyzing data. Each subsequent chapter provides step-by-step guidance for working through the data examples in the corresponding chapter of SDA, with code, output, and interpretation. Tips and warnings help you develop good programming practices and avoid common survey data analysis errors. R features and functions are introduced as they are needed so you can see how each type of sample is selected and analyzed. Each chapter builds on the knowledge developed earlier for simpler designs; after finishing the book, you will know how to use R to select and analyze almost any type of probability sample"-- Provided by publisher.

Identifiers: LCCN 2021039318 (print) | LCCN 2021039319 (ebook) | ISBN 9781032135946 (paperback) | ISBN 9781032132150 (hardback) | ISBN 9781003228196 (ebook)

Subjects: LCSH: R (Computer program language) | Sampling (Statistics)
Classification: LCC QA276.45.R3 L8 2022 (print) | LCC QA276.45.R3 (ebook)
| DDC 519.5/202855133--dc23

LC record available at <https://lccn.loc.gov/2021039318>

LC ebook record available at <https://lccn.loc.gov/2021039319>

ISBN: 978-1-032-13215-0 (hbk)

ISBN: 978-1-032-13594-6 (pbk)

ISBN: 978-1-003-22819-6 (ebk)

DOI: 10.1201/9781003228196

Access the Support Material: <https://www.routledge.com/9781032135946>

To Guoyi and Lynn, and to Doug

Contents

Preface	xi
1 Getting Started	1
1.1 Obtaining the Software	2
1.2 Installing R Packages	2
1.3 R Basics	4
1.4 Reading Data into R	5
1.5 Saving Output	7
1.6 Integrating R Output into L ^A T _E X Documents	10
1.7 Missing Data	12
1.8 Summary, Tips, and Warnings	13
2 Simple Random Sampling	15
2.1 Selecting a Simple Random Sample	15
2.2 Computing Statistics from a Simple Random Sample	18
2.3 Additional Code for Exercises	24
2.4 Summary, Tips, and Warnings	25
3 Stratified Sampling	27
3.1 Allocation Methods	27
3.2 Selecting a Stratified Random Sample	30
3.3 Computing Statistics from a Stratified Random Sample	32
3.4 Estimating Proportions from a Stratified Random Sample	36
3.5 Additional Code for Exercises	37
3.6 Summary, Tips, and Warnings	38
4 Ratio and Regression Estimation	41
4.1 Ratio Estimation	41
4.2 Regression Estimation	44
4.3 Domain Estimation	46
4.4 Poststratification	48
4.5 Ratio Estimation with Stratified Sampling	49
4.6 Model-Based Ratio and Regression Estimation	50
4.7 Summary, Tips, and Warnings	54
5 Cluster Sampling with Equal Probabilities	57
5.1 Estimates from One-Stage Cluster Samples	57
5.2 Estimates from Multi-Stage Cluster Samples	59
5.3 Model-Based Design and Analysis for Cluster Samples	63
5.4 Additional Code for Exercises	65
5.5 Summary, Tips, and Warnings	67

6	Sampling with Unequal Probabilities	69
6.1	Selecting a Sample with Unequal Probabilities	69
6.1.1	Sampling with Replacement	69
6.1.2	Sampling without Replacement	70
6.2	Selecting a Two-Stage Cluster Sample	71
6.3	Computing Estimates from an Unequal-Probability Sample	77
6.3.1	Estimates from with-Replacement Samples	77
6.3.2	Estimates from without-Replacement Samples	79
6.4	Summary, Tips, and Warnings	83
7	Complex Surveys	85
7.1	Selecting a Stratified Two-Stage Sample	85
7.2	Estimating Quantiles	88
7.3	Computing Estimates from Stratified Multistage Samples	89
7.4	Univariate Plots from Complex Surveys	92
7.5	Scatterplots from Complex Surveys	95
7.6	Additional Code for Exercises	103
7.7	Summary, Tips, and Warnings	105
8	Nonresponse	107
8.1	How R Functions Treat Missing Data	107
8.2	Poststratification and Raking	108
8.3	Imputation	110
8.4	Summary, Tips, and Warnings	112
9	Variance Estimation in Complex Surveys	113
9.1	Replicate Samples and Random Groups	113
9.2	Constructing Replicate Weights	116
9.2.1	Balanced Repeated Replication	117
9.2.2	Jackknife	120
9.2.3	Bootstrap	122
9.2.4	Replicate Weights and Nonresponse Adjustments	124
9.3	Using Replicate Weights from a Survey Data File	126
9.4	Summary, Tips, and Warnings	127
10	Categorical Data Analysis in Complex Surveys	129
10.1	Contingency Tables and Odds Ratios	129
10.2	Chi-Square Tests	131
10.3	Loglinear Models	133
10.4	Summary, Tips, and Warnings	137
11	Regression with Complex Survey Data	139
11.1	Straight Line Regression with a Simple Random Sample	139
11.2	Linear Regression for Complex Survey Data	142
11.3	Using Regression to Compare Domain Means	145
11.4	Logistic Regression	149
11.5	Additional Resources and Code	151
11.6	Summary, Tips, and Warnings	152
12	Additional Topics for Survey Data Analysis	155
12.1	Two-Phase Sampling	155
12.2	Estimating the Size of a Population	157

<i>Contents</i>	ix
12.2.1 Ratio Estimation of Population Size	157
12.2.2 Loglinear Models with Multiple Lists	159
12.3 Small Area Estimation	161
12.4 Summary	162
A Data Set Descriptions	163
Bibliography	197
Index	205

Preface

R Companion for Sampling: Design and Analysis, Third Edition shows how to use the R statistical software environment to perform the calculations in the textbook *Sampling: Design and Analysis, Third Edition* (SDA) by Sharon L. Lohr. It is intended to be read in conjunction with SDA and is not a standalone text. The parallel book by Lohr (2022) shows how to perform the computations for the examples using SAS[®] software, and could be read together with this book and SDA to learn how to perform the analyses in each software package.

All code and data sets can be downloaded from any of the following websites:

<https://math.unm.edu/~luyan/rbook.html>

<https://www.sharonlohr.com>

<https://www.routledge.com/9781032135946>

The first two websites also contain additional programs, not discussed in this book, that you can adapt for some of the SDA exercises. The data sets used in this book have also been saved in R format in the contributed R package `SDAResources` (Lu and Lohr, 2021).

In this book, we give step-by-step guidance for using functions from base R and contributed packages to select samples and analyze the data sets discussed in Chapters 1–13 of SDA. The software, however, can do much more than analyze the examples presented in this book. You can find information on advanced capabilities for the `survey` and `sampling` contributed packages in the documentation for those packages by Lumley (2020) and Tillé and Matei (2021); the books and articles by Lumley (2004, 2010) and Tillé and Matei (2010) provide additional information about the packages. Goga (2018) gives an overview of using R for survey sampling.

For easy reference, the index at the back of the book gives page numbers for the examples in SDA. To locate the code and output for Example 2.5, for example, look up the subentry “Example 02.05” under “Examples in SDA” in the index. The book also gives code and suggestions for some of the exercises in SDA, and these are listed under index entry “Exercises in SDA.”

Each chapter ends with a summary section containing tips and warnings for the analyses discussed in that chapter. These provide ways of avoiding common survey data analysis errors and checking whether you did the analysis correctly.

Although prior experience with R is helpful, it is not needed to read this book. Chapter 1 tells how to obtain the software and do basic operations in R. It also lists resources for learning more about programming in R and tells how to obtain help.

This book makes use of functions that exist in base R and contributed packages, and does not discuss how to write R functions. One of R’s most valuable features, however, is the capacity for writing functions to carry out new tasks. Advanced R users may want to write their own functions to select samples or analyze data from a complex survey. When teaching

survey sampling to students who have R programming experience, we have sometimes asked them to write their own functions to carry out various sampling tasks. This helps solidify their knowledge of the material and allows them to do computations not available in existing functions. For example, we have asked students to write R functions to perform allocation for and analyze data from a stratified random sample, select a with-replacement unequal-probability sample using Lahiri's method, compute the Sen–Yates–Grundy estimate of the variance, simulate the sampling distribution of a statistic, and find empirical estimates of the coverage probability of a confidence interval for a biased estimator.

All code, data sets, and output in this book are provided for educational purposes only and without warranty. Base R does not contain functions for survey data, and this book relies heavily on contributed packages that have been developed. These packages are in widespread use and have been quality-checked by their authors and other users. We have verified that the calculations from the R functions used for the examples in this book agree with calculations by the formulas and with calculations performed in other survey software packages.

Other R packages may not be checked as carefully, however. Although R contributed packages undergo some consistency and functionality tests when they are submitted (see Wickham, 2015, for a description of checks that are performed), no central authority reviews the packages to make sure that the functions do what they claim or that the algorithms perform computations accurately. Most R contributed packages are not peer-reviewed, and you should be aware that some may contain errors.

The code and output in this book were developed using version 4.0.4 of R for Windows (R Core Team, 2021) and the versions of the packages listed in their respective bibliography entries, and all code in the book works with those versions. But R is a dynamic language, and the R Core Team and authors of contributed packages can change or remove functions at any time. Although most authors who revise a package try to avoid changes that will affect previously written code, functions in R are not guaranteed to be backward compatible—it is possible that R code you write today may not work the same way with future versions of the software. If backward compatibility is important to you—for example, if you will be using the same code to produce estimates each year for an annual survey—you may want to perform or check your computations in a package that is backward compatible, such as SAS software. If a function changes in a subsequent version of an R package, you can either:

- Read the documentation and change your code so that it works with the modified function, or
- Download and use the older version of the package. You can find previous versions on the package's web page under the heading "Old sources."

Acknowledgments. Many thanks to John Kimmel, our editor at CRC Press, for encouraging us to write this book, and to the CRC Press production team for all their support and help. We are grateful to Yves Tillé and Thomas Lumley for answering questions about the **sampling** and **survey** packages. Students in Yan Lu's sampling class at the University of New Mexico provided helpful suggestions for clarifying the material. We also want to thank Lynn Zhang for helping with the preparation of the **SDAResources** package.

Getting Started

The R statistical software environment is a powerful and flexible platform for performing statistical analyses. The basic package contains thousands of functions for computing statistics, and user-contributed packages for this open-source software provide thousands more. Advanced users can write their own functions to implement new methods for statistical analyses.

Best of all, the base R package and all user-contributed packages are available free of charge to anyone with an internet connection.

This chapter tells you how to obtain R software and contributed packages and introduces you to some basic R functions. It also shows you how to read data sets into R and save output and graphics produced while you are using the package.

Conventions used in this book. This book is intended to be read in conjunction with *Sampling: Design and Analysis, Third Edition* by Sharon L. Lohr, henceforth referred to as SDA. Many of the examples in this book refer to figures, tables, examples, or exercises in SDA. To avoid confusion, we refer to figures in SDA as “Figure x.x in SDA.” We refer to figures in *this* book as “Figure x.x” with no qualifier.

The names of external data files and programs, such as `agsrs.csv` and `ch02.R`, are in **typewriter font**, as are the names of R packages and code we type. Variable names, function names, and internal R data set names are in *italic type*.

Much of this book consists of R commands and output, set in light shaded boxes such as the following:

```
# This is a comment
# Enter data values into vector 'myvec'
myvec <- c(1,2,3,9,14,27,5,21,pi)
# Print the vector 'myvec'
myvec
## [1]  1.000000  2.000000  3.000000  9.000000 14.000000 27.000000  5.000000
## [8] 21.000000  3.141593
# Calculate summary statistics
summary(myvec)
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.00    3.00    5.00    9.46   14.00   27.00
```

To distinguish between the code and the output that is produced, a command that we typed into R is displayed flush against the left margin. The output from that command is preceded by `##`. Our comments are preceded by a single `#`. You can obtain the commands and comments (without the output) for all code in files `ch01.R`, `ch02.R`, etc., on the book website (see the Preface for the website address).

1.1 Obtaining the Software

The R system is available for FREE download from the Comprehensive R Archive Network at

<https://cran.r-project.org/>.

If R is not already installed on your system, download the package now. Click on the link for the operating system of your choice: Windows, Mac, or Linux. Install the package (you can do this for Windows or Mac by double-clicking on the icon after the download is finished). You should be able to click “Next” for all dialogs to finish the installation.

Good! Now you can use R directly or through the integrated development environment RStudio®. RStudio adds a larger number of statistical packages to R and provides a user-friendly graphical interface. You can obtain RStudio free of charge for Windows, Mac, or Linux from <https://rstudio.com>. As with R, you can install RStudio by double-clicking on the icon after download and can click on “Next” for all dialogs. If you have any trouble, numerous videos are available online that demonstrate the installation of R and RStudio.

Figure 1.1 is a screenshot of RStudio. It shows four windows:

1. Upper-left: Program editor window, where you will write code and programs. You can save the contents of this window so you have a record of what you have done.
2. Lower-left: R console, where you will run R commands.
3. Upper-right: The Workspace window tells you what you have created so far. Environment lists the variables in memory. History lists the commands you have submitted in the R console. The Tutorial window provides a tutorial for using R and RStudio.
4. Lower-right: Displays the files in the current working directory, plots you have produced, packages installed and whether they are loaded, and help for commands.

How to get things done in RStudio

1. Work out commands in the program editor window, and then either copy/paste the commands into the console or use Ctrl-Enter or click the icon “Run” to submit the current line or selection. You can open a new program editor window by clicking “R Script” under the “New File” icon in the upper left of the main RStudio toolbar (this can also be accessed through the File menu).
2. The program editor will keep a history of what you have done.
3. Make comments using # to help your future self recall what you did and why you did it.

Updating R and RStudio. If you already have R and RStudio installed, you may update them to the latest versions. You can either reinstall the latest version or (only for Windows) use the *updateR* function from contributed package *installr*. To update RStudio, select “Check for Updates” from the Help menu within the program.

1.2 Installing R Packages

The base R software contains few functions for analyzing survey data. In this book we rely on user-contributed packages that provide functions for selecting probability samples and for graphing and analyzing survey data.




FIGURE 1.1: Screenshot of RStudio.

The main contributed packages used in this book are:

- **survey** (Lumley, 2020): Used to calculate statistics from complex surveys with functions such as *svydesign*, *svymean*, and *svytotal*.
- **sampling** (Tillé and Matei, 2021): Used to select probability samples. The **sampling** package will also analyze data from complex samples, although in this book we primarily use the **survey** package for data analysis.
- **SDAResources** (Lu and Lohr, 2021): Includes all the data sets of SDA in R format and additional functions for analyzing and graphing probability samples.

Install these packages now. You only need to install each package once for a particular version of R (you may need to reinstall packages if you update to a new R version).

Copy and paste the following code into the R Console window and press [Enter]. Each package will be downloaded (along with package dependencies) and installed.

```

install.packages("survey")
install.packages("sampling")
install.packages("SDAResources")

```

You must load a package into the program with the *library* function each time you open a new R session. When they are needed, run (in this order):

```

library(survey)
library(sampling)
library(SDAResources)

```

1.3 R Basics

Getting help. There are many sources for help about R and RStudio. The web pages <https://www.r-project.org/> and <https://education.rstudio.com/> contain links to online tutorials and resources. For those who prefer a printed book, Horton and Kleinman (2015), Horton et al. (2018), Ismay and Kennedy (2019), and Kabacoff (2021) provide useful introductions to getting started with R and RStudio. The book by Wickham (2019) is an advanced guide to R.

The *help* function in R provides access to the documentation pages for R functions, data sets, and other objects. For example, if you want to get help with the function *mean*, type `help("mean")`, or type `help.search("mean")` into the R console to search the help system for documentation matching the given character string “mean”. You can also display the code of a function (for some packages) by typing its name. For example, typing `ls` in the R console window displays the code for the function *ls*, which lists the objects in the directory; to list the objects, type `ls()`.

For another example, suppose that you have not installed package *survey* yet, and you want to get help with the function *svymean*. Typing `help("svymean")` will show that “No documentation for ‘svymean’ in specified packages and libraries”. Typing `help.search("svymean")` will direct you to the page with information that function *svymean* is under package *survey*.

To find help for a package, for example, package *survey*, type `help(package = "survey")`. To find help for data *agpop*, type `help("agpop")`.

In RStudio, you can find help on a function by typing the function name into the search window of the Help menu.

Frequently used commands. Tables 1.1 and 1.2 list frequently used operators, functions, and commands on vectors, matrices, and data frames. Note that R is case-sensitive: *myvar*, *MYVAR*, and *Myvar* are all different variables.

The R functions we use to analyze survey data operate on data frames. A data frame is like a matrix, where each row contains the data values for an observation, and each column contains the data values for one variable. You can turn a matrix *A* into a data frame with the *as.data.frame* function.

```
A<-matrix(0,2,2) # create a 2*2 matrix with all 0s
A[1,1]<-2
A[1,2]<-5
A[2,1]<-1
A[2,2]<-4
colnames(A)<-c("y","x")
A
##           y x
## [1,]  2 5
## [2,]  1 4
A.data <- as.data.frame(A) # turn matrix A into a data frame A.data
A.data
##           y x
## 1  2 5
## 2  1 4
```


TABLE 1.1

Frequently used commands in R: Vectors, matrices, and mathematical functions.

Command	Purpose	Example	Output
Assignment and Comment			
<-	Assign value	a <- 1	1
=	Assign value	b = 1 + 2	3
#	Comment	# My comment	
Mathematical Functions			
+	Addition	1 + 1	2
-	Subtraction	1 - 1	0
*	Scalar multiplication	1*2	2
/	Division	1/2	0.5
^	Exponentiation	2^3	8
abs	Absolute value	abs(-3)	3
exp	Exponential function	exp(1)	2.718282
log	Natural log	log(2.718282)	1
sqrt	Square root	sqrt(4)	2
Vector/Matrix Operations			
c	Combine values into vector	vec1<-c(1,2,3) vec2<-c(4,5,6)	1 2 3 4 5 6
seq	Sequence or, when by=1	seq(from=1,to=10,by=2) 1:5	1 3 5 7 9 1 2 3 4 5
sequence	Vector of sequences	sequence(vec1)	1 1 2 1 2 3
rep	Replicate	rep(1,4) vec3<-rep(vec1,2)	1 1 1 1 1 2 3 1 2 3
cbind	Column bind	mat1<-cbind(vec1,vec2)	1 4 2 5 3 6
rbind	Row bind	mat2<-rbind(vec1,vec2)	1 2 3 4 5 6
length	Length of vector	length(vec1)	3
sort	Sort a vector	sort(vec3)	1 1 2 2 3 3
order	Indices to sort a vector	order(vec3)	1 4 2 5 3 6
unique	List unique objects	unique(vec3)	1 2 3
sum	Summation	sum(c(1,2,4))	7
prod	Product	prod(c(1,2,4))	8
%*%	Matrix multiplication	mat1 %*% mat2	17 22 27 22 29 36 27 36 45
t	Transpose	t(mat1)	1 2 3 4 5 6

1.4 Reading Data into R

The first step for using R to analyze data from a survey is to read the data into the system. There are four basic ways to do this.

Enter the data directly into R. You can use the *c* function to enter a vector directly from the R console. Type

```
myvec <- c(5,2,8,4)
```

TABLE 1.2

Frequently used commands in R: Extracting data elements and computing statistics.

Command	Purpose	Example	Output
Working with Data			
as.data.frame	Create data frame	A<-as.data.frame(mat1)	vec1 vec2 1 4 2 5 3 6
names(data)	Extract variable names	names(A)	"vec1" "vec2"
\$	Extract a column by name	A\$vec1	1 2 3
[, j]	Extract a column by number	A[,1]	1 2 3
[i ,]	Extract a row by number	A[1,]	1 4
nrow(data)	# of rows in <i>data</i>	nrow(A)	3
ncol(data)	# of columns in <i>data</i>	ncol(A)	2
head(data,n=i)	First i rows	head(A,n=2)	1 4 2 5
tail(data, n=i)	Last i rows	tail(A,n=2)	2 5 3 6
Statistics			
max/min	Maximum/Minimum of a vector or matrix	max(vec1) min(mat1)	3 1
quantile	Calculate quantiles	quantile(1:101,probs=c(0,0.25,0.5))	0% 25% 50% 1 26 51
mean	Sample mean	mean(vec1)	2
var	Sample variance	var(vec1)	1
cor	Correlation of two vectors	cor(vec1,vec2)	1

to store the vector of values (5, 2, 8, 4) in variable *myvec*. You can also use the *data.frame* function to enter data into a data frame from the R console:

```
# read matrix into data frame "mydata"
mydata <- data.frame(x = c(1,4,3,2), y = c(2,4,2,4), z = c(3,7,2,3))
# show the data
mydata
##   x y z
## 1 1 2 3
## 2 4 4 7
## 3 3 2 2
## 4 2 4 3
```

Read data from a comma-delimited (.csv) or text file. With a longer data set, it is often more convenient to store the data in an external file and then read it in through the data step. The following code will read data in a text file from a web page.

```
ex.data <- read.table(file="https://math.unm.edu/**/**/*.txt", header=T)
# Use header=T if there is a header
ex.data # print the data
```

Suppose the folder on your computer containing the data sets is C:\MyFilePath\. To read a csv file, use *read.csv*. Note that to specify a file path in R, you should replace each backslash ‘\’ with either a forward slash ‘/’ or a double backslash ‘\\’.

```
ex.data <- read.csv(file="C:/MyFilePath/exampladata.csv",header=F)
```

```
# header=F is the default
colnames(ex.data) <- c("y","x1","x2") #add column names
ex.data
```

Import data with RStudio. Use the “Import Dataset” dropdown from the “Environment” window (top right panel in RStudio). The import formats are grouped into 3 categories: text data, spreadsheet data, and statistical data.

- import data from webpage
Select “From Text (readr)”, enter URL address, and click “Import”.
- import data from local file
Select “From Text (base)” or “From Excel”, select a local file and click “Import”

Luraschi (2021) gives a detailed description of how to import data in RStudio.

Read R data sets. This is the easiest method of all—provided someone else has already saved the file as an R data set. All the data sets in SDA are loaded in R package *SDAResources*. Here is an example to read data *agpop*.

```
library(SDAResources)
data(agpop) # Load the agpop data set
N <- nrow(agpop)
N
## [1] 3078
head(agpop)
##               county state acres92 acres87 acres82 farms92 farms87 farms82
## 1 ALEUTIAN ISLANDS AREA   AK  683533  726596  764514      26      27      28
## 2     ANCHORAGE AREA     AK   47146   59297  256709     217     245     223
## 3     FAIRBANKS AREA     AK  141338  154913  204568     168     175     170
## 4         JUNEAU AREA     AK    210    214    127      8      8      12
## 5    KENAI PENINSULA AREA   AK   50810   85712   98035     93     119     137
## 6     AUTAUGA COUNTY      AL  107259  116050  145044     322     388     453
##  largef92 largef87 largef82 smallf92 smallf87 smallf82 region
## 1      14      16      20      6      4      1      W
## 2       9      10      11     41     52     38      W
## 3      25      28      21     12     18     25      W
## 4       0       0       0      5      4      8      W
## 5       9      18      17     12     18     19      W
## 6      25      32      32      8     19     17      S
```

1.5 Saving Output

Section 1.4 gave four methods for reading data into R. How do you save the output from the program? Here are several methods that will allow you to save the output or paste it into another document. Section 1.6 will show you how to incorporate R output directly into a $\text{\LaTeX}$ document.

Use the *sink* function to save console input and output. First, use *sink* to create a file in local drive. Next, run the R program. Last, close *sink*. Below is an example.

```
sink(file = "C:/MyOutputPath/output.txt",append = TRUE,
     type = c("output","message"))
```

```

cat("boot.samples", sep="\n") #add name of the R output
#R program
set.seed (244)
B <- 6
n <- 5
boot.samples <- matrix(sample(agpop$acres92, size=B*n, replace=TRUE),B,n)
boot.samples
sink(file = NULL) #close sink

```

The file C:\MyOutputPath\output.txt in the local drive appears as follows:

```

boot.samples
      [,1]      [,2]      [,3]      [,4]      [,5]
[1,]    831 1449976  98142 161724 138986
[2,]   86856 215222 245681 576468 80905
[3,]   96474  45214  51836 284888 302352
[4,]  137426  17392 141260 117768 177189
[5,]  342653 338529  73654 118291 297984
[6,]   50319   2531 223429 710546 216638

```

In some cases you may also want to save the R commands, so you can look at them later. You may use the function `source` to read in an R program with option `echo=TRUE` to include the R code, so that R output together with R code will be saved in `output.txt`.

```

sink(file = "MyOutputPath/output.txt",append = TRUE,
      type = c("output","message"))
source("MyInputFilePath/homework3.R", echo=TRUE)
sink(file = NULL)

```

Convert a matrix to $\text{\LaTeX}$ code in table format. In case you want to include a matrix as part of a $\text{\LaTeX}$ document, there is an easy way to convert the matrix to $\text{\LaTeX}$ code in table format. The following code, using the `stargazer` package, converts the matrix `boot.samples` to a $\text{\LaTeX}$ table.

```

# save output as latex
library(stargazer)
set.seed (244)
B <- 6
n <- 5
boot.samples <- matrix(sample(agpop$acres92, size=B*n, replace=TRUE),B,n)
# Now use stargazer to produce Latex commands for table
# Add header=FALSE to omit the initial comments in the Latex code output.
stargazer(boot.samples, digits = 2,header=FALSE)
##
## \begin{table}[!htbp] \centering
## \caption{}
## \label{}
## \begin{tabular}{@{\extracolsep{5pt}} ccccc}
## \\\[-1.8ex]\hline
## \hline \\\[-1.8ex]
## $831$ & $1,449,976$ & $98,142$ & $161,724$ & $138,986$ \\\
## $86,856$ & $215,222$ & $245,681$ & $576,468$ & $80,905$ \\\
## $96,474$ & $45,214$ & $51,836$ & $284,888$ & $302,352$ \\\
## $137,426$ & $17,392$ & $141,260$ & $117,768$ & $177,189$ \\\
## $342,653$ & $338,529$ & $73,654$ & $118,291$ & $297,984$ \\\
## $50,319$ & $2,531$ & $223,429$ & $710,546$ & $216,638$ \\\

```

```
## \hline \[-1.8ex]
## \end{tabular}
## \end{table}
```

Saving graphs. R and RStudio both allow you to save a graph using menu options. In standalone R, you can save a graph by choosing “Save as” from the File menu. In RStudio, select the Export dropdown from the plot panel (lower right-panel):

Plots → Export → Save as Image or Save as PDF.

Alternatively, you can specify files to save the image using functions such as *jpeg*, *png*, and *pdf*, which save the plot in an external file with the designated format. All of these functions have multiple options for sizing and formatting the graph. To save a graph as a pdf file, for example, first open the pdf file, then create the plot, and then close the file:

```
# Open a pdf file
pdf("~/MyOutputPath/rplot.pdf")
# Create a plot
boxplot(acres92/10^6 ~ region, xlab = "Region", ylab = "Millions of Acres",
        data = agpop)
# Close the pdf file
dev.off()
```

Or, to save as a jpeg file,

```
# save as jpeg file
jpeg("~/MyOutputPath/rplot.jpg",width=350,height=350)
boxplot(acres92/10^6 ~ region, xlab = "Region", ylab = "Millions of Acres",
        data = agpop)
dev.off()
```

Saving a data set. You can use function *write.table* or *write.csv* to save a data set to an external file. In the following, we will use data *classes* to create a new dataset in a long format called *classeslong*, and use *write.csv* to save the data in file *classeslong.csv*.

```
# read in data classes
data(classes)
# change to a long format by creating a record for each student
# create new data frame with each row repeated as many times as number of students
classeslong<-classes[rep(1:nrow(classes),times=classes$class_size),]
# add column of student ids, goes from 1 to number of students in each class
classeslong$studentid <- sequence(classes$class_size)
nrow(classeslong)
## [1] 647
head(classeslong)
##      class class_size studentid
## 1      1         44         1
## 1.1    1         44         2
## 1.2    1         44         3
## 1.3    1         44         4
## 1.4    1         44         5
## 1.5    1         44         6
# you can save classeslong in your local drive (or other file path) using write.csv
write.csv(classeslong, file=~ /classeslong.csv", row.names = FALSE)
```

1.6 Integrating R Output into L^AT_EX Documents

The **Sweave** system within R allows you to embed R code and output within L^AT_EX documents to generate a pdf file that includes narrative, graphics, code (if needed), and the results of the computations in R. The R package **knitr** serves as an engine for **Sweave** and allows you to generate dynamic reports with reproducible research (that is, someone else can produce the same results when given access to the data and code) using R. Help for **knitr** is available on the web (<https://yihui.org/knitr/>) and in the book by Xie (2015). We used the **knitr** package to produce this book.

Let's start with a simple example, using RStudio.

Step 1. Open a new Rnw script by clicking the “R Sweave” icon under the “New File” icon in the upper left of the main RStudio toolbar. Your first Rnw file will appear as in Figure 1.2. Save the file with extension **Rnw**—for example, you might name this file **test.Rnw**.




FIGURE 1.2: Example of a Rnw file.

Step 2. Weave the Rnw file using **knitr**. You will be asked to install package **knitr** if this is the first time you have used it.

Click Tools → Global options → Sweave → Weave Rnw files using → knitr.

Step 3. Click “Compile PDF” to generate a pdf file. The file **test.pdf** will appear in the same directory as **test.Rnw**, along with the L^AT_EX file, log file, and other auxiliary files.

Embedding R code within L^AT_EX. **Sweave** allows you to embed R code within L^AT_EX documents so you can incorporate statistical results in your pdf document. Type your L^AT_EX document in the program editor window in RStudio. To embed a “chunk” of R code in

knitr, start with `<< .. >>=` (all on one line) and close with `@`. Below is an illustration (we had to break the first line so it fit on the page, but you should leave it unbroken):

```
<<exercise0, echo=TRUE, size='footnotesize', include=TRUE, fig.width=6,
    fig.height=5, out.width='0.80\\textwidth'>>=
# Insert R code here!
x <- 1:10
y <- seq(from=20,to=2,by=-2)
x
y
x*y
cor(x,y)
# End the chunk with an @
@
```

This example, which has label *exercise0*, includes the following options to customize the R output.

echo = TRUE (print source code in the document) or FALSE (do not print source code in the document).

size = font size of the code text, if echoed. Options include ‘normalsize,’ ‘footnotesize,’ ‘large,’ and other font sizes in L^AT_EX.

include = TRUE if it is desired to include the chunk output in the document. If FALSE, the output is not produced, but the R code is still executed and the plots are generated.

fig.width = figure width in inches

fig.height = figure height in inches

out.width = scaling figures to fit on page (can be inches or cm, or relative to page size).

Here, the width is scaled to be 0.8 times the value of the *textwidth* in the L^AT_EX document.

To learn more about chunk options, please refer to https://yihui.org/knitr/options/#chunk_options.

Here is what is displayed when we include the code for *exercise0* in the document:

```
# Insert R code here!
x <- 1:10
y <- seq(from=20,to=2,by=-2)
x
## [1] 1 2 3 4 5 6 7 8 9 10
y
## [1] 20 18 16 14 12 10 8 6 4 2
x*y
## [1] 20 36 48 56 60 60 56 48 36 20
cor(x,y)
## [1] -1
# End the chunk with an @
```

Now, try a slightly more complicated example in `hw0.Rnw` (included on the book webpage) to see how to embed R chunks (including graphs) within latex, and compile it. You can see the pdf document that is produced by the example in file `hw0.pdf`.

Many excellent books are available for learning and using L^AT_EX. The books by Talbot (2012), Oetiker et al. (2021), and Wikibooks contributors (2021) may be downloaded free of charge.

1.7 Missing Data

Many survey data sets have observations that are missing. For example, the data file `agpop.csv`, from which the sample used in Chapter 2 is drawn, has missing data. As is common in data sets intended to be readable by multiple programs, a designated number is used to indicate that the data value is missing. In this data set, the value “-99” indicates that the data value is missing. This value must be recoded to the R symbol that indicates missing data before performing calculations. Otherwise, if, say, you want to calculate the mean of a variable, a function will treat all the “-99”s as if they were observations with the value -99 instead of missing values—this could lead to embarrassing results such as computing a negative value for the average number of acres per farm.

Missing values in R are coded by the symbol `NA`. If your data set has missing values that are coded as a number (such as -99), you should recode those to `NA` before starting your analysis. Missing data in the R data sets in package `SDAResources` (Lu and Lohr, 2021) have been recoded to `NA`, but the data sets for SDA that are in `.csv` format use other codes, which are described in Appendix A. If you want to read the data set `agpop.csv` into R from the `.csv` file (instead of loading it from `SDAResources`), you should recode the missing data as follows:

```
agpop1 <- read.csv(file="~/agpop.csv",header=TRUE)
# some variables have missing values, coded as -99
sum(agpop1==-99) # 59 missing values altogether
## [1] 59
# look at a row with missing data
agpop1[200,]
##               county state acres92 acres87 acres82 farms92 farms87 farms82
## 200 SAN FRANCISCO COUNTY   CA      7    -99     19      6      5      7
##      largef92 largef87 largef82 smallf92 smallf87 smallf82 region
## 200         0         0         0         6         5         7      W
agpop2 <- agpop1
# recode missing values to NA
agpop2[agpop1==-99] <- NA
# look at row with missing data again
agpop2[200,]
##               county state acres92 acres87 acres82 farms92 farms87 farms82
## 200 SAN FRANCISCO COUNTY   CA      7      NA     19      6      5      7
##      largef92 largef87 largef82 smallf92 smallf87 smallf82 region
## 200         0         0         0         6         5         7      W
# count missing values in recoded data
sum(is.na(agpop2))
## [1] 59
```

Different functions in R treat missing data in different ways. For many functions, the treatment of missing values can be specified using the *na.rm* argument. The `sum` function, for example, will compute the sum of the non-missing values in a vector if you include `na.rm=TRUE`; otherwise, it returns `NA` if there are missing values.

```
sum(agpop2$acres87)
## [1] NA
sum(agpop2$acres87,na.rm=TRUE)
## [1] 963466689
```


1.8 Summary, Tips, and Warnings

Tables 1.1 and 1.2 describe commonly used R functions for working with data. Table 1.3 lists other functions used in this chapter to read and write data. The **base**, **stats**, and **utils** packages are automatically included when you install R on your system.

TABLE 1.3

Functions used for Chapter 1.

Function	Package	Usage
library	base	Load an R package that you have installed on your system. You need to load a package each time you start a new R session.
sink	base	Send R output to an external file
source	base	Read and execute R commands from an external file
install.packages	utils	Install an R package
data	utils	Load a specified data set
read.table	utils	Read data from an external file into a data frame. You can specify the character that separates fields, row and column names, and whether the first line of the file gives the variable names.
read.csv	utils	This function is like read.table for an external file that is in comma-delimited (.csv) format.
write.table	utils	Write an R data set to an external file
write.csv	utils	Write an R data set to an external file in .csv format

Tips and Warnings

- R commands can be tricky, and sometimes the result of using functions on vectors and matrices may not be what you were expecting. When creating new variables or matrices, print out the first few rows or compute summary statistics to check that you have created the object you want.
- Use vector and matrix commands whenever possible when performing calculations with R. Although we occasionally use *for* loops in this book where this will make the code easier to understand, in general, it is much more efficient to perform operations on vectors.
- Many surveys contain missing data. Check how these are coded in the data set, and recode missing values to NA before starting your analysis.
- Sometimes multiple packages contain functions or data sets with the same names. For example, the **sampling** and **survival** packages (the **survival** package is loaded when you load the **survey** package) both have functions named *strata* and *cluster*. If you load **sampling** and then load **survival**, you will get the function from the **survival** package when you call *strata*—the function *strata* from the **sampling** package is masked. To avoid this problem, load the **sampling** package after loading the **survival** package or access the function from **sampling** by typing `sampling::strata`. For this book, we suggest loading the packages in the order given in Section 1.2.

Simple Random Sampling

Data from a simple random sample (SRS) can be analyzed using R functions that are designed for data that can be considered as independent and identically distributed, and an SRS can be selected using the R *sample* function. For other types of probability samples, however, you either need to write your own function to account for the survey design or employ functions that have been written by other R users in contributed packages. This chapter reviews how to select a sample and compute estimates from an SRS using functions in base R. It also introduces you to the *srswor* and *srsur* functions from the **sampling** package (Tillé and Matei, 2021) to select an SRS, with or without replacement; and the *svydesign*, *svymean*, and *svytotal* functions from the **survey** package (Lumley, 2020) to calculate the statistics.

All code in this chapter can be found in file `ch02.R` on the book website (see the Preface for the website address). The data sets are available from the book website and in the R package **SDAResources** (Lu and Lohr, 2021). The variables in the data sets are described in Appendix A.

In this and future chapters, load the following three packages before starting the computations. You installed these packages in Chapter 1. In the R console, type:

```
library(survey)
library(sampling)
library(SDAResources)
```

Before calculating statistics, let's first look at how to use R functions to select an SRS from a population.

2.1 Selecting a Simple Random Sample

Example 2.5 of SDA. *Selecting an SRS from a population.* SDA used a random number table to select an SRS of size 4 from a population of size 10. There are several options for selecting an SRS in R.

Using the *sample* function in base R. Base R contains the function *sample* that can be used to select an SRS. We can select an SRS (without replacement) of size 4 from a population of size 10 as follows:

```
# Set the seed for random number generation
set.seed(108742)
# Select an SRS of size n=4 from a population of size N=10 without replacement
srs4 <- sample(1:10, 4, replace = FALSE)
# Print the sample to see
# Can print an object by typing its name, or typing "print(object)"
```

```
srs4
## [1] 1 8 9 5
```

The first line of the code uses the function

```
set.seed(seed)
```

where *seed* is an integer that you supply to the function. If you want to be able to reproduce your sample later, call *set.seed* immediately before calling the function that generates a sample, and record the value of *seed* that you used. You will then get the same sample the next time you call the *sample* function with the same value of *seed*.¹ If you do not call *set.seed* during your R session, the starting point will be generated by the program.

For this example, the *sample* function provided a sample containing units 1, 8, 9, and 5. Running the sample function again, without using *set.seed*, gives a different sample. But when we reset the seed to the original value of 108742, we obtain the first sample {1, 8, 9, 5} again.

```
# Run again, without setting a new seed.
sample(1:10, 4, replace = FALSE)
## [1] 9 7 3 10
# Now go back to original seed.
set.seed(108742)
sample(1:10, 4, replace = FALSE)
## [1] 1 8 9 5
```

The *sample* function is called with

```
sample(x,size,replace=FALSE)
```

to select an SRS without replacement of *size* observations from the population in *x*. We sample 4 observations from the population in the vector [1, 2, ..., 10].

The *sample* function will also select a simple random sample with replacement by calling it with *replace* = TRUE. Unit “9” appears twice in the following with-replacement sample.

```
# Using the sample function to select an SRS with replacement
set.seed(101)
srswr4 <- sample(1:10, 4, replace = TRUE)
srswr4
## [1] 9 9 7 1
```

Using the *srswor* or *srswr* function from the sampling package. An alternative is to use function *srswor* or *srswr* to select an SRS without or with replacement, respectively. These are in the **sampling** package (Tillé and Matei, 2021), which you installed in Chapter 1. Now load the package and select a sample by calling:

```
srswor(n,N)
```

where *n* is the sample size and *N* is the population size.

```
# Load the sampling package if you have not already done so.
# library(sampling)
set.seed(1329)
# Select an SRS of size n=4 from a population of size N=10 without replacement.
s1<-srswor(4,10)
# List the units in the sample (the population units having s1=1).
```

¹Samples generated by the *sample* function in R versions 3.6.0 and later, however, will differ from samples generated by earlier versions of R because the *sample* function was revised in version 3.6.0 to fix a bug.

```
s1
## [1] 0 0 1 1 1 0 0 0 1 0
(1:10)[s1==1]
## [1] 3 4 5 9
```

The function *srswr* returns of vector of length N , with ones in the positions of the units selected for the sample. The sample in *s1* consists of units 3, 4, 5, and 9.

Function *srswr* for drawing a with-replacement sample is similar, but returns a vector containing the number of times each unit is in the sample.

```
# Select an SRS of size n=4 from a population of size N=10 with replacement.
set.seed(35882)
s2<-srswr(4,10)
# the selected units are 2 and 9
s2
## [1] 0 2 0 0 0 0 0 0 2 0
(1:10)[s2!=0]
## [1] 2 9
# number of replicates, units 2 and 9 both appear twice
s2[s2!=0]
## [1] 2 2
# can use the getdata function to extract the sample from data frame with population
popdf<-data.frame(popid=1:10)
getdata(popdf,s2)
## [1] 2 2 9 9
```

The *getdata* function in the **sampling** package extracts the sample from the population listing. For this example, since units 2 and 9 are each selected twice, *getdata* repeats each of these units twice.

Note that the *sample* function has the sample size as the second argument, while the *srswr* and *srswr* functions have the sample size as the first argument. When you type *srswr*(4,10), R assigns the values to variables in the order given in the function definition (see Usage in the help file for the function). If you want to assign values in a different order (or even if you want to use the same order but want to be able to read the call easily), name the variables when calling the function. All of the following are equivalent:

```
# Call a function using variable names
set.seed(35882)
srswr(4,10)
## [1] 0 2 0 0 0 0 0 0 2 0
set.seed(35882)
srswr(n=4,N=10)
## [1] 0 2 0 0 0 0 0 0 2 0
set.seed(35882)
srswr(N=10,n=4)
## [1] 0 2 0 0 0 0 0 0 2 0
```

Example 2.6 of SDA. In SDA, the sample in *agsrs.csv* was selected from the population *agpop.csv* using random numbers generated in a spreadsheet. In the following, we show how to use R code to select a different SRS of size 300 from the 3078 counties in data set *agpop*. The function *getdata* is used to extract the sampled units from the population data, and the first 6 observations are printed.

```
# Select a different sample of size 300 from agpop
# Load the SDAResources package containing all the data in SDA book
```

```
# library(SDAResources) # we comment this since we already loaded the package
data(agpop) # Load the data set agpop
N <- nrow(agpop)
N # 3078 observations
## [1] 3078
# Select an SRS of size n=300 from agpop
set.seed(8126834)
index <- srswor(300,N)
# each unit k is associated with index 1 or 0, with 1 indicating selection
index[1:10]
## [1] 0 0 0 1 0 0 0 0 0 0
# agsrs2 is an SRS with size 300 selected from agpop
# extract the sampled units from the data frame containing the population
agsrs2 <- getdata(agpop,index)
agsrs2 <- agpop[(1:N)[index==1],] # alternative way to extract the sampled units
head(agsrs2)
##           county state acres92 acres87 acres82 farms92 farms87 farms82
## 4      JUNEAU AREA   AK      210      214      127         8         8        12
## 30 DE KALB COUNTY   AL  210733  213440  221502      1894      2047      2228
## 38    HALE COUNTY   AL  167583  154581  179618       382       441       481
## 46    LEE COUNTY    AL   67962   79836  100949       336       402       407
## 50 MADISON COUNTY   AL  224370  235478  292873       871       977      1101
## 62 RUSSELL COUNTY   AL  112620  143568  141048       213       276       314
##   largef92 largef87 largef82 smallf92 smallf87 smallf82 region
## 4         0         0         0         5         4         8        W
## 30        13         5         6        114        133       168        S
## 38        38        33        39         12         22        17        S
## 46        10        10        20         15         22        20        S
## 50        59        59        61         46         76        89        S
## 62        25        30        33         14         14        25        S
```

The functions *sample*, *srswor*, and *srsur* select a sample but do not provide sampling weights. After drawing the sample, you need to create a variable of sampling weights for the sampled units. For *agsrs2*, the weight variable *sampwt* has the value 3078/300 for all units.

```
# Create the variable of sampling weights
n <- nrow(agsrs2)
agsrs2$sampwt <- rep(3078/n,n)
# Check that the weights sum to N
sum(agsrs2$sampwt)
## [1] 3078
```

2.2 Computing Statistics from a Simple Random Sample

In this section, we look at two ways of computing estimates from an SRS: by using standard functions in R to calculate estimates from the formulas and by using functions in the **survey** package. We shall use the **survey** package to calculate estimates for the other chapters of this book, but show how to compute estimates for an SRS using the formulas so you can see that the two methods give the same numbers. Since R is a highly flexible language, you can also write your own functions to compute estimates using the formulas.

Examples 2.6, 2.7, and 2.11 of SDA. This section analyzes variables in data set `agsrs.csv`, described in Example 2.6 of SDA and in Appendix A of this book. The primary response variable for these examples is `acres92` (acreage devoted to farms in 1992).

First, we draw a histogram of the variable `acres92`, shown in Figure 2.1. The optional argument `breaks=20` tells R to use 20 bins for the histogram. The `col` (specifies color of bars), `xlab` (gives label for the x-axis), and `main` (specifies the title for the graph) arguments are also optional but make the picture look nicer.

```
# Draw a histogram
hist(agsrs$acres92,breaks=20,col="gray",xlab="Acres devoted to farms, 1992",
     main="Histogram: Number of acres devoted to farms, 1992")
```




FIGURE 2.1: Histogram of farm acreage in 1992 (data *agsrs*).

Using the formulas in SDA. Statistics from an SRS can be computed using functions supplied in the base R package with the formulas in SDA. Functions such as `t.test` will compute confidence intervals without a finite population correction (fpc).

```
# Base R functions such as t.test will calculate statistics for an SRS,
# but without the fpc.
t.test(agsrs$acres92)
##
##  One Sample t-test
##
## data:  agsrs$acres92
## t = 14.975, df = 299, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##  258749.6 337044.5
## sample estimates:
## mean of x
##  297897
```

The confidence interval from the *t.test* function is wider than that in Example 2.11 of SDA, however. Standard errors and confidence intervals that incorporate the fpc can be calculated directly with the formulas. If you are familiar with writing functions in R, you could also write your own function to do these calculations.

```
# Calculate the statistics by direct formula
n <- length(agsrs$acres92)
ybar <- mean(agsrs$acres92)
ybar
## [1] 297897
hatvybar<-(1-n/3078)*var(agsrs$acres92)/n
seybar<-sqrt(hatvybar)
seybar
## [1] 18898.43
# Calculate confidence interval by direct formula using t distribution
Mean_CI <- c(ybar - qt(.975, n-1)*seybar, ybar + qt(.975, n-1)*seybar)
names(Mean_CI) <- c("lower", "upper")
Mean_CI
##      lower      upper
## 260706.3 335087.8
# To obtain estimates for the population total,
# multiply each of ybar, seybar, and Mean_CI by N = 3078
seybar*3078
## [1] 58169381
Mean_CI*3078
##      lower      upper
## 802453859 1031400361
# Calculate coefficient of variation of mean
seybar/ybar
## [1] 0.06343948
```

Using functions in the *survey* package. Most of the statistics discussed in Chapter 2 can be computed in R using functions *svydesign*, *svymean*, and *svytotal* from the *survey* package (Lumley, 2020).

The data set *agsrs* does not contain a variable of sampling weights, so we need to create one. Also define the variable *lt200k*, which takes on value 1 if *acres92* < 200,000 and the value 0 if *acres92* ≥ 200,000. The mean of variable *lt200k* estimates the proportion of farms that have fewer than 200,000 acres.

```
# Create the variable of sampling weights
n <- nrow(agsrs)
agsrs$sampwt <- rep(3078/n,n)
# Create variable lt200k
agsrs$lt200k <- rep(0,n)
agsrs$lt200k[agsrs$acres92 < 200000] <- 1
# look at the first 10 observations with column 3 (acres92) and column 17 (lt200k)
agsrs[1:10,c(3,17)]
##      acres92 lt200k
## 1    175209      1
## 2    138135      1
## 3     56102      1
## 4    199117      1
## 5     89228      1
## 6     96194      1
## 7     57253      1
```



```
## 8    210692    0
## 9    78498    1
## 10   219444    0
```

Now specify the survey design with function *svydesign*. The function has numerous optional arguments, so we call it using the variable names in the arguments. The arguments used for an SRS are:

id In general survey designs, *id* specifies the cluster identifiers. For an SRS we use `id = ~1`, which tells *svydesign* that there is no clustering. The tilde (`~`) is used by R to specify a formula, and the syntax for formulas will be discussed in later chapters.

weights Names the variable in the data frame that contains the sampling weights. The weights argument can actually be omitted for calculating means in an SRS (the function will calculate weights from the *fpc* argument if it is supplied, or set all weights equal to 1 if neither *weights* nor *fpc* is included), but it is good to get in the habit of using a weight variable, so we include it here.

fpc Information for calculating the finite population correction. For an SRS, we can use `fpc = rep(N, n)` with *N* the population size and *n* the sample size. For *agsrs*, *N* = 3078, *n* = 300, and `fpc = rep(3078, 300)`.

data Name of the data frame containing the variables to be analyzed.

```
# If you did not load the survey package at the beginning of the chapter, do it now
# library(survey)
# Specify the survey design.
# This is an SRS, so the only design features needed are the weights
#   or information used to calculate the fpc.
dsrs <- svydesign(id = ~1, weights = ~sampwt, fpc = rep(3078,300), data = agsrs)
dsrs
## Independent Sampling design
## svydesign(id = ~1, weights = ~sampwt, fpc = rep(3078, 300), data = agsrs)
```

When we print *dsrs*, we are told that this is an “Independent Sampling design”—that is, there is no stratification or clustering. We will come back to the function *svydesign* in later chapters as we encounter other survey designs.

Now that the survey design has been specified, we can calculate estimated means and totals using the *svymean* and *svytotal* functions. For each of these, the first argument contains the name(s) of the variable(s) to be analyzed, and the second argument is the name of the design object that was created by *svydesign*. The function *confint* will construct a 95% confidence interval (you can specify other confidence levels with the optional *level* argument, but we omit this since we always use 95% intervals in this book) using a *t* distribution with *df* degrees of freedom. If you do not specify the *df*, a normal distribution will be used for the confidence intervals. For an SRS, the *t* distribution has *n* − 1 *df*, where *n* is the sample size.

```
# Calculate the mean for acres92 and its standard error using the svymean function.
smean <- svymean(~acres92,dsrs)
smean
##           mean      SE
## acres92 297897 18898
# Use the confint function to compute a 95% confidence interval from
# the information in smean, df = n-1 = 300-1 = 299
confint(smean, df=299)
##           2.5 %    97.5 %
```

```
## acres92 260706.3 335087.8
# Repeat these steps with the svytotal function to obtain estimated totals.
stotal <- svytotal(~acres92,dsrs)
stotal
##           total           SE
## acres92 916927110 58169381
confint(stotal, df=299)
##           2.5 %           97.5 %
## acres92 802453859 1031400361
# Calculate the CV of the mean
SE(smean)/coef(smean)
##           acres92
## acres92 0.06343948
# or
smean<-as.data.frame(smean)
smean[[2]]/smean[[1]]
## [1] 0.06343948
```

You can analyze multiple variables at a time by putting them in a formula. The following code estimates the population means for variables *acres92* and *lt200k*.

```
# Estimate population means for multiple variables
agsrs_means <- svymean(~acres92+lt200k,dsrs)
agsrs_means
##           mean           SE
## acres92 297897.05 18898.4344
## lt200k    0.51    0.0275
confint(agsrs_means, df=299)
##           2.5 %           97.5 %
## acres92 2.607063e+05 3.350878e+05
## lt200k  4.559508e-01 5.640492e-01
```

Estimating proportions from an SRS. For a binary numeric variable (taking on values 0 or 1), the estimated proportion is the mean of the variable, and the proportion of the population having *lt200k* = 1 is the mean of variable *lt200k*. From the above output, we can see that the value of $\hat{p} = 0.51$ is the estimated proportion where *lt200k* takes on the value 1. The standard error is 0.0275, and a 95% confidence interval of p is [0.456, 0.564].

Sometimes you want to estimate the proportion of the population that falls in each of multiple categories. The variable *region* in the *agsrs* data describes the census region for each county in the sample and takes on the values “NE,” “NC,” “S,” and “W”. Running *svymean* with the variable *region* gives the estimated proportion in each category.

```
# Analyzing a categorical variable that is coded as characters
# First, display the category names and counts
table(agsrs$region)
##
##  NC  NE   S   W
## 107  24 130  39
# Find the estimated proportions in each category
region_prop <- svymean(~region,dsrs)
region_prop
##           mean           SE
## regionNC 0.35667 0.0263
## regionNE 0.08000 0.0149
## regionS  0.43333 0.0272
```

```
## regionW 0.13000 0.0185
confint(region_prop,df=299)
##           2.5 %    97.5 %
## regionNC 0.30487557 0.4084578
## regionNE 0.05066780 0.1093322
## regionS  0.37975605 0.4869106
## regionW  0.09363889 0.1663611
region_total <- svytotal(~region,dsrs)
region_total
##           total      SE
## regionNC 1097.82 81.005
## regionNE  246.24 45.878
## regionS  1333.80 83.799
## regionW   400.14 56.872
confint(region_total,df=299)
##           2.5 %    97.5 %
## regionNC  938.4070 1257.2330
## regionNE  155.9555  336.5245
## regionS  1168.8891 1498.7109
## regionW   288.2205  512.0595
```

Numeric and categorical variables. Numeric variables are variables for which you want to calculate statistics such as means (for example, *acres92* is a numeric variable). Categorical variables are those for which the values represent categories. Region is a categorical variable. We want to estimate the proportion of the population in each region, but we cannot calculate an “average” region. Here, *region* is automatically recognized as a categorical variable because it contains characters other than numbers.

Some surveys code categories as numbers; be careful to treat such variables as categorical rather than numeric. For example, a survey variable *haircolor* might take values 1–6, where 1 represents black, 2 represents brown, 3 represents blond, 4 represents red, 5 represents bald, and 6 represents other. You can calculate the mean for the variable *haircolor*, but it has no meaning. If you want to estimate the population proportion with each hair color, you can either (1) define binary variables for each category, for example, *redhair* = 1 if *haircolor* = 4 and 0 otherwise and find the mean of each variable, or (2) declare the variable *haircolor* to be categorical.

In R, you specify that a variable is categorical with the *factor* function. You can either declare the variable to be a factor variable in the data set or in the function call of *svymean*.

```
# Analyzing a categorical variable that is coded as numbers
# First, analyze lt200k as a numeric variable (works only if all values are 0 or 1)
# This gives the mean of variable lt200k, which is the proportion with lt200k = 1.
svymean(~lt200k,dsrs)
##           mean      SE
## lt200k 0.51 0.0275
# Now, analyze lt200k as a factor variable. This gives the proportion
# in each category
svymean(~factor(lt200k),dsrs)
##           mean      SE
## factor(lt200k)0 0.49 0.0275
## factor(lt200k)1 0.51 0.0275
```

2.3 Additional Code for Exercises

This section contains additional code and references to functions used in three of the exercises in Chapter 2 of SDA. Some of these make use of advanced features of R; you should skip this section if you are new to R.

Exercise 2.27 of SDA asks you to estimate the sampling distribution of $\bar{y}$ by repeatedly taking samples of size n with replacement from the sample in *agsrs*, where y is the variable *acres92*. The following code constructs the histogram of the statistics from the bootstrap replicates that is shown in Figure 2.2. We use the *apply* function to calculate the mean value of *acres92* from each bootstrap replicate.

```
# Calculating bootstrap means for Exercise 2.27 in SDA
set.seed (244)
B = 1000
n = length(agsrs$acres92)
boot.samples = matrix(sample(agsrs$acres92, size = B * n, replace = TRUE),B, n)
boot.statistics = apply(boot.samples, 1, mean)
hist(boot.statistics, main="Estimated Sampling Distribution of ybar",
     xlab="Mean of acres92 from Bootstrap Replicate",
     col="gray",border="white")
```




FIGURE 2.2: Histogram of means from bootstrap replicates.

Exercise 2.32 of SDA. The function *srswor1* in the **sampling** package will select an SRS using the algorithm in this exercise.

Exercise 2.34 of SDA. The function *binom.test* will calculate Clopper-Pearson confidence intervals for an SRS. For complex surveys, the function *svyciprop* in the **survey** package will calculate a variety of asymmetric confidence intervals for a proportion (see Section 3.4).

2.4 Summary, Tips, and Warnings

Table 2.1 lists the major functions used in this chapter to select or analyze data from an SRS. The base, graphics, stats, and utils packages are automatically included when you install R on your system.

TABLE 2.1

Functions used for Chapter 2.

Function	Package	Usage
sample	base	Select a simple random sample with or without replacement
mean	base	Calculate the mean of a vector
var	base	Calculate the variance of a vector
table	base	Compute the counts in each category of a vector
factor	base	Convert a vector to a factor object
set.seed	base	Initialize the seed for the random number generator
hist	graphics	Draw a histogram of data from an SRS
qt	stats	qt(.975,df) calculates the 0.975 quantile of a t distribution with the specified df
t.test	stats	Calculate a t confidence interval for an SRS (assumed to be with-replacement)
confint	stats	Calculate confidence intervals
data	utils	Loads the data set enclosed in parentheses; <code>data()</code> lists the available data sets
srswor	sampling	Select a simple random sample without replacement
srswr	sampling	Select a simple random sample with replacement
getdata	sampling	Extract the sampled data from a data frame containing the population units
svydesign	survey	Specify the survey design, for example, SRS
svymean	survey	Calculate mean and standard error of mean (if the variable is numeric), or proportion in each category (if variable is categorical)
svytotal	survey	Calculate total and standard error of total

Tips and Warnings

- If you want to be able to draw the same sample from a population at a later date, use the *set.seed* function before calling the sample-generating function.
- Although the *weights* argument is optional in the *svydesign* function for an SRS, it is a good practice to include it for computing estimates in *svymean* or *svytotal*. Most data sets from complex surveys come with weights that must be used to compute estimates correctly.
- When computing weights for an SRS, check that the sum of the weights equals the population size.
- If you are analyzing categorical variables with the *svymean* or *svytotal* functions, make sure to declare them to be *factor* variables. Otherwise, the functions will treat them as numeric variables and calculate the mean instead of computing the proportion in each category.

3

Stratified Sampling

In a stratified random sample, the population is divided into subgroups called strata. An SRS is selected independently from each stratum. In this chapter, we look at methods for allocating and selecting stratified random samples using functions in base R and the **sampling** package (Tillé and Matei, 2021). We then discuss the usage of functions *svydesign*, *svymean*, *svytotal*, and *svyby* from the **survey** package (Lumley, 2020) in a stratified random sample.

All code in this chapter can be found in file **ch03.R** on the book website. As always, load the packages **survey**, **sampling**, and **SDAResources** before starting your work.

3.1 Allocation Methods

Data *agpop* contains a stratum variable *region* that describes the census region for each county in the population and takes on the values North Central (NC), Northwest (NE), South (S), and West (W). The following code calculates the population counts (N_h) for the variable *region* with the *table* function.

```
data(agpop) # load the data set
names(agpop) # list the variable names
## [1] "county" "state" "acres92" "acres87" "acres82" "farms92"
## [7] "farms87" "farms82" "largef92" "largef87" "largef82" "smallf92"
## [13] "smallf87" "smallf82" "region"
head(agpop) # take a look at the first 6 obsns
##               county state acres92 acres87 acres82 farms92 farms87 farms82
## 1 ALEUTIAN ISLANDS AREA AK  683533  726596  764514      26      27      28
## 2 ANCHORAGE AREA AK    47146   59297  256709     217     245     223
## 3 FAIRBANKS AREA AK   141338  154913  204568     168     175     170
## 4 JUNEAU AREA AK      210     214     127       8       8      12
## 5 KENAI PENINSULA AREA AK   50810   85712   98035      93     119     137
## 6 AUTAUGA COUNTY AL  107259  116050  145044     322     388     453
##   largef92 largef87 largef82 smallf92 smallf87 smallf82 region
## 1      14       16      20        6        4        1      W
## 2       9       10      11       41       52       38      W
## 3      25      28      21       12       18       25      W
## 4       0       0       0        5        4        8      W
## 5       9      18      17       12       18       19      W
## 6      25      32      32        8       19       17      S
nrow(agpop) #number of rows, 3078
## [1] 3078
unique(agpop$region) # take a look at the four regions, NC, NE, S, W
## [1] "W" "S" "NE" "NC"
table(agpop$region) # number of counties in each stratum
```

```
##
##   NC   NE    S    W
## 1054  220 1382  422
```

We can use the information about *region* to allocate a stratified sample.

Proportional allocation. With proportional allocation, the stratum sample sizes are proportional to the population stratum sizes N_h . A proportional allocation is easy to calculate in R; simply multiply N_h/N by the desired sample size. For example, region NC has 1054 counties and the population has 3078 counties. For a sample with $n = 300$, proportional allocation will select $300 \cdot 1054/3078 = 103$ counties from region NC. The values in *propalloc* are fractions, so we round these to the nearest integers to obtain the sample size.

```
popsiz<= table(agpop$region)
propalloc <- 300*popsiz/sum(popsiz)
propalloc
##
##      NC      NE      S      W
## 102.7290  21.4425 134.6979  41.1306
# Round to nearest integer
propalloc_int <- round(propalloc)
propalloc_int
##
##   NC   NE    S    W
## 103  21 135  41
sum(propalloc_int) # check that stratum sample sizes sum to 300
## [1] 300
```

Neyman allocation. For Neyman allocation, you need to provide additional information about the stratum variances. Sometimes you have information about a variable that is related to key survey responses from the sampling frame, or sometimes you have information on variances from a pilot study or from similar surveys that have been done. In other cases, you may need to make a conjecture about the stratum variances.

In the following example, we assume that the survey planner does not have the true population variances available, and we enter conjectures for the relative variances of the strata. For example, the variance in the West is set at twice the variance for the South. Using the *popsiz* vector that was calculated in the previous code, we have:

```
stratvar <- c(1.1,0.8,1.0,2.0)
# Make sure the stratum variances in stratvar are in same
# order as the table in popsiz
neymanalloc <- 300*(popsiz*sqrt(stratvar))/sum(popsiz*sqrt(stratvar))
neymanalloc
##
##      NC      NE      S      W
## 101.07640  17.99204 126.36327  54.56828
neymanalloc_int <- round(neymanalloc)
neymanalloc_int
##
##   NC   NE    S    W
## 101  18 126  55
sum(neymanalloc_int)
## [1] 300
```


Optimal allocation. Optimal allocation can be done similarly by defining costs or relative costs for sampling in each stratum.

```
relcost <- c(1.4,1.0,1.0,1.8)
# Make sure the relative costs in relcost are in same
# order as the table in popsize
optalloc <- 300*(popsize*sqrt(stratvar/relcost))/sum(popsize*sqrt(stratvar/relcost))
optalloc
##
##      NC      NE      S      W
## 94.75776 19.95766 140.16833 45.11626
optalloc_int <- round(optalloc)
optalloc_int
##
##  NC  NE   S   W
##  95  20 140  45
sum(optalloc_int)
## [1] 300
```

Table 3.1 summarizes the results of these three allocation methods for the *agpop* population. Of course, the Neyman and optimal allocations are only optimal under the assumed variances and costs used in the calculations. If those variances or costs are wrong, then these allocations will not be optimal for the variable of interest. And an allocation that is optimal for one response variable might not be optimal for another.

TABLE 3.1

Proportional, Neyman, and optimal allocation for the four regions.

Number of Counties in Stratum	NC	NE	S	W	Total
Population	1054	220	1382	422	3078
Sample with proportional allocation	103	21	135	41	300
Sample with Neyman allocation	101	18	126	55	300
Sample with optimal allocation	95	20	140	45	300

Other allocation methods. The sample sizes specified by the proportional, Neyman, and optimal methods are just guidelines. You can set the stratum sample sizes to any values that meet your research needs. For example, if you want to have high precision for comparing stratum means, you may want to select the same number of observations from each stratum.

There are other functions in R that you can use for allocation with stratified data such as functions *strAlloc* from the *PracTools* package (Valliant et al., 2020) and *optiallo* from the *optimStrat* package (Bueno, 2020). The package *SamplingStrata* (Barcaroli, 2014; Barcaroli et al., 2020) provides R functions for determining the optimal stratification and allocation that will achieve predetermined precisions for multiple y variables. For example, you can use the package to design a stratification that will ensure that the coefficients of variation for five key variables do not exceed 0.05. The package *stratification* (Bailargeon and Rivest, 2011; Rivest and Baillargeon, 2017) contains functions for determining stratum boundaries when the stratifying variable is continuous.

3.2 Selecting a Stratified Random Sample

The sample in Example 3.2 of SDA was selected using a spreadsheet but let's look at how to select a similar sample using R, with the *sample* and *strata* functions (this will, of course, give a different sample than obtained in Example 3.2 of SDA).

Using the *sample* function in base R. As we have discussed in Chapter 2, the *sample* function can be used to select an SRS. To select a stratified random sample, we select an SRS independently from each stratum.

Data `agpop.csv` contains a stratum variable *region* that describes the census region for each county in the population. In the following example, we use the proportional allocation from Table 3.1 to divide the $n = 300$ units among the four strata, that is, selecting 103, 21, 135, and 41 counties from regions NC, NE, S, and W, respectively.

```
# Select an SRS without replacement from each region with proportional allocation
# with total size n=300
regionname <- c("NC","NE","S","W")
# Make sure sampsize has same ordering as regionname
sampsize <- c(103,21,135,41)
# Set the seed for random number generation
set.seed(108742)
index <- NULL
for (i in 1:length(sampsize)) {
  index <- c(index,sample((1:N)[agpop$region==regionname[i]],
                        size=sampsize[i],replace=F))
}
strsample<-agpop[index,]
# Check that we have the correct stratum sample sizes
table(strsample$region)
##
##  NC  NE   S   W
## 103  21 135  41
# Print the first six rows of the sample to see
strsample[1:6,]
##
##               county state acres92 acres87 acres82 farms92 farms87
## 1316          ISANTI COUNTY    MN  131563  142998  153003    680    817
## 2034          DEFIANCE COUNTY    OH  196759  206905  210781    830    987
## 864           ATCHISON COUNTY    KS  245099  233619  234730    686    694
## 553           DES MOINES COUNTY   IA  192467  210843  224770    681    753
## 1738           DUNN COUNTY       ND 1352738 1358843 1397141    650    733
## 1325 LAKE OF THE WOODS COUNTY    MN  103665  118959  119296    176    222
##
## farms82 largef92 largef87 largef82 smallf92 smallf87 smallf82 region
## 1316    947    18    14    8    14    26    34    NC
## 2034   1033    25    20    18    40    50    50    NC
## 864    768    55    42    41    48    48    65    NC
## 553    815    33    30    24    56    56    72    NC
## 1738    697   358   368   361    19    13    34    NC
## 1325    230    30    35    26    4    4    1    NC
```

This simple code used a *for* loop to select an SRS from each stratum (defined by the subset having *region* equal to the stratum name) in turn; alternatively, one could use the *apply* function, or write a custom R function, to do this without looping. The vector containing the sample sizes must be in the same order as the vector giving the stratum names.

Using the *strata* function from the sampling package. An alternative is to use function *strata* to select a stratified random sample. This function is in the *sampling* package (Tillé and Matei, 2021), which you installed in Chapter 1.

First, sort the data by the stratification variable *region* before selecting the sample. Next, call the *strata* function with sorted data *agpop2* and the stratification variable *region* with first argument *agpop2* and second argument *stratanames="region"*. You can also use a vector of variables to define the strata, such as *stratanames=c("A","B")*, if the strata are formed from multiple variables. Add the information on number of counties to be selected within each stratum by *size=c(103,21,135,41)* in the *strata* function. Finally, choose the method to select the sample within each stratum; for this chapter we use either SRS without replacement (*method="srswor"*) or SRS with replacement (*method="srswr"*).

```
# Sort the population by stratum
agpop2<-agpop[order(agpop$region),]
# Use the strata function to select the units for the sample
# Make sure size argument has same ordering as the stratification variable
index2<-strata(agpop2,stratanames=c("region"),size=c(103,21,135,41),
               method="srswor")
table(index2$region) # look at number of counties selected within each region
##
##  NC  NE   S   W
## 103  21 135  41
head(index2)
##      region ID_unit      Prob Stratum
## 2      NC      2 0.09772296      1
## 9      NC      9 0.09772296      1
## 27     NC     27 0.09772296      1
## 36     NC     36 0.09772296      1
## 42     NC     42 0.09772296      1
## 43     NC     43 0.09772296      1
strsample2<-getdata(agpop2,index2) # extract the sample
head(strsample2)
##      county state acres92 acres87 acres82 farms92 farms87 farms82
## 526  ADAMS COUNTY   IA  239800  243607  254071    643    688    737
## 533  BREMER COUNTY   IA  236668  235086  250402   1058   1140   1287
## 551  DECATUR COUNTY   IA  261494  278714  300684    648    715    769
## 560  FREMONT COUNTY   IA  302352  308796  306786    596    719    771
## 566  HARDIN COUNTY   IA  332358  337990  355823    986   1065   1208
## 567  HARRISON COUNTY   IA  399155  387190  408601    919   1024   1192
##      largef92 largef87 largef82 smallf92 smallf87 smallf82 region ID_unit
## 526      38      32      21      40      50      33      NC      2
## 533      25      18      11      96     116     109      NC      9
## 551      52      54      56      20      34      37      NC     27
## 560      91      72      51      37      59      50      NC     36
## 566      56      36      42      90     115     132      NC     42
## 567      88      62      51      60      60      66      NC     43
##      Prob Stratum
## 526 0.09772296      1
## 533 0.09772296      1
## 551 0.09772296      1
## 560 0.09772296      1
## 566 0.09772296      1
## 567 0.09772296      1
```

The data frame *index2* contains the stratum variables, the identifiers of the units selected to be in the sample, and the inclusion probability for each unit in the sample. The function *getdata* then extracts the sampled units from the population data.

The *strata* function gives the inclusion probabilities for the sample units but not the weights. You can calculate the sampling weights by taking the reciprocal of the inclusion probabilities. When calculating weights for a stratified random sample, always check that the weights sum to the stratum population sizes. If they do not sum to the stratum population sizes, you have made a mistake somewhere in the weight calculations.

```
# Calculate the sampling weights
# First check that no probabilities are 0
sum(strsample2$Prob<=0)
## [1] 0
strsample2$sampwt<-1/strsample2$Prob
# Check that the sampling weights sum to the population sizes for each stratum
tapply(strsample2$sampwt,strsample2$region,sum)
##   NC   NE    S    W
## 1054 220 1382 422
```

3.3 Computing Statistics from a Stratified Random Sample

Examples 3.2 and 3.6 of SDA. As in Chapter 2, function *svydesign* from the *survey* package can be used to enter the stratified random sample information, and functions *svymean* and *svytotal* will calculate estimated means and totals from a stratified random sample. The data set *agstrat* is a stratified random sample taken from the population data *agpop* with proportional allocation. First, let's look at the data.

```
data(agstrat)
names(agstrat) # list the variable names
## [1] "county" "state" "acres92" "acres87" "acres82" "farms92"
## [7] "farms87" "farms82" "largef92" "largef87" "largef82" "smallf92"
## [13] "smallf87" "smallf82" "region" "rn" "strwt"
agstrat[1:6,1:8] # take a look at the first 6 obsns from columns 1 to 8
##   county state acres92 acres87 acres82 farms92 farms87 farms82
## 1 PIERCE C NE 297326 332862 319619 725 857 865
## 2 JENNINGS IN 124694 131481 139111 658 671 751
## 3 WAYNE CO OH 246938 263457 268434 1582 1734 1866
## 4 VAN BURE MI 206781 190251 197055 1164 1278 1464
## 5 OZAUKEE WI 78772 85201 89331 448 483 527
## 6 CLEARWAT MN 210897 229537 213105 583 699 693
nrow(agstrat) # number of rows, 300
## [1] 300
unique(agstrat$region) # take a look at the four regions, NC, NE, S, W
## [1] "NC" "NE" "S" "W"
table(agstrat$region) # number of counties in each stratum
##
## NC NE S W
## 103 21 135 41
# check that the sum of the weights equals the population size
sum(agstrat$strwt) #3078
## [1] 3078
```

Figure 3.1 gives a boxplot for variable *acres92* (scaled to millions of acres). We can see that the West region has the highest median and largest variability, while the Northeast region has the lowest median and smallest variability. Note that we can use the *boxplot* function in the following code because an SRS is taken within each stratum (and, because of proportional allocation, the sample is approximately self-weighting); for other designs, one should incorporate the weights into the plot as shown in Chapter 7.

```
boxplot(acres92/10^6 ~ region, xlab = "Region", ylab = "Millions of Acres",
        data = agstrat)
# notice the large variability in western region
```




FIGURE 3.1: Boxplot of 1992 acreage by region (data *agstrat*).

Now let's calculate the estimates from the stratified random sample. We use function *svydesign* to input the design information and functions *svymean* and *svytotal* to calculate the survey statistics. The following gives the code to find estimates for data *agstrat* along with the output of the statistics calculated. These estimates were given in Examples 3.2 and 3.6 of SDA.

First, we set up the information for the survey design.

```
# create a variable containing population stratum sizes, for use in fpc (optional)
# popsize_recode gives popsize for each stratum
popsize_recode <- c('NC' = 1054, 'NE' = 220, 'S' = 1382, 'W' = 422)
# next statement substitutes 1054 for each 'NC', 220 for 'NE', etc.
agstrat$popsize <- popsize_recode[agstrat$region]
table(agstrat$popsize) #check the new variable
##
```

```
## 220 422 1054 1382
## 21 41 103 135
# input design information for agstrat
dstr <- svydesign(id = ~1, strata = ~region, weights = ~strwt, fpc = ~popsize,
                data = agstrat)

dstr
## Stratified Independent Sampling design
## svydesign(id = ~1, strata = ~region, weights = ~strwt, fpc = ~popsize,
## data = agstrat)
```

The syntax is similar to that for an SRS. The only difference is in the arguments to the *svydesign* function. The arguments used for stratification are as follows:

id As for an SRS, we use `id = ~1` to indicate that there is no clustering.

strata The *strata* argument gives the variable name(s) containing the stratification information (here, the stratification variable is *region*).

weights For stratified sampling, we use the weights associated with the selection probabilities in each stratum. This example has a stratified sample with proportional allocation, where the weights are almost identical for all strata. In samples with disproportionate stratification, however, the weights will vary across strata, and estimates calculated without weights will be biased.

fpc The *fpc* argument specifies the variable that contains information for calculating the finite population correction (fpc) in each stratum. The easiest way to do that is to create a new variable in the data frame that contains the population stratum sizes. Our code defines a variable *popsize_recode* that associates each stratum name with its population size; alternatively, the variable *popsize* could be created with the *merge* or *match* function or by assigning values separately to each region in a *for* loop.

If you omit the *fpc* argument (and still include the *weights* argument), the estimates of means and totals are the same, but standard error estimates are without the finite population correction.

All of the work specifying the design information is done in the *svydesign* function; after you have defined the design there, the *svymean* and *svytotal* functions are used exactly as in Chapter 2 for an SRS.

```
# calculate mean, SE and confidence interval
smean<-svymean(~acres92, dstr)
smean
##          mean      SE
## acres92 295561 16380
confint(smean, level=.95, df=degf(dstr)) # note that df = n-H = 300-4
##          2.5 %    97.5 %
## acres92 263325 327796.5
# calculate total, SE and CI
stotal<-svytotal(~acres92, dstr)
stotal
##          total      SE
## acres92 909736035 50417248
degf(dstr) # Show the degrees of freedom for the design
## [1] 296
# calculate confidence intervals using the degrees of freedom
confint(stotal, level=.95,df= degf(dstr))
##          2.5 %    97.5 %
```

```
## acres92 810514350 1008957721
```

The output is pretty self-explanatory. Note that 296 degrees of freedom ($df = n - H$) are used for the confidence intervals. The df can also be found by applying function *degf* to the design object *dstr*, that is, *degf(dstr)*. If you want to calculate confidence intervals that are based on the normal distribution, simply omit the *df* argument in the *confint* function. If the sample has few observations, however, we need to specify the degrees of freedom and use the *t* distribution to calculate confidence intervals.

Weights and fpc arguments. We supplied both *weights* and *fpc* arguments to the *svydesign* function in this example, but for a stratified random sample with no nonresponse, the *svydesign* function will calculate weights from the *fpc* information and the sample sizes in the data set. The design object *dstrfpc* in the following code results in the same statistics as the design object *dstr* (with the *weights* and *fpc* arguments) that we used earlier. Including the *weights* argument but omitting the *fpc* argument results in standard errors that are calculated without the *fpc*. (Do not omit both *weights* and *fpc*; then the *svydesign* function will assume all weights are equal.)

```
# Alternative design specifications
# Get same result if omit weights argument since weight = popsize/n_h
dstrfpc <- svydesign(id = ~1, strata = ~region, fpc = ~popsize, data = agstrat)
svymean(~acres92, dstrfpc)
##          mean      SE
## acres92 295561 16380
# If you include weights but not fpc, get SE without fpc factor
dstrwt <- svydesign(id = ~1, strata = ~region, weights = ~strwt, data = agstrat)
svymean(~acres92, dstrwt)
##          mean      SE
## acres92 295561 17241
```

Calculating stratum means and variances. Function *svyby* will calculate statistics and their standard errors for subgroups of the data. Here we use it to calculate the stratum means and totals. The first argument of *svyby* is the formula for the variable(s) for which statistics are desired, and the second argument (*by=*) is the variable that defines the groups. Then list the design object and the name of the function that calculates the statistics. Set *keep.var=TRUE* to display the standard errors for the statistics.

```
# calculate mean and se of acres92 by regions
svyby(~acres92, by=~region, dstr, svymean, keep.var = TRUE)
##    region  acres92      se
## NC      NC 300504.16 16107.59
## NE      NE  97629.81 18149.49
## S       S 211315.04 18925.35
## W       W 662295.51 93403.65
# calculate total and se of acres92 by regions
svyby(~acres92, ~region, dstr, svytotal, keep.var = TRUE)
##    region  acres92      se
## NC      NC 316731380 16977399
## NE      NE  21478558  3992889
## S       S 292037391 26154840
## W       W 279488706 39416342
```

If you want to check the calculations by formula, you can also calculate summary statistics directly for each stratum using the *tapply* function and then use the formulas from SDA to calculate the standard errors for each estimated stratum mean or total. The variances

of the stratum means are calculated with the formula $(1 - n_h/N_h)s_h^2/n_h$, where n_h and N_h are the sample and population sizes, and s_h^2 is the sample variance within stratum h .

```
# formula calculations, using tapply
# variables sampsize and popsize were calculated earlier in the chapter
# calculate mean within each region
strmean<-tapply(agstrat$acres92,agstrat$region,mean)
strmean
##          NC          NE          S          W
## 300504.16  97629.81 211315.04 662295.51
# calculate variance within each region
strvar<-tapply(agstrat$acres92,agstrat$region,var)
strvar
##          NC          NE          S          W
## 29618183543  7647472708 53587487856 396185950266
# verify standard errors by direct formula
strse<- sqrt((1-sampsize/popsize)*strvar/sampsize)
# same standard errors as from svyby
strse
##
##          NC          NE          S          W
## 16107.59 18149.49 18925.35 93403.65
```

3.4 Estimating Proportions from a Stratified Random Sample

A proportion is a special case of a mean of a variable taking on values 1 and 0. As defined in Chapter 2, variable *lt200k* takes on value 1 if *acres92* < 200,000 and takes on value 0 if *acres92* ≥ 200,000. The mean of variable *lt200k* estimates the proportion of farms that have fewer than 200,000 acres. The total of variable *lt200k* estimates the number of farms that have fewer than 200,000 acres.

```
# Create variable lt200k
agstrat$lt200k <- rep(0,nrow(agstrat))
agstrat$lt200k[agstrat$acres92 < 200000] <- 1
# Rerun svydesign because the data set now has a new variable
dstr <- svydesign(id = ~1, strata = ~region, fpc = ~popsize,
  weights = ~strwt, data = agstrat)
# calculate proportion, SE and confidence interval
smeanp<-svymean(~lt200k, dstr)
smeanp
##          mean          SE
## lt200k 0.51391 0.0248
confint(smeanp, level=.95,df=degf(dstr))
##          2.5 %          97.5 %
## lt200k 0.4651188 0.5627107
# calculate total, SE and CI
stotalp<-svytotals(~lt200k, dstr)
stotalp
##          total          SE
## lt200k 1581.8 76.318
confint(stotalp, level=.95,df=degf(dstr))
##          2.5 %          97.5 %
## lt200k 1431.636 1732.024
```


You can also calculate proportions and totals of categorical variables by defining them to be factors, either by declaring the variable to be a factor variable in the data set or in the function call of *svymean*. Here we define variable *lt200kf* to be a factor variable in the data set.

```
# Create a factor variable lt200kf
agstrat$lt200kf <- factor(agstrat$lt200k)
# Rerun svydesign because the data set now has a new variable
dstr <- svydesign(id = ~1, strata = ~region, fpc = ~popsize,
                 weights = ~strwt, data = agstrat)
# calculate proportion, SE and confidence interval
smeanp2<-svymean(~lt200kf, dstr)
smeanp2
##               mean      SE
## lt200kf0 0.48609 0.0248
## lt200kf1 0.51391 0.0248
confint(smeanp2, level=.95,df=degf(dstr))
##           2.5 %    97.5 %
## lt200kf0 0.4372893 0.5348812
## lt200kf1 0.4651188 0.5627107
# calculate total, SE and CI
stotalp2<-svyttotal(~lt200kf, dstr)
stotalp2
##           total      SE
## lt200kf0 1496.2 76.318
## lt200kf1 1581.8 76.318
confint(stotalp2, level=.95,df=degf(dstr))
##           2.5 %    97.5 %
## lt200kf0 1345.976 1646.364
## lt200kf1 1431.636 1732.024
```

Note that the *svyttotal* function gives the estimated total for each category of variable *lt200kf*.

The *survey* package will also estimate asymmetric confidence intervals for survey data (Korn and Graubard, 1998), which may have more accurate coverage probabilities for proportions that are near 0 or 1 than the symmetric confidence intervals based on the normal approximation. This is done with the *svyciprop* function, choosing *method="beta"* to obtain a version of the Clopper-Pearson confidence interval (the function will also compute asymmetric confidence intervals using other methods). We illustrate with binary variable *lt200k*. Note that you need to list the formula as *~I(lt200k)* or *~I(lt200k==1)*.

```
# calculate proportion and confidence interval with svyciprop
svyciprop(~I(lt200k==1), dstr, method="beta")
##               2.5% 97.5%
## I(lt200k == 1) 0.514 0.464 0.56
```

3.5 Additional Code for Exercises

Some of the exercises in Chapter 3 ask you to find an ANOVA table. Here's how to do that for *agstrat* using the *lm* function, which performs regression and analysis of variance. The first argument of *lm* is the formula for the regression model, of the form $y \sim x$. We specify *region* to be a factor so that the function will treat it as a categorical variable. (We'll see

the function that conducts regression analyses for survey data in Chapter 4, and it will have a similar structure.)

```
myfit <- lm(acres92~factor(region), data=agstrat)
anova(myfit)
## Analysis of Variance Table
##
## Response: acres92
##           Df      Sum Sq   Mean Sq F value    Pr(>F)
## factor(region)  3 7.2976e+12  2.4325e+12   27.48 1.048e-15 ***
## Residuals      296 2.6202e+13  8.8521e+10
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

3.6 Summary, Tips, and Warnings

Table 3.2 lists the major functions used in this chapter to select or analyze data from a stratified random sample.

Tips and Warnings

- When calculating optimal allocations, make sure that the variables containing the cost and variance information are in the same order as the variable(s) containing the stratum identifiers.
- Sort the population data set by the stratification variable(s) before calling the *strata* function to select a stratified sample.
- When calculating sampling weights for a stratified random sample, check that the sum of the sampling weights for each stratum equals the population size for that stratum.
- When analyzing data from a stratified random sample, first create the design object in the *svydesign* function, using the **strata=** argument. Then call the *svymean* and *svytotal* function with that design object.
- Functions *svymean*, *svytotal*, and *svyby* can be used to calculate statistics for two or more variables simultaneously. For example, *svymean(~acres92 + acres87, dstr)* will display statistics of both variables *acres92* and *acres87*.

TABLE 3.2

Functions used for Chapter 3.

Function	Package	Usage
order	base	Give indices for data sorted according to the specified variable
sample	base	Select a simple random sample with or without replacement
tapply	base	Apply a function to each group of values; groups are defined by the second argument
confint	stats	Calculate confidence intervals; add df for t confidence interval
lm	stats	Fit a linear model to a data set (not using survey methods)
anova	stats	Calculate an analysis of variance table from a model object
boxplot	graphics	Draw boxplot of data (used to display strata in a stratified random sample)
strata	sampling	Select a stratified random sample
getdata	sampling	Extract the sampled units from the population
svydesign	survey	Specify the survey design; add stratum information for stratified random sample
degf	survey	Find degrees of freedom based on design information
svymean	survey	Calculate mean and standard error of mean (if the variable is numeric), or proportion in each category (if variable is categorical)
svytotal	survey	Calculate total and standard error of total
svyby	survey	Calculate survey statistics on subsets of a survey defined by factors
svyciprop	survey	Compute confidence intervals for proportions using various methods (if estimated proportions are close to 0 or 1, sometimes an asymmetric confidence interval is preferred to the symmetric confidence interval produced by <i>svymean</i>)

Ratio and Regression Estimation

Ratio and regression estimation both use auxiliary information to increase the precision of survey estimates. This chapter shows how to incorporate auxiliary information into survey data analyses using R. The code in this chapter is in file `ch04.R` on the book website.

4.1 Ratio Estimation

Examples 4.2 and 4.3 of SDA. The *svyratio* function in the `survey` package (Lumley, 2020) computes ratios from survey data. Let's see how it works for Examples 4.2 and 4.3 of SDA. As the correlation coefficient between variables *acres87* and *acres92* is 0.995806, *acres87* would be an excellent auxiliary variable for ratio estimation. The code and output to estimate the ratio $\bar{y}_U/\bar{x}_U$, where $\bar{y}_U$ is the population mean of *acres92* and $\bar{x}_U$ is the population mean of *acres87*, are given in the following.

```
data(agsrs)
n<-nrow(agsrs) #300
agsrs$sampwt <- rep(3078/n,n)
agdsrs <- svydesign(id = ~1, weights=~sampwt, fpc=rep(3078,300), data = agsrs)
agdsrs
## Independent Sampling design
## svydesign(id = ~1, weights = ~sampwt, fpc = rep(3078, 300), data = agsrs)
# correlatIon of acres87 and acres92
cor(agsrs$acres87,agsrs$acres92)
## [1] 0.995806
# estimate the ratio acres92/acres87
sratio<-svyratio(numerator = ~acres92, denominator = ~acres87,design = agdsrs)
sratio
## Ratio estimator: svyratio.survey.design2(numerator = ~acres92,
##      denominator = ~acres87, design = agdsrs)
## Ratios=
##      acres87
## acres92 0.9865652
## SEs=
##      acres87
## acres92 0.005750473
confint(sratio, df=degf(agdsrs))
##      2.5 %      97.5 %
## acres92/acres87 0.9752487 0.9978818
```

The sample in *agsrs* is an SRS, so we specify the survey design object in *svydesign* exactly as we did in Chapter 2. The only new feature is the *svyratio* function, which calculates the ratio $\hat{B} = \bar{y}/\bar{x}$ and its standard error.

Now that we have estimated the ratio from the data, we can use the *predict* function to obtain the ratio estimates of the population mean and total of y . The population total of x is $t_x = 964,470,625$ and the population mean of x is $\bar{x}_U = t_x/N$. Note that the value of t_x came from the official U.S. Census of Agriculture statistics for 1987 (U.S. Bureau of the Census, 1995). This is greater than the sum of all x values in data set *agpop* because some counties in the population have missing values for *acres87*, as we saw in Section 1.7.

```
# provide the population total of x
xpopttotal <- 964470625
# Ratio estimate of population total
predict(sratio,total=xpopttotal)
## $total
##          acres87
## acres92 951513191
##
## $se
##          acres87
## acres92 5546162
# Ratio estimate of population mean
predict(sratio,total=xpopttotal/3078)
## $total
##          acres87
## acres92 309133.6
##
## $se
##          acres87
## acres92 1801.872
```

Examples 4.2 and 4.3 of SDA also explore the scatterplot of *acres92* versus *acres87*. Because all of the weights are the same ($=3078/300$), we can use the base R function *plot* to display the data in Figure 4.1 (see Chapter 7 for how to draw scatterplots for samples with unequal weights).

We scale the x and y variables so that the plot shows millions of acres instead of acres, and specify the axis labels in the *xlab* and *ylab* arguments. The function *abline* draws the line through the origin with slope $\hat{B}$.

```
par(las=1) # make tick mark labels horizontal (optional)
plot(x=agsrs$acres87/1e6,y=agsrs$acres92/1e6,
     xlab="Millions of Acres Devoted to Farms (1987)",
     ylab = "Millions of Acres Devoted to Farms (1992)",
     main = "Acres Devoted to Farms in 1987 and 1992")
# draw line through origin with slope Bhat
abline(0,coef(sratio))
```

Example 4.5 of SDA. Variables of interest in this example are the number of woody seedlings in pig-protected areas under each of ten sampled oak trees in 1992 (*seed92*) and 1994 (*seed94*) on Santa Cruz Island, California. The code below draws the scatterplot (shown in Figure 4.4 of SDA and not reproduced here) of *seed94* versus *seed92*. It also calculates the correlation of the two variables.

```
#scatterplot of seed92 and seed94
data(santacruz)
plot(santacruz$seed92,santacruz$seed94,
     main="Number of seedlings in 1994 and 1992",
     xlab="Number of seedlings in 1992",ylab="Number of seedlings in 1994")
```




FIGURE 4.1: Scatterplot of acres devoted to farms in 1987 and 1992 (data *agsrs*).

```
cor(santacruz$seed92,santacruz$seed94)
## [1] 0.6106537
```

Now let's calculate the ratio of the number of seedlings in 1994 to the number of seedlings in 1992. Because the number of trees in the population is unknown, we define *sampwt* to be 1 for each observation, and we omit the *fpc* argument in the *svydesign* function (alternatively, one could set the population size to be a large number for the weights and *fpc*).

```
nrow(santacruz) #10
## [1] 10
santacruz$sampwt <- rep(1,nrow(santacruz))
design0405 <- svydesign(ids = ~1, weights = ~sampwt, data = santacruz)
design0405
## Independent Sampling design (with replacement)
## svydesign(ids = ~1, weights = ~sampwt, data = santacruz)
#Ratio estimation using number of seedlings of 1992 as auxiliary variable
```

```

sratio3<-svyratio(~seed94, ~seed92,design = design0405)
sratio3
## Ratio estimator: svyratio.survey.design2(~seed94, ~seed92, design = design0405)
## Ratios=
##      seed92
## seed94 0.2961165
## SEs=
##      seed92
## seed94 0.1152622
confint(sratio3, df=10-1)
##              2.5 %      97.5 %
## seed94/seed92 0.03537532 0.5568577

```

4.2 Regression Estimation

Example 4.7 of SDA. Function *svyglm* calculates regression coefficients and regression estimators from survey data. It is the survey analog of the R function *glm*, which fits generalized linear models.

```

data(deadtrees)
head(deadtrees)
##   photo field
## 1     10    15
## 2     12    14
## 3      7     9
## 4     13    14
## 5     13     8
## 6      6     5
nrow(deadtrees) # 25
## [1] 25
# Fit with survey regression
dtree<- svydesign(id = ~1, weight=rep(4,25), fpc=rep(100,25), data = deadtrees)
myfit1 <- svyglm(field~photo, design=dtree)
summary(myfit1) # displays regression coefficients
##
## Call:
## svyglm(formula = field ~ photo, design = dtree)
##
## Survey design:
## svydesign(id = ~1, weight = rep(4, 25), fpc = rep(100, 25), data = deadtrees)
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.0593      1.3930   3.632   0.0014 **
## photo         0.6133      0.1259   4.870 6.44e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 5.548341)
##
## Number of Fisher Scoring iterations: 2
confint(myfit1,df=23) # df = 25-2

```



```
##           2.5 %    97.5 %
## (Intercept) 2.3291420 7.7894421
## photo      0.3664593 0.8600894
# Regression estimate of population mean field trees
newdata <- data.frame(photo=11.3)
predict(myfit1, newdata)
##      link      SE
## 1 11.989 0.418
confint(predict(myfit1, newdata),df=23)
##      2.5 %    97.5 %
## 1 11.12455 12.85404
# Estimate total field tree, add population size in total= argument
newdata2 <- data.frame(photo=1130)
predict(myfit1, newdata2, total=100)
##      link      SE
## 1 1198.9 41.802
confint(predict(myfit1, newdata2, total=100),df=23)
##      2.5 %    97.5 %
## 1 1112.455 1285.404
```

The regression estimation uses the following functions:

svydesign As before, use *svydesign* to describe the design of the survey, which, in this case, is an SRS with sample size $n = 25$ and population size $N = 100$. Each observation has a weight of $100/25 = 4$.

svyglm The function `svyglm(field~photo, design=dtree)` tells which variables to analyze in the regression statement. The dependent (y) variable is before the $\sim$ sign, and the independent (x) variables follow it. In this example, the dependent variable is *field* and there is one independent variable, *photo*. The *design* argument tells the name of the survey design object (here, *dtree*) to use in calculations.

summary The *summary* function gives the estimates of the regression coefficients, their associated standard errors, and the t statistic and p -value for testing whether each regression parameter equals 0 (the standard errors and tests will be discussed in Chapter 11).

confint As before, the *confint* function requests confidence limits for the regression parameters. You can specify the degrees of freedom with the *df* argument if desired: here the *df* equal the sample size minus 2: $25 - 2 = 23$. If you omit the *df* argument, the function uses the normal distribution to produce confidence intervals.

predict The *predict* function allows you to obtain estimates for predicted values from the estimated regression equation. For regression estimation of the mean, we first define a new data frame with `photo = 11.3` because we want to calculate the predicted value of the regression function at $\bar{x}_U = 11.3$. The statement `predict(myfit1, newdata)` gives $\hat{B}_0 * (1) + \hat{B}_1 * (11.3)$. Regression estimation of the total multiplies each of these by the population size N (here, $N = 100$) by including the argument `total=100` in the *predict* function (called as `predict(myfit1, newdata2, total=100)`), which estimates $\hat{B}_0 * (100) + \hat{B}_1 * (1130)$, where $t_x = 1130$.

Note that the output gives slightly different standard errors and confidence intervals for the regression estimates of the mean and total than SDA because the R functions use a slightly different (although asymptotically equivalent) formula to calculate the standard error. See Section 11.6 of SDA for a discussion of the two variance estimates used.

4.3 Domain Estimation

A domain is a subset of the population for which estimates are desired. Because estimated domain means and totals are ratio estimates, they can be calculated with the *svyratio* function. It is usually easier, however, to compute them using *subset* or *svyby*.

The procedure to calculate estimates for domains is essentially the same as that to calculate estimates for the full sample, but you need to redefine the design for the domain with the *subset* function so that standard errors are calculated correctly. Type:

```
newdesign<-subset(original_design, domain)
```

Example 4.8 of SDA. The following code uses the *subset* function to request design information for each level of the variable *farmcat*, which is defined to equal “large” when *farms92* ≥ 600 and “small” otherwise.

```
agsrsnew<-agsrs #copy agsrs as agsrsnew, since we want to create a new column
# we calculated sampwt in the first code in this chapter
# define new variable farmcat
agsrsnew$farmcat<-rep("large",n)
agsrsnew$farmcat[agsrsnew$farms92 < 600] <- "small"
head(agsrsnew)
##               county state acres92 acres87 acres82 farms92 farms87 farms82
## 1      COFFEE COUNTY   AL  175209  179311  194509    760    842    944
## 2    COLBERT COUNTY   AL  138135  145104  161360    488    563    686
## 3      LAMAR COUNTY   AL   56102   59861   72334    299    362    447
## 4    MARENGO COUNTY   AL  199117  220526  231207    434    471    622
## 5     MARION COUNTY   AL   89228  105586  113618    566    658    748
## 6 TUSCALOOSA COUNTY   AL   96194  120542  134616    436    521    650
##  largef92 largef87 largef82 smallf92 smallf87 smallf82 region sampwt farmcat
## 1         29      28      21        57        47        66      S   10.26   large
## 2         37      41      42        12        44        47      S   10.26   small
## 3          4       4       3         16        20        30      S   10.26   small
## 4         48      66      62         14        11        28      S   10.26   small
## 5          7       9       9         11        23        27      S   10.26   small
## 6         20      17      23         18        32        29      S   10.26   small
dsrsnew <- svydesign(id = ~1, weights=~sampwt, fpc=rep(3078,300), data=agsrsnew)
# domain estimation for large farmcat with subset statement
dsub1<-subset(dsrsnew,farmcat=='large') # design info for domain large farmcat
smean1<-svymean(~acres92,design=dsub1)
smean1
##               mean      SE
## acres92 316566 21553
df1<-sum(agsrsnew$farmcat=='large')-1 #calculate domain df if desired
df1
## [1] 128
confint(smean1, level=.95,df=df1) # CI
##               2.5 %    97.5 %
## acres92 273918.9 359212.4
stotal1<-svytotal(~acres92,design=dsub1)
stotal1
##               total      SE
## acres92 418987302 38938277
confint(stotal1, level=.95,df=df1)
```

```
##           2.5 %    97.5 %
## acres92 341941269 496033335
# domain estimation for small farmcat
dsub2<-subset(dsrsnew,farmcat=='small') # design info for domain small farmcat
smean2<-svymean(~acres92,design=dsub2)
smean2
##           mean      SE
## acres92 283814 28852
df2<-sum(agsrsnew$farmcat=='small')-1 #calculate domain df if desired
confint(smean2, level=.95,df=df2) #CI
##           2.5 %    97.5 %
## acres92 226858.9 340768.5
stotal2<-svytotal(~acres92,design=dsub2)
stotal2
##           total      SE
## acres92 497939808 55919525
confint(stotal2, level=.95,df=df2)
##           2.5 %    97.5 %
## acres92 387553732 608325884
```

You can also calculate statistics for all domains defined by a factor variable at the same time, using the *svyby* function. Here, we estimate the population total and mean for both domains defined by *factor(farmcat)*. The first argument of *svyby* contains the variable(s) to analyze, and the second argument is the factor variable that defines the domains. The last argument gives the name of the function that is to be applied to each group in the *by* argument.

```
bothtot<-svyby(~acres92,by=~factor(farmcat),design=dsrsnew,svytotal)
bothtot
##           factor(farmcat)  acres92      se
## large           large 418987302 38938277
## small           small 497939808 55919525
confint(bothtot,level=.95)
##           2.5 %    97.5 %
## large 342669682 495304922
## small 388339553 607540062
bothmeans<-svyby(~acres92,by=~factor(farmcat),design=dsrsnew,svymean)
bothmeans
##           factor(farmcat)  acres92      se
## large           large 316565.7 21553.21
## small           small 283813.7 28852.24
confint(bothmeans,level=.95)
##           2.5 %    97.5 %
## large 274322.1 358809.2
## small 227264.4 340363.1
```

Note that confidence intervals here are slightly smaller than those given from the calculations with the *subset* function and in Example 4.8 of SDA. Because we did not specify the *df* in the *confint* function, it uses a normal distribution to calculate the intervals; the previous code, using the *subset* function, calculated the confidence intervals using a *t* distribution having $n_d - 1$ *df*, where n_d is the sample size of domain *d*.

Warning. In SRSs, you can calculate domain means and their standard errors by first forming a new, subsetting data set that consists of the observations in the domain and then calculating statistics on the subsetting data set. In complex surveys, however, that

method can result in incorrect standard errors (see Section 11.3 of SDA). To obtain correct statistics for domains, first define the survey design object using the function *svydesign* with the entire data set. Then use the function *subset* or *svyby* with the survey design object to obtain correct inferences for domains.

4.4 Poststratification

Example 4.9 of SDA. The *postStratify* function computes poststratification weights and uses them to estimate population means and totals, along with their standard errors (discussed in Chapter 11 of SDA). Let's poststratify the SRS in *agsrs* by variable *region*.

```
data(agsrs)
dsrs <- svydesign(id = ~1, weights=rep(3078/300,300), fpc=rep(3078,300),
  data = agsrs)
# Create a data frame that gives the population totals for the poststrata
pop.region <- data.frame(region=c("NC","NE","S","W"), Freq=c(1054,220,1382,422))
# create design information with poststratification
dsrsp<-postStratify(dsrs, ~region, pop.region)
summary(dsrsp)
## Independent Sampling design
## postStratify(dsrs, ~region, pop.region)
## Probabilities:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.09242 0.09407 0.09407 0.09771 0.10152 0.10909
## Population size (PSUs): 3078
## Data variables:
## [1] "county" "state" "acres92" "acres87" "acres82" "farms92"
## [7] "farms87" "farms82" "largef92" "largef87" "largef82" "smallf92"
## [13] "smallf87" "smallf82" "region"
1/unique(dsrsp$prob) # See the poststratified weight for each region
## [1] 10.630769 10.820513 9.850467 9.166667
svymean(~acres92, dsrsp)
##           mean      SE
## acres92 299778 17513
svytotal(~acres92, dsrsp)
##           total      SE
## acres92 922717031 53906392
```

The only new feature here is the *postStratify* function:

```
postStratify(design=dsrs, strata=region, population=pop.region)
```

The *postStratify* function tells R to construct poststratification weights, using poststrata in the variable *region*. The third argument is the name of the data frame (here, *pop.region*) that gives the population totals for the poststrata. The poststratified estimates of the population mean and total of *acres92*, when calculated with poststratified design object *dsrsp*, are reported together with standard errors when the *svymean* or *svytotal* function is called.

Note that the standard errors reported by the **survey** package differ slightly from those in Example 4.9 of SDA because a slightly different (although asymptotically equivalent) estimator for the variance is used (see Section 11.6 of SDA).

4.5 Ratio Estimation with Stratified Sampling

The *svyratio* function will compute either separate or combined ratio estimates. The default is combined ratio estimation, which calculates the ratio $\hat{y}/\hat{x}$, where $\hat{y}$ is the estimate of the mean of y using the stratified design and $\hat{x}$ is the estimated mean of x . All we need to do is to include the stratification information in the design structure formed by the *svydesign* function.

Combined ratio estimator. The following shows how to compute the ratio of *acres92* to *acres87* and the ratio estimator of the total for the stratified sample in *agstrat*, using the combined ratio estimator.

```
data(agstrat)
popsize_recode <- c('NC' = 1054, 'NE' = 220, 'S' = 1382, 'W' = 422)
agstrat$popsize <- popsize_recode[agstrat$region]
# input design information for agstrat
dstr <- svydesign(id = ~1, strata = ~region, fpc = ~popsize, weight = ~strwt,
                 data = agstrat)
# now compute the combined estimator of the ratio
combined<-svyratio(~ acres92,~acres87,design = dstr)
combined
## Ratio estimator: svyratio.survey.design2(~acres92, ~acres87, design = dstr)
## Ratios=
##          acres87
## acres92 0.9899971
## SEs=
##          acres87
## acres92 0.006187757
# we can get the combined ratio estimator of the population total
# with the predict function
predict(combined,total=964470625)
## $total
##          acres87
## acres92 954823130
##
## $se
##          acres87
## acres92 5967910
```

Separate ratio estimator. You can calculate ratios separately for each stratum by including *separate=TRUE* in the *svyratio* function.

```
separate<-svyratio(~acres92,~acres87,design = dstr,separate=TRUE)
separate
## Stratified ratio estimate: svyratio.survey.design2(~acres92, ~acres87,
##          design = dstr, separate = TRUE)
## Ratio estimator: Stratum == "NC"
## Ratios=
##          acres87
## acres92 0.9750666
## SEs=
##          acres87
## acres92 0.005483458
## Ratio estimator: Stratum == "NE"
```

```
## Ratios=
##          acres87
## acres92 0.8956073
## SEs=
##          acres87
## acres92 0.008853011
## Ratio estimator: Stratum == "S"
## Ratios=
##          acres87
## acres92 0.9935483
## SEs=
##          acres87
## acres92 0.01418835
## Ratio estimator: Stratum == "W"
## Ratios=
##          acres87
## acres92 1.011974
## SEs=
##          acres87
## acres92 0.01169809
# Define the stratum totals for acres87 as a list:
stratum.xtotals <- list(NC=350474227,NE=22033421,S=280631939,W=311331038)
predict(separate,stratum.xtotals)
## $total
##          acres87
## acres92 955349448
##
## $se
##          acres87
## acres92 5731438
```

4.6 Model-Based Ratio and Regression Estimation

This section is optional and need only be read if covering Section 4.6 of SDA.

Example 4.11 of SDA. A model-based analysis of data from an SRS uses the same techniques taught in an introductory statistics class. Since the model-based analysis does not make use of the sampling weights, the *lm* or *glm* functions, which fit linear and generalized linear models for non-survey data, are used to fit the regression models and obtain the residuals. Here we use the *lm* function.

The format for fitting a regression model with *lm* is similar to *svyglm*, but with one important difference: the weights mean different things in the two functions. In the *svyglm* function, **weights=** tells how many population units are represented by each sample unit. In the *lm* function, the weight variable contains relative weights for a weighted least squares fit.

The model used is $Y_i = \beta x_i + \varepsilon_i$, with $V(\varepsilon_i) = \sigma^2 x_i$. The model has variance proportional to x_i , so obtaining the best linear unbiased estimates under this model would use a weight value proportional to the reciprocal of the variances. This is specified by defining the variable $recacr87 = 1/acres87$ when $acres87 > 0$ and $recacr87 = \text{NA}$ when $acres87 = 0$ (to avoid division by zero).

```

data(agsrs)
# define weights to use for weighted least squares analysis
agsrs$recacr87<-agsrs$acres87
agsrs$recacr87[agsrs$acres87!=0] <- 1/agsrs$acres87[agsrs$acres87!=0]
agsrs$recacr87[agsrs$acres87==0] <- NA
# fit weighted least squares model without intercept
fit<-lm(acres92~acres87-1,weights=recacr87,data=agsrs)
summary(fit)
##
## Call:
## lm(formula = acres92 ~ acres87 - 1, data = agsrs, weights = recacr87)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -369.9  -22.2   -5.8   10.8  311.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## acres87  0.986565    0.004844   203.7   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 46.1 on 298 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.9929, Adjusted R-squared:  0.9928
## F-statistic: 4.149e+04 on 1 and 298 DF,  p-value: < 2.2e-16
anova(fit)
## Analysis of Variance Table
##
## Response: acres92
##           Df    Sum Sq Mean Sq F value    Pr(>F)
## acres87     1 88168461 88168461   41487 < 2.2e-16 ***
## Residuals 298   633307     2125
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
# find predicted value at population total for x
newdata3 <- data.frame(acres87=964470625)
predict(fit, newdata3, se.fit=TRUE)
## $fit
##           1
## 951513191
##
## $se.fit
## [1] 4671509
##
## $df
## [1] 298
##
## $residual.scale
## [1] 46.0998

```

The *weights* argument in *lm* specifies that a weighted least squares analysis is performed with weights *recacr87*, minimizing the weighted sum of squares $\sum_{i \in S} (y_i - \beta x_i)^2 / x_i$. The “-1” in *lm(acres92~acres87-1)* tells that the model is to be fit without an intercept. The *summary* function displays the regression coefficient $\hat{\beta} = 0.986565$ and the *anova* function

displays the ANOVA table. The model is fit to the 299 observations that have $acres87 > 0$.

The *predict* function requests the predicted value from the regression model when $acres87$ takes on the value $t_x = 964,470,625$, giving $\hat{t}_{yM} = \hat{\beta}t_x = 951,513,191$. The standard error, without the *fpc*, is $\sigma t_x / \sqrt{\sum_{i \in S} x_i} = 4,671,509$.

Note that the sum of squares for error in the ANOVA table, 633,307, is the sum of squares of the weighted residuals, so the mean squared error in the ANOVA table gives $\hat{\sigma}^2 = 2125.19$.

The residuals produced by *lm* are $e_i = y_i - \hat{y}_i$. For a ratio model, the weighted residuals $e_{iw} = e_i / \sqrt{x_i}$ should be plotted instead of e_i , because if the model variance structure is correct, the e_{iw} should all have approximately equal variances and a plot of e_{iw} versus the predicted values or x_i will show no patterns.

```
# plot weighted residual versus predicted values
wresid<-fit$residuals*sqrt(fit$weights)
par(las=1)
plot(fit$fitted.values, wresid,
     main="Plot of weighted residuals versus predicted values",
     xlab="Predicted value from regression model",
     ylab="Weighted residuals")
```




FIGURE 4.2: Plot of weighted residuals versus fitted values.

Figure 4.2 shows a couple of potential outliers, but no other indications that the model is inappropriate.

Example 4.12 of SDA. The *lm* function is also used to fit a regression model and to obtain residuals for the dead tree data from Example 4.7 of SDA.


```

data(deadtrees)
# Fit with lm
fit2 <- lm(field~photo, data=deadtrees)
summary(fit2)
##
## Call:
## lm(formula = field ~ photo, data = deadtrees)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.0319 -1.8053  0.1947  1.4212  3.8080
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.0593     1.7635   2.869 0.008676 **
## photo         0.6133     0.1601   3.832 0.000854 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.406 on 23 degrees of freedom
## Multiple R-squared:  0.3896, Adjusted R-squared:  0.3631
## F-statistic: 14.68 on 1 and 23 DF,  p-value: 0.0008538
# Estimate mean field trees
newdata <- data.frame(photo=11.3)
predict(fit2, newdata, se.fit=TRUE)
## $fit
##      1
## 11.98929
##
## $se.fit
## [1] 0.4941007
##
## $df
## [1] 23
##
## $residual.scale
## [1] 2.406153

```

Because the regression model is fit under the assumption that $V(\varepsilon_i) = \sigma^2$ for all observations, no *weights* argument is used in *lm*. The *predict* function gives the regression estimate of the population mean $\hat{\beta}_0 + \hat{\beta}_1 \bar{x}_{\mathcal{U}} = 11.9893$, and its standard error (without fpc) of 0.494. These are the values calculated in Example 4.12 of SDA. Typing `summary(fit2)` gives the regression coefficients, their standard errors, and other information about the fit.

We have applied both *lm* and *svyglm* to analyze the tree data from Example 4.7. Table 4.1 compares the estimates and standard errors from the two functions. All the point estimates are the same, but the standard errors from *svyglm* differ from those calculated by *lm*; Sections 4.6 and 11.4 of SDA discuss why that occurs.


TABLE 4.1

Comparison of the estimates and standard errors for model `field~photo` from the *lm* and *svyglm* functions.

	Intercept		Slope		Predicted Value, $x = 11.3$	
	lm	svyglm	lm	svyglm	lm	svyglm
Estimate	5.0593	5.0593	0.6133	0.6133	11.989	11.989
SE	1.7635	1.3930	0.1601	0.1259	0.494	0.418

Requesting `plot(fit2)` produces a collection of residual and diagnostic plots from the *lm* model object. Figure 4.3 displays the plot of the residuals versus *photo* (the x variable), which shows no pattern.

```
# plot residuals versus predicted values
plot(deadtrees$photo, fit2$residuals,
     main="Plot of residuals versus photo values",
     xlab="Photo values (x variable)",
     ylab="Residuals")
```

**FIGURE 4.3:** Plot of residuals versus x variable.

4.7 Summary, Tips, and Warnings

Table 4.2 lists the major functions used in this chapter to compute ratio and regression estimates, calculate statistics for domains, and create poststratification weights.

To calculate a ratio or ratio estimator for an SRS or stratified sample, use the *svyratio* function from the **survey** package. The function *svyglm* fits regression models to survey data.

Tips and Warnings

- Draw a scatterplot of your data when fitting a ratio or regression model, so you can see whether ratio or regression estimation is likely to improve efficiency.
- For domain estimation, first define a design object for the entire sample with the *svydesign* function. Then use the *subset* function to define a domain of interest, or calculate statistics for all domains with the *svyby* function. It may be tempting to calculate statistics for a subset of the sample by creating a data set containing only that subset, but doing that can result in incorrect standard errors for domain statistics.
- Poststratification can be done using the *postStratify* function.

TABLE 4.2

Functions used for Chapter 4.

Function	Package	Usage
subset	base	Work with a subset of a vector, matrix, or data frame
confint	stats	Calculate confidence intervals
cor	stats	Calculate the correlation of vectors (not using survey methods)
lm	stats	Fit a linear model to a data set (not using survey methods)
anova	stats	Compute an analysis of variance table from a model object
predict	stats	Obtain predicted values from a model object
plot	graphics	Draw a scatterplot of data
abline	graphics	Add a straight line to a plot
svydesign	survey	Specify the survey design
svymean	survey	Calculate mean and standard error of mean
svyratio	survey	Calculate a ratio or ratio estimate from survey data
svyglm	survey	Fit a regression model to survey data. The coefficients may then be used to calculate regression estimates
svytotal	survey	Calculate total and standard error of total
svyby	survey	Calculate statistics for subsets of a survey defined by a factor variable
postStratify	survey	Adjust the sampling weights using poststratification

Cluster Sampling with Equal Probabilities

This chapter shows how to use R to compute estimates from one- and two-stage cluster samples when an SRS is selected at each stage. Chapter 6 will tell how to select a one-stage or two-stage cluster sample with equal or unequal probabilities—the syntax is similar for both, and deferring the sample selection examples to Chapter 6 allows us to look at the general case. The code in this chapter is in file `ch05.R` on the book website.

5.1 Estimates from One-Stage Cluster Samples

Example 5.2 of SDA. The following code and output produces estimates of the population mean and total for the GPA data, using functions from the `survey` package (Lumley, 2020). Variable `suite` identifies the clusters in the data, and variable `wt` is the sampling weight for the persons in the sample, defined as $100/5 = 20$ for every person.

```
data(gpa)
# define one-stage cluster design
# note that id is suite instead of individual student as we take an SRS of suites
dgpa<-svydesign(id=~suite,weights=~wt,fpc=~rep(100,20),data=gpa)
dgpa
## 1 - level Cluster Sampling design
## With (5) clusters.
## svydesign(id = ~suite, weights = ~wt, fpc = ~rep(100, 20), data = gpa)
# estimate mean and se
gpamean<-svymean(~gpa,dgpa)
gpamean
##      mean      SE
## gpa 2.826 0.1637
degf(dgpa)
## [1] 4
# n=5, t-approximation is suggested for CI
confint(gpamean,level=.95,df=4) # use t-approximation
##      2.5 %    97.5 %
## gpa 2.371593 3.280407
# confint(gpamean,level=.95) # uses normal approximation, if desired (for large n)
# estimate total and se (if desired)
gpatotal<-svytotal(~gpa,dgpa)
gpatotal
##      total      SE
## gpa 1130.4 65.466
confint(gpatotal,level=.95,df=4)
##      2.5 %    97.5 %
## gpa 948.6374 1312.163
```

The following features of the code and output deal with the cluster sampling:

- The `id=~suite` argument in *svydesign* tells all functions that use the design object that variable *suite* contains the primary sampling unit (psu) identifiers. If you omit the psu variable and instead use `id=~`, the data will be (incorrectly) analyzed as an SRS.
- The `weights=~wt` argument says that variable *wt* contains the sampling weights for the observation units (here, the students). If `weights=` is omitted but `fpc=` is supplied, selection probabilities are calculated from the population sizes, assuming an SRS of psus. That results in the same weights for this survey, but we recommend including the `weights=` argument as routine practice because most surveys have some type of adjustment that causes the final weights to differ from the sampling weights. (There is one exception, and that is for the without-replacement variance calculations discussed in Section 5.2.)
- The `fpc=~rep(100,20)` argument indicates that the total number of psus in the population is 100. When the *fpc* argument is included, functions use the finite population correction (fpc) when calculating variances. Omit `fpc=` if you do not want an fpc (but then make sure you include the `weights=` argument).
- Typing `dgpa` shows that this is a “1 - level Cluster Sampling design With (5) clusters”. The *svydesign* function recognizes this as a one-stage design because one clustering variable is included in the *id* argument. This is the only indication in the output that the clustering was used in the analysis. Otherwise, the form of the statistics output is the same as for simple random or stratified sampling. Always check that the number of clusters in the design object equals the number of psus in your sample.
- The degrees of freedom (df), from function *degf*, equals the number of psus minus 1.
- The *svymean* and *svytotal* functions produce standard errors (SEs) for estimated population means and totals that account for the clustering in the design.

You can verify the calculations using the formulas given in Section 5.2 of SDA if desired.

```
# you can also calculate SEs by direct formula
suitesum<-tapply(gpa$gpa,gpa$suite,sum) #sum gpa for each suite
# variability comes from among the suites
st2<-var(suitesum)
st2
## [1] 2.25568
# SE of t-hat, formula (5.3) of SDA
vthat <-100^2*(1-5/100)*st2/5
sqrt(vthat)
## [1] 65.46596
# SE of ybar, formula (5.6) of SDA
sqrt(vthat)/(4*100)
## [1] 0.1636649
```

The variability *st2* is coming from the suite totals, and the fpc $(1 - 5/100)$ is applied to calculate the variance of $\hat{t}$. The SE of $\hat{t}$ is 65.46596, which is the same as calculated by *svytotal*.

The procedure for calculating estimates from one-stage cluster samples is exactly the same when the psus have unequal sizes.

Example 5.6 of SDA. Data set *algebra* has 12 classes (clusters) with unequal sizes selected from 187 classes. The syntax for analyzing the data with unequal-sized clusters is exactly

the same as for Example 5.2. Again, note that we use the number of clusters minus 1 (= 11) as the df.

```
data(algebra)
algebra$sampwt<-rep(187/12,299)
# define one-stage cluster design
dalg<-svydesign(id=~class,weights=~sampwt,fpc=~rep(187,299), data=algebra)
dalg
## 1 - level Cluster Sampling design
## With (12) clusters.
## svydesign(id = ~class, weights = ~sampwt, fpc = ~rep(187, 299),
##      data = algebra)
# estimate mean and se
svymean(~score,dalg)
##          mean          SE
## score 62.569 1.4916
# n=12, t-distribution is suggested for CI
degf(dalg)
## [1] 11
confint(svymean(~score,dalg),level=.95,df=11) #use t-approximation
##          2.5 %   97.5 %
## score 59.28562 65.8515
# estimate total and se if desired
svytotal(~score,dalg)
##          total          SE
## score 291533 19893
confint(svytotal(~score,dalg),level=.95,df=11)
##          2.5 %   97.5 %
## score 247749.4 335316.6
```

5.2 Estimates from Multi-Stage Cluster Samples

Calculating estimates from a multi-stage cluster sample is similar. The clustering structure is specified for the design object in the *svydesign* function, and then all functions called with that design object account for the clustering in the variance calculations.

There are several ways to estimate variances using the **survey** package. Let's start with Example 5.8 of SDA, where we calculate the with-replacement variance, and then discuss the issues involved for calculating variances for without-replacement samples using the schools data in Example 5.7 of SDA.

Example 5.8 of SDA. The *coots* data come from Arnold's (1991) work on egg size and volume of American coot eggs in Minnedosa, Manitoba, with a sample of 184 clutches (nests of eggs). Variable *csize* gives the number of eggs in the clutches. Two eggs (secondary sampling unit, ssu) are randomly selected from each clutch (psu). Since we do not have information on the total number of psus N , we use the relative weights *relwt* defined by $csize/2$ to calculate the mean volume of eggs and its standard error.

```
data(coots)
# Want to estimate the mean egg volume
nrow(coots) #368
## [1] 368
```

```

coots$ssu<-rep(1:2,184) # index of ssu
coots$relwt<-coots$csz/2
head(coots)
##   clutch csz length breadth   volume tmt ssu relwt
## 1      1   13  44.30   31.10 3.7957569   1   1   6.5
## 2      1   13  45.90   32.70 3.9328497   1   2   6.5
## 3      2   13  49.20   34.40 4.2156036   1   1   6.5
## 4      2   13  48.70   32.70 4.1727621   1   2   6.5
## 5      3    6  51.05   34.25 0.9317646   0   1   3.0
## 6      3    6  49.35   34.40 0.9007362   0   2   3.0
dcoots<-svydesign(id=~clutch+ssu,weights=~relwt,data=coots)
dcoots
## 2 - level Cluster Sampling design (with replacement)
## With (184, 368) clusters.
## svydesign(id = ~clutch + ssu, weights = ~relwt, data = coots)
svymean(~volume,dcoots) #ratio estimator
##           mean      SE
## volume 2.4908 0.061
confint(svymean(~volume,dcoots),level=.95,df=183)
##           2.5 %    97.5 %
## volume 2.370423 2.611134
# now only include psu information, results are the same
dcoots2<-svydesign(id=~clutch,weights=~relwt,data=coots)
dcoots2
## 1 - level Cluster Sampling design (with replacement)
## With (184) clusters.
## svydesign(id = ~clutch, weights = ~relwt, data = coots)
svymean(~volume,dcoots2)
##           mean      SE
## volume 2.4908 0.061

```

In *svydesign*, the two stages of the cluster sampling are given as `id=~clutch+ssu`. The formula lists the psus first and then the ssus. When the with-replacement variance is calculated, however, as is done here, you need only specify the psus—the point and variance estimates are the same whether you specify just the psus or you specify all stages of sampling. The *weight* argument must be included for this design because the weights are unequal.

Note that the confidence interval uses a *t* critical value with 183 df (number of psus minus 1). Also note that *svydesign* does not contain the *fpc* argument. This is because the total number of clutches in the population, N , is unknown. As a result, the *svymean* does not use an *fpc* when calculating estimates. In general, we recommend omitting the *fpc* argument for multi-stage cluster sampling even when N is known, and the remainder of this section discusses this issue.

Variance estimation for without-replacement two-stage cluster samples. Here are two options for estimating the variance of estimated means and totals in without-replacement two-stage sampling, where an SRS is selected at each stage.

Option 1. Calculate the with-replacement variance (recommended). As shown in Sections 5.3 and 6.6 of SDA, the estimated variability among estimated psu totals, s_t^2 , also includes variability from the subsequent stages of sampling. If you estimate the with-replacement variance (at the psu level), the variance estimator incorporates *all* the variability from subsequent stages of sampling. The expected value of the with-replacement variance estimator is larger than the true variance of the without-replacement sample, but the difference is small if the sampling fraction at the psu level, n/N , is small.

Chapter 6 of SDA outlines additional benefits of ignoring the fpcs when the psus are selected with unequal probabilities.

To estimate the with-replacement variance for a multi-stage cluster sample, simply call the *svydesign* function as:

```
svydesign(id=~psuid,weights=~weightvariable,data=dataset)
```

where the *id* formula consists only of the variable giving the psu membership. Do not include the *fpc* argument when calling *svydesign*. Then *svytotal*, *svymean*, and other functions will calculate the with-replacement variance.

Option 2. Calculate the without-replacement variance. When an SRS or stratified random sample is taken at all stages of sampling, you can specify all stages of sampling in the *svydesign* function and calculate without-replacement variances. For a two-stage sample, call the *svydesign* function as:

```
svydesign(id=~psuid+ssuid,fpc=~psufpc+ssufpc,data=dataset).
```

The *id* formula gives the variable identifying the psu membership followed by the variable identifying the ssu membership. The *fpc* formula has *psufpc*, the variable giving the population number of psus ($= \text{rep}(N, \text{nrow}(\text{dataset}))$), followed by *ssufpc*, the variable giving the values of the psu size for each observation ($= M_i$ for ssus in psu *i*).

No *weights* argument is included. When the *weights* argument is omitted, it is assumed that an SRS is taken at both stages, and the inclusion probabilities are calculated from the population sizes given in *fpc* and the sample sizes in the data set. Thus, the weights are assumed to be $(NM_i)/(nm_i)$, where the values of *n* and *m_i* are counted from the data set and the values of *N* and *M_i* are given in the *fpc* arguments.

If there are more than two stages of sampling, and an exactly unbiased estimate of the variance is desired, you need to include terms for all stages of sampling in the *id* and *fpc* arguments of the *svydesign* function. If a survey has three stages, the without-replacement variance estimate requires knowledge of the psu membership, population size, and sample size; the ssu membership, population size, and sample size; and the tertiary sampling unit membership, population size, and sample size.

It can be complicated to keep track of all this information. In addition, calculations are done under the assumption that the final weights are the same as the sampling weights (computed as the inverse of the inclusion probabilities)—that is, there are no nonresponse adjustments or other modifications of the sampling weights.

Example 5.7 of SDA: With-replacement variance. Let's look at the with- and without-replacement variance calculations for the *schools* data. The following code calculates the with-replacement variance. Note that only the psu-level clustering is specified in the *id* argument and that the vector of student-level weights is provided. We can also estimate the proportion and the total number of students having *mathlevel*=2 by treating *mathlevel* as a factor variable.

```
data(schools)
head(schools)
##   schoolid gender  math  reading mathlevel readlevel  Mi finalwt
## 1         9     F   42    42         2         2 163  61.125
## 2         9     F   29    30         1         1 163  61.125
## 3         9     M   31    25         1         1 163  61.125
## 4         9     F   22    33         1         2 163  61.125
## 5         9     M   35    36         1         2 163  61.125
```

```
## 6          9          F    30          17          1          1 163  61.125
# calculate with-replacement variance; no fpc argument
# include psu variable in id; include weights
dschools<-svydesign(id=~schoolid,weights=~finalwt,data=schools)
# dschools tells you this is treated as a with-replacement sample
dschools
## 1 - level Cluster Sampling design (with replacement)
## With (10) clusters.
## svydesign(id = ~schoolid, weights = ~finalwt, data = schools)
mathmean<-svymean(~math,dschools)
mathmean
##          mean          SE
## math 33.123  1.7599
degf(dschools)
## [1] 9
# use t distribution for confidence intervals because there are only 10 psus
confint(mathmean,df=degf(dschools))
##          2.5 %   97.5 %
## math 29.14179 37.1041
# estimate proportion and total number of students with mathlevel=2
svymean(~factor(mathlevel),dschools)
##          mean          SE
## factor(mathlevel)1 0.71231 0.0542
## factor(mathlevel)2 0.28769 0.0542
svytotal(~factor(mathlevel),dschools)
##          total          SE
## factor(mathlevel)1 12303.4 2244.14
## factor(mathlevel)2  4969.1  676.26
```

Example 5.7 of SDA: Without-replacement variance. The *svymean* function will calculate without-replacement variances when simple or stratified random sampling is used at each stage. (As of this writing, it does not do so for all designs and thus will not compute the without-replacement variance for most of the unequal-probability samples that are discussed in Chapter 6.) To use it with the *schools* data, put both stages of clustering in the *id* argument and put both the psu and the ssu population sizes in the *fpc* argument. Do not include the *weights* argument.

```
# create a variable giving each student an id number
schools$studentid<-1:(nrow(schools))
# calculate without-replacement variance
# specify both stages of the sample in the id argument
# give both sets of population sizes in the fpc argument
# do not include the weight argument
dschoolwor<-svydesign(id=~schoolid+studentid,fpc=~rep(75,nrow(schools))+Mi,
                    data=schools)
dschoolwor
## 2 - level Cluster Sampling design
## With (10, 200) clusters.
## svydesign(id = ~schoolid + studentid, fpc = ~rep(75, nrow(schools)) +
##          Mi, data = schools)
mathmeanwor<-svymean(~math,dschoolwor)
mathmeanwor
##          mean          SE
## math 33.123  1.6605
confint(mathmeanwor,df=degf(dschoolwor))
##          2.5 %   97.5 %
```

```
## math 29.36667 36.87923
# estimate proportion and total number of students with mathlevel=2
svymean(~factor(mathlevel),dschoolwor)
##               mean      SE
## factor(mathlevel)1 0.71231 0.0516
## factor(mathlevel)2 0.28769 0.0516
svytotal(~factor(mathlevel),dschoolwor)
##               total      SE
## factor(mathlevel)1 12303.4 2097.83
## factor(mathlevel)2  4969.1  657.69
```

In the *schools* data, variable M_i gives the population number of students in each school. This information must be available in the data set to be able to calculate the without-replacement variance. The design object *dschoolwor* repeats that this is a “2-level Cluster Sampling Design” with 10 psus and 200 ssus.

Even with the relatively large sampling fractions in this example, the with- and without-replacement standard errors are similar. For variable *math*, the with-replacement standard error is 1.76, and the without-replacement standard error is 1.66.

In general, we recommend calculating the with-replacement variance (omitting the *fpc* argument) for multi-stage cluster sampling. It produces a variance estimate whose expectation is slightly larger than the true variance, but if n/N is small, the difference is negligible. If forced to choose between a standard error that is slightly too large and one that is too small, we usually prefer the former because a too-small standard error leads to claiming that estimates are more precise than they really are.

The most important thing to keep in mind for computing standard errors for cluster samples is that ssus in the same psu are usually more homogeneous than randomly selected ssus from the population. Thus, the essential feature for calculating standard errors is to capture that homogeneity by including the *id=psuid* argument in *svydesign*. The issue of “to fpc or not to fpc” is minor compared with the effect of clustering.

5.3 Model-Based Design and Analysis for Cluster Samples

We often use models when designing a cluster sample, as shown in Section 5.4 of SDA. Data from a previous survey or pilot sample may be used to estimate the optimal subsampling or psu size. This often involves estimating the value of R^2 or R_a^2 , which can be obtained from an Analysis of Variance (ANOVA) table.

Example 5.12 of SDA. The following shows how to derive an ANOVA table for the schools data.

```
# run lm with schoolid as a factor
fit5.12<-lm(math~factor(schoolid), data=schools)
# print ANOVA table
anova(fit5.12)
## Analysis of Variance Table
##
## Response: math
##               Df  Sum Sq Mean Sq F value    Pr(>F)
## factor(schoolid)   9  7018.5   779.83   7.5834 1.785e-09 ***
```

```
## Residuals          190 19538.4  102.83
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
# extract the value of R-squared and adjusted R-squared
summary(fit5.12)$r.squared
## [1] 0.264281
summary(fit5.12)$adj.r.squared
## [1] 0.2294312
```

Example 5.14 of SDA. This example employs a random effects model, in which the school means are assumed to be normally distributed random variables with mean μ . In packages `nlme` (Pinheiro et al., 2021) and `lme4` (Bates et al., 2015, 2020), the *lme* (short for linear mixed effects) and *lmer* functions, respectively, calculate estimates from random effects models.

We use function *lme* for this example. For the one-way random effects model, the only fixed effect is the mean, so the *fixed* formula is `fixed=math~1`. Random effects are specified in the *random* argument so `factor(schoolid)`, the factor variable describing the psu membership, is placed behind the vertical bar in the *random* argument.

```
library(nlme)
fit5.14 <- lme(fixed=math~1,random=~1|factor(schoolid),data=schools)
summary(fit5.14)
## Linear mixed-effects model fit by REML
## Data: schools
##      AIC      BIC    logLik
## 1516.259 1526.139 -755.1295
##
## Random effects:
## Formula: ~1 | factor(schoolid)
##      (Intercept) Residual
## StdDev:      5.818064 10.14069
##
## Fixed effects:  math ~ 1
##              Value Std.Error DF  t-value p-value
## (Intercept) 34.66  1.974628 190 17.55267      0
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -2.26713655 -0.74262324 -0.09451607  0.79142521  2.18576500
##
## Number of Observations: 200
## Number of Groups: 10
# extract the variance components
VarCorr(fit5.14)
## factor(schoolid) = pdLogChol(1)
##      Variance StdDev
## (Intercept) 33.84987  5.818064
## Residual    102.83368 10.140694
```

A model-based analysis predicts the values of observations that are not observed in the data. For this data set, the unobserved values are the students who are not measured in the sampled schools, as well as the unsampled schools in the population. The estimated mean 34.66 from the output under "Fixed effects: `math ~ 1`" does not account for the population sizes of the schools, and gives a different estimate than in Example 5.7 of SDA.

The residuals and predicted values from the model can be obtained by requesting `resid(fit5.14)` and `predict(fit5.14)`. You can also type `plot(fit5.14)` to obtain a plot of standardized residuals versus fitted values.

5.4 Additional Code for Exercises

Exercise 5.40 of SDA. The exercise uses the function *intervals_ex40*, available from the book website and R package *SDAResources* (Lu and Lohr, 2021). To run the function, load the *SDAResources* package and type

```
intervals_ex40(groupcorr=0, numintervals=100, groupsize=5, sampgroups=10,
               popgroups=5000, mu=0, sigma=1)
```

using the desired values for the arguments. The function call given above uses the default values of the arguments and will give the same results as running `intervals_ex40()`.

The arguments of function *intervals_ex40* are given in Table 5.1.

TABLE 5.1

Arguments of the *intervals_ex40* function.

Argument	Description
groupcorr	Desired intraclass correlation coefficient, must be between 0 and 1 (default is 0).
numintervals	Number of confidence intervals to be generated (default is 100).
groupsize	Number of observations, M , in each population cluster (default is 5).
sampgroups	Number of clusters to be sampled (default is 10).
popgroups	Number of clusters in population (default is 5000). This should be set to be at least 200 times as large as the value of <i>sampgroups</i> so that the fpc is negligible.
mu	Population mean (default is 0).
sigma	Population standard deviation (default is 1).

For the exercise, you are asked to generate 100 intervals of 50 observations each, taken in 10 clusters of size 5. This uses the default values of all arguments except for ICC. When running the function, you only need to specify the arguments that differ from the default values, so that you can generate 100 intervals with $ICC = 0.3$ by running the function with argument `groupcorr=0.3`.

```
set.seed(9231)
# generate intervals for cluster sample with groupcorr = 0.3
intervals_ex40(groupcorr = 0.3) # leave other parameters unchanged
##   Number_of_intervals   SRS_cover_prob   Cluster_cover_prob
##           100.0000000           0.8000000           0.9500000
##   SRS_mean_CI_width Cluster_mean_CI_width
##           0.5556272           0.9111856
##   Replicate mu sample_mean   srs_lci   srs_uci in_srs_ci SRS_CI_width
## [1,]       1  0 -0.07311322 -0.35806539 0.21183894       1  0.5699043
## [2,]       2  0  0.25038517 -0.05127437 0.55204471       1  0.6033191
## [3,]       3  0  0.03499401 -0.27232646 0.34231448       1  0.6146409
## [4,]       4  0 -0.18948478 -0.45038074 0.07141119       1  0.5217919
```

```
## [5,]          5  0  0.14058713 -0.08585203 0.36702629          1  0.4528783
##          clus_lci  clus_uci in_clu_ci clus_CI_width
## [1,] -0.6526766  0.5064502          1  1.1591268
## [2,] -0.2716028  0.7723731          1  1.0439759
## [3,] -0.4941334  0.5641214          1  1.0582548
## [4,] -0.5814274  0.2024579          1  0.7838853
## [5,] -0.2344528  0.5156271          1  0.7500799
```




FIGURE 5.1: Interval estimates created assuming SRS and using clustering formulas.

We initialized the random number seed so that we could reproduce the intervals later, but if you are repeating the exercise, you may want to let the computer generate your starting seed (otherwise, you may get the same set of samples each time).

- Function `intervals_ex40` calculates two sets of interval estimates: a set that uses SRS formulas and hence has coverage probability (proportion of intervals that include the true population mean 0) `SRS_cover_prob` of 0.80 that is less than 0.95, and a second set that calculates the correct confidence intervals using the formulas for one-stage cluster sampling with coverage probability `Cluster_cover_prob` of 0.95.
- It also prints the average width of the interval estimates for the two methods: `SRS_mean_CI_width` of 0.5556272, which is less than `Cluster_mean_CI_width` of 0.9111856.
- The first five replicates and their summary statistics are printed, where `srs_lci` is the lower limit, and `srs_uci` is the upper limit from the SRS estimate. Similarly, `clus_lci` and `clus_uci` are the lower and upper confidence limits from the estimate calculated using the clustering. If desired, the function can be modified so that information from all replicates is stored in a data set.

- Three graphs are produced, similar to those in Figure 5.1. The first graph shows a scatterplot of the last simulated sample, the second graph shows the interval estimates produced for each sample if analyzed as an SRS, and the third shows the interval estimates produced for each sample when analyzed as a cluster sample. When the graph is produced in color, intervals that include the true value of the population mean are black, and those that do not include the true value are red.

The estimated coverage probability for each procedure is the proportion of intervals that include the true population mean. In Figure 5.1, the estimated coverage probability of the procedure that (incorrectly) treats the data as an SRS is 0.80—substantially smaller than the nominal 0.95 probability of a confidence interval.

5.5 Summary, Tips, and Warnings

Table 5.2 lists the main R functions used in this chapter. We have seen most of these before; the main difference is how the survey design is specified in the *svydesign* function to indicate the clustering.

TABLE 5.2

Functions used for Chapter 5.

Function	Package	Usage
tapply	base	Apply a function to each group of values; groups are defined by the second argument
confint	stats	Calculate confidence intervals, add df for <i>t</i> confidence interval
lm	stats	Fit a linear model to a data set (not using survey methods)
lme	nlme	Fit a random-effects or mixed-effects model to a data set (not using survey methods)
svydesign	survey	Specify the survey design
svymean	survey	Calculate mean and standard error of mean
svytotal	survey	Calculate total and standard error of total
intervals_ex40	SDAResources	Show differences between (incorrect) SRS formulas and (correct) cluster formulas applied to cluster samples

In the **survey** package, clusters are identified in the `id=` argument of the *svydesign* function. The general form of the *svydesign* function for a one-stage cluster sample (without stratification) is:

```
svydesign(id=~psuvar,weights=~weightvar,fpc=rep(N,nobs),data=dataset)
```

where *psuvar* is the name of the variable in *dataset* containing the psu identifiers. The variable *weightvar* contains the weights for the observation units in the data. If an *fpc* is desired (and often it is not), *N* is the number of psus in the population, and *nobs* is the number of psus in the sample.

For multi-stage cluster sampling (again without stratification), the following form of the *svydesign* function will calculate point estimates with the weights in *weightvar* and without replacement standard errors (note the absence of the *fpc* argument):

```
svydesign(id=~psuvar,weights=~weightvar,data=dataset).
```

After specifying the survey design, the *svymean* and *svytotal* functions are used exactly as in other chapters. The only difference for cluster sampling is that you must list the cluster variable(s) in the *svydesign* function.

Tips and Warnings

- Use the *id* argument to specify the clustering, and check that the number of clusters listed when you print the survey design object equals the number of psus in your sample.
- When calculating estimates for one-stage cluster samples, or for two-stage cluster samples using with-replacement variance estimates, include the *weights* argument when specifying the survey design. The weight variable should contain the final weights at the observation level. Check that the sum of the weights approximately equals the number of observation units in the population.
- In general, we recommend calculating with-replacement variances, but the **survey** package functions will also calculate without-replacement variances for the designs discussed in this chapter, where an SRS is taken at each stage of sampling. If you calculate the without-replacement variances for a two-stage cluster design, it is useful to check these by also calculating the with-replacement variance (they should be close).

Sampling with Unequal Probabilities

In this chapter, we discuss how to select a sample with equal or unequal probabilities, and how to compute estimates from an unequal probability sample. The code is in file `ch06.R` on the book website.

Let's start with sample selection. Section 6.1 shows how to select a one-stage cluster sample with equal or unequal probabilities, and Section 6.2 presents two methods for selecting a two-stage sample. We'll look at code for computing estimates for unequal-probability samples in Section 6.3.

6.1 Selecting a Sample with Unequal Probabilities

This section shows how to select a sample of primary sampling units (psus) with unequal probabilities with the *sample* function and with functions from the **sampling** package (Tillé and Matei, 2021). Subsampling all secondary sampling units (ssus) in the selected psus will give a one-stage cluster sample.

6.1.1 Sampling with Replacement

Example 6.2 of SDA. In Chapters 2 through 5, units, whether observation units or clusters, were selected with equal probabilities. In Example 6.2, 5 classes are sampled from 15 classes with probability proportional to size (pps) and with replacement.

Section 2.1 showed how to use the *sample* function to select a simple random sample, with or without replacement. It can also be used to select a with-replacement sample with unequal probabilities by including the optional *prob* argument. Call the function as

```
sample(1:N,n,replace=TRUE,prob=probvar)
```

where N is the number of psus in the population, n is the desired sample size of psus, and *probvar* is a vector of length N that gives the size measures or selection probabilities for the psus. In this example, psus are classes, and *class_size* gives the number of students in the class.

```
data(classes)
classes[1:2,]
##   class class_size
## 1     1         44
## 2     2         33
N<-nrow(classes)
set.seed(78065)
# select 5 classes with probability proportional to class size and with replacement
sample_units<-sample(1:N,5,replace=TRUE,prob=classes$class_size)
```

```

sample_units
## [1] 5 14 6 14 6
mysample<-classes[sample_units,]
mysample
##      class class_size
## 5         5         76
## 14        14        100
## 6         6         63
## 14.1       14        100
## 6.1        6         63
# calculate ExpectedHits and sampling weights
mysample$ExpectedHits<-5*mysample$class_size/sum(classes$class_size)
mysample$SamplingWeight<-1/mysample$ExpectedHits
mysample$psuid<-row.names(mysample)
mysample
##      class class_size ExpectedHits SamplingWeight psuid
## 5         5         76    0.5873261     1.702632     5
## 14        14        100    0.7727975     1.294000    14
## 6         6         63    0.4868624     2.053968     6
## 14.1       14        100    0.7727975     1.294000   14.1
## 6.1        6         63    0.4868624     2.053968    6.1
# check sum of sampling weights
sum(mysample$SamplingWeight)
## [1] 8.398568

```

- Note that classes 6 and 14 both appear twice in the sample. When collecting data in one stage, each student within classes 6 and 14 must be included twice. Otherwise, estimates will be biased. If collecting data in two stages, you would take two independent subsamples from class 6 and two independent subsamples from class 14.

When analyzing the data, make sure you use different psu names for the multiple instances of psus that appear more than once. For this example, you might want to use `row.names(mysample)` as the psu identifier, since it gives a unique name to each sampled psu.

- After selecting the sample, we need to calculate the sampling weights. The *ExpectedHits* variable gives $n\psi_i$, where n is the sample size and ψ_i is the draw-by-draw selection probability that is proportional to *class_size* (note that the values of ψ_i sum to 1). This is the number of times we expect the unit to be in the sample. For example, class 5 has $ExpectedHits = 5*76/647 = 0.5873261$.

Then, the *SamplingWeight* is $1/ExpectedHits = 1/(n\psi_i)$ for with-replacement sampling.

- The sum of the sampling weights for the sample is an unbiased estimator of N , and for large samples it should be close to N . For this example, however, the weight sum has a large standard error because the sample size is so small (see Exercise 6.45 of SDA).

6.1.2 Sampling without Replacement

There are several functions in the `sampling` package that will select unequal-probability samples without replacement. When the list of units in the sampling frame is in random order, systematic sampling is likely to produce a sample that behaves like an SRS without replacement. The *cluster* function in the `sampling` package can select a pps sample using systematic sampling.

The following code shows how the *cluster* function is called.

```
set.seed(330582)
cluster(data=classes, clustertype=c("class"), size=5, method="systematic",
        pik=classes$class_size,description=TRUE)
## Number of selected clusters: 5
## Number of units in the population and number of selected units: 15 5
##   class ID_unit   Prob
## 1     1       1 0.3400309
## 2     5       5 0.5873261
## 3     8       8 0.3400309
## 4    11      11 0.3554869
## 5    14      14 0.7727975
```

The following arguments of the *cluster* function are used:

- **pik** is the vector of inclusion probabilities (or a vector of relative unit sizes that can be used to compute the probabilities). In this example, *pik* is *class_size*, the auxiliary variable that gives the size of each class.
- **size=5** requests a sample of 5 units.
- **method="systematic"** describes the method used to select the sample. The *cluster* function can also be called with methods **"srswor"** (simple random sampling without replacement), **"srswr"** (simple random sampling with replacement), or **"poisson"** (Poisson sampling). The *pik* argument is not needed with methods *srswor* and *srswr*.
- **description=TRUE** asks the function to print the number of population and sampled units.

The *cluster* function creates variable *Prob* that includes the final inclusion probabilities for the units in the sample. You can compute sampling weights as $1/Prob$.

Other functions are also available for selecting unequal-probability samples. Table 6.1 lists sample selection functions in the **sampling** package that correspond to methods discussed in Chapters 5 and 6 of SDA. Tillé (2006) describes these and additional methods in the **sampling** package (the functions that select unequal-probability samples have names that begin with *UP*) for selecting samples. Another resource is the **pps** package (Gambino, 2021), which contains several functions for selecting unequal-probability samples.

You may want to consider writing your own function or using a different software package if none of the methods implemented in R meet your sample selection needs. The SURVEYSELECT procedure in SAS software provides additional options for selecting without-replacement probability samples, including the Hanurav–Vijayan method (Hanurav, 1967; Vijayan, 1968); see SAS Institute Inc. (2021) and Lohr (2022) for details.

6.2 Selecting a Two-Stage Cluster Sample

There are several ways to select a two-stage cluster sample in R. The *mstage* function from the **sampling** package will select both stages at once for simple random, systematic, or Poisson sampling. Alternatively, you can select the units at each stage separately: First select the psus, then select a sample of ssus from the selected psus.

TABLE 6.1

Some functions for selecting a probability sample in the **sampling** package.

Function	Description
UPbrewer(pik)	Select an unequal-probability sample without replacement containing 2 psus per stratum using Brewer's (1963, 1975) method. The <i>pik</i> argument contains the desired inclusion probabilities π_i . Note that there is no argument for the sample size. For all of the UP sample selection methods, the sample size is assumed to be implicit in the <i>pik</i> vector because $\sum_{k=1}^N \pi_k = n$. The function <code>inclusionprobabilities(size,n)</code> will compute <i>pik</i> from a vector <i>size</i> of positive numbers and desired sample size <i>n</i> .
UPpoisson(pik)	Select a sample (of variable size) using Poisson sampling. The <i>pik</i> argument contains the desired inclusion probabilities π_k for each unit, and these should be between 0 and 1.
UPsystematic(pik)	Select an unequal-probability sample via systematic sampling.
UPSampford(pik)	Select an unequal-probability sample using Sampford's (1967) method, an extension of Brewer's method that allows drawing more than 2 psus per stratum.
srswor(n,N)	Select an SRS of size <i>n</i> without replacement from a population of size <i>N</i> .
srswr(n,N)	Select an SRS of size <i>n</i> with replacement from a population of size <i>N</i> .

Example 6.11 of SDA: Selection with *mstage* function. In Chapter 1, we expanded data classes to a long format that includes student information. Let's redo that here.

```
# create data frame classeslong
data(classes)
classeslong<-classes[rep(1:nrow(classes),times=classes$class_size),]
classeslong$studentid <- sequence(classes$class_size)
nrow(classeslong)
## [1] 647
table(classeslong$class) # check class sizes
##
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15
## 44 33 26 22 76 63 20 44 54 34 46 24 46 100 15
head(classeslong)
##      class class_size studentid
## 1      1         1         44      1
## 1.1    1         1         44      2
## 1.2    1         1         44      3
## 1.3    1         1         44      4
## 1.4    1         1         44      5
## 1.5    1         1         44      6
```

We now use the *mstage* function to select a pps systematic sample of 5 classes and take an SRS without replacement of 4 students from each class. Each class in the psu sample can therefore be considered as a stratum for sample selection purposes, and an independent sample of size 4 is taken from each stratum. We call the function as

```
mstage(classeslong,stage=list("cluster","stratified"),
       varnames=list("class","studentid"),
       size=numbersselect, method=list("systematic","srswor"),pik=prob)
```

In the function *mstage*, sampling specifications for the different stages are given in list objects. Lists in R allow you to combine structures of different type; here, the lists consist of vectors that have different lengths. If there are four stages of sampling, each list will have four components, each in the order of the stages of sampling. This example is for two-stage sampling, so each list has two components.

- The two stages are denoted by `stage=list("cluster","stratified")`, and the corresponding list naming the stratification or clustering variables at the stages is `varnames=list("class","studentid")`.

For this example, a cluster sample of psus, identified by *class*, is desired for the first stage of sampling. After the psus are selected, the frame for the second stage of sampling consists of the listing of ssus for the sample of 5 psus. An SRS is selected from each of those 5 psus, so the method used is "stratified".

- The desired sample sizes are given in the *size* argument as a list containing two levels. Here, we set *size* equal to the list `numbersselect<-list(5,rep(4,5))`. We want to select a sample of 5 psus and then a subsample of 4 ssus from each sampled psu.
- The sampling methods at the two stages, systematic and simple random sampling respectively, are denoted by `method=list("systematic","srswor")`.
- The selection probabilities are given in `pik=prob`, where

```
prob<-list(classes$class_size/647,4/classeslong$class_size)
or
prob<-list(classes$class_size/647)
# srswor is with probability of 4/classeslong$class_size by default
# since ssu sample size of 4 is supplied in numbersselect argument
```

As always, you can set the seed to any integer you like; this allows you to re-create the sample later. Note that the sample in this book has different psus than the sample in SDA, which was selected using SAS software (SAS Institute Inc., 2021).

```
# select a two-stage cluster sample, psu: class, ssu: studentid
# number of psus selected: n = 5 (pps systematic)
# number of students selected: m_i = 4 (srs without replacement)
# problist<-list(classes$class_size/647) # same results as next command
problist<-list(classes$class_size/647,4/classeslong$class_size) #selection prob
problist[[1]] # extract the first object in the list. This is pps, size M_i/M
## [1] 0.06800618 0.05100464 0.04018547 0.03400309 0.11746522 0.09737249
## [7] 0.03091190 0.06800618 0.08346213 0.05255023 0.07109737 0.03709428
## [13] 0.07109737 0.15455951 0.02318393
problist[[2]][1:5] # first 5 values in second object in list
## [1] 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909
# number of psus and ssus
n<-5
numbersselect<-list(n,rep(4,n))
numbersselect
## [[1]]
## [1] 5
##
## [[2]]
```

```
## [1] 4 4 4 4 4
# two-stage sampling
set.seed(75745)
tempid<-mstage(classeslong,stage=list("cluster","stratified"),
               varnames=list("class","studentid"),
               size=numbersselect, method=list("systematic","srswor"),pik=problast)
```

The output *tempid* contains two objects

```
sample1<-getdata(classeslong,tempid)[[1]]
sample2<-getdata(classeslong,tempid)[[2]].
```

Here, *sample1* contains the classes selected at stage 1 along with the selection probabilities *Prob_1_stage*, and *sample2* contains the selected ssus and the second-stage sampling probabilities, as well as the final selection probabilities *Prob* (which equals the product of *Prob_1_stage* and *Prob_2_stage*). We only need *sample2* but also show *sample1* so you can see the first-stage probabilities.

```
# get data
sample1<-getdata(classeslong,tempid)[[1]]
# sample 1 contains the ssus of the 5 psus chosen at the first stage
# Prob_1_stage has the first-stage selection probabilities
head(sample1)
##      class_size studentid class ID_unit Prob_1_stage
## 4.21          22         22    4     125    0.1700155
## 4.20          22         21    4     124    0.1700155
## 4.6           22          7    4     110    0.1700155
## 4            22          1    4     104    0.1700155
## 4.7           22          8    4     111    0.1700155
## 4.8           22          9    4     112    0.1700155
nrow(sample1)
## [1] 285
table(sample1$class) # lists the psus selected in the first stage
##
##  4   6   9  13  14
## 22 63 54 46 100
sample2<-getdata(classeslong,tempid)[[2]]
# sample 2 contains the final sample
# Prob_2_stage has the second-stage selection probabilities
# Prob has the final selection probabilities
head(sample2)
##      class class_size studentid ID_unit Prob_2_stage      Prob
## 4.21      4          22         22     125    0.18181818 0.0309119
## 4.7       4          22          8     111    0.18181818 0.0309119
## 4.5       4          22          6     109    0.18181818 0.0309119
## 4.19      4          22         20     123    0.18181818 0.0309119
## 6.48      6          63         49     250    0.06349206 0.0309119
## 6.53      6          63         54     255    0.06349206 0.0309119
nrow(sample2) # sample of 20 ssus altogether
## [1] 20
table(sample2$class) # 4 ssus selected from each psu
##
##  4   6   9  13  14
##  4   4   4   4   4
# calculate final weight = 1/Prob
sample2$finalweight<-1/sample2$Prob
```

```
# check that sum of final sampling weights equals population size
sum(sample2$finalweight)
## [1] 647
sample2[,c(1,2,3,6,7)] # print variables from final sample
##      class class_size studentid      Prob finalweight
## 4.21      4         22        22 0.0309119        32.35
## 4.7       4         22         8 0.0309119        32.35
## 4.5       4         22         6 0.0309119        32.35
## 4.19      4         22        20 0.0309119        32.35
## 6.48      6         63        49 0.0309119        32.35
## 6.53      6         63        54 0.0309119        32.35
## 6.23      6         63        24 0.0309119        32.35
## 6.33      6         63        34 0.0309119        32.35
## 9.50      9         54        51 0.0309119        32.35
## 9.29      9         54        30 0.0309119        32.35
## 9.31      9         54        32 0.0309119        32.35
## 9.36      9         54        37 0.0309119        32.35
## 13.10     13         46        11 0.0309119        32.35
## 13       13         46         1 0.0309119        32.35
## 13.45     13         46        46 0.0309119        32.35
## 13.39     13         46        40 0.0309119        32.35
## 14.4      14        100         5 0.0309119        32.35
## 14.78     14        100        79 0.0309119        32.35
## 14.98     14        100        99 0.0309119        32.35
## 14.63     14        100        64 0.0309119        32.35
```

The final sampling weight *finalweight* is the reciprocal of final selection probability. The psus were selected with probabilities $5M_i/647$ and the ssus for psu i were selected with probability $4/M_i$, so the final weight is $[647/(5M_i)](M_i/4) = 647/20 = 32.35$ for each ssu in the sample.

Example 6.11 of SDA: Selection in two steps. Another option, if you want to use a method other than systematic sampling to select the psus with unequal probabilities, is to select the first-stage units and the second-stage units in separate steps. This is often more convenient for populations where the psu sizes M_i are known only for units in the sample. For example, if nursing homes are psus, you may have to find out the value of M_i directly from each home and thus would know these values only after the first-stage sample is selected.

Let's select a sample from data set *classes* in two stages. In this example we use function *UPsampford* (see Table 6.1) instead of systematic sampling to select 5 classes (psus) at the first stage of sample selection. We call

```
UPsampford(pik)
```

where *pik* is a vector of length N containing the desired inclusion probabilities. The function has no argument for the sample size; it is assumed that $\text{sum}(\mathbf{pik})=n$. The function *inclusionprobabilities* will compute *pik*, having sum n , from the desired sample size and a vector of positive numbers that gives the relative sizes of the units.

```
# select a cluster sample in two stages, psu: class, ssu: studentid
# number of psu selected n =5 (Sampford's method)
# first, convert the measure of size to a vector of probabilities
classes$stage1prob<-inclusionprobabilities(classes$class_size,5)
sum(classes$stage1prob) # inclusion probabilities sum to n
## [1] 5
```

```
# select the psus
set.seed(29385739)
stage1.units<-UPsampford(classes$stage1prob)
stage1.sample<-getdata(classes,stage1.units)
stage1.sample
##      ID_unit class class_size stage1prob
## 1         1     1         44  0.3400309
## 3         3     3         26  0.2009274
## 7         7     7         20  0.1545595
## 13        13    13         46  0.3554869
## 14        14    14        100  0.7727975
```

The data frame *stage1.sample* contains the psus (classes) for the sample. Variable *stage1prob* contains the first-stage selection probabilities that we computed from the *class_size* variable. Now we can use the function to select the second-stage units. Since we have already formed the data frame *stage1.sample* that consists only of the sampled psus, the *strata* function provides a convenient way to select an SRS from each sampled psu.

To draw an SRS of students from each sampled psu, we first create data frame *stage1.long* with a data record for each student in the sampled classes, and then take an SRS of 4 students from each class.

```
# first-stage units are in stage1.sample
# now select the second-stage units (students)
# convert the psus in the sample to long format and assign student ids
npsu<-nrow(stage1.sample)
stage1.long<-stage1.sample[rep(1:npsu,times=stage1.sample$class_size),]
stage1.long$studentid<-sequence(stage1.sample$class_size)
head(stage1.long)
##      ID_unit class class_size stage1prob studentid
## 1         1     1         44  0.3400309         1
## 1.1       1     1         44  0.3400309         2
## 1.2       1     1         44  0.3400309         3
## 1.3       1     1         44  0.3400309         4
## 1.4       1     1         44  0.3400309         5
## 1.5       1     1         44  0.3400309         6
# use strata function to select 4 ssus from each psu
stage2.units<-strata(stage1.long, stratanames=c("class"),
                     size=rep(4,5), method="srswor")
nrow(stage2.units)
## [1] 20
# get the data for the second-stage sample
ssusample<-getdata(stage1.long,stage2.units)
head(ssusample)
##      class_size stage1prob studentid class ID_unit      Prob Stratum
## 1.3          44  0.3400309         4     1         4 0.09090909         1
## 1.13         44  0.3400309        14     1        14 0.09090909         1
## 1.21         44  0.3400309        22     1        22 0.09090909         1
## 1.26         44  0.3400309        27     1        27 0.09090909         1
## 3.11         26  0.2009274        12     3        56 0.15384615         2
## 3.18         26  0.2009274        19     3        63 0.15384615         2
```

The last step is computing the final selection probability, accounting for both stages of sampling, and the final sampling weight. In data frame *ssusample*, variable *stage1prob* contains the psu-level sampling probability (we defined this variable earlier and used it to select the psus) and variable *Prob* contains the ssu-level sampling probability (this is computed by the

strata function and for this example equals $4/M_i$). Thus, the selection probability for each student in the sample is the product $stage1prob \times Prob$. The final weight is the reciprocal of the final selection probability, which for this example equals 32.35 for all students in the sample because the first-stage sample was selected with probability proportional to M_i and the same subsample size (here, $m_i = 4$) was selected from each psu.

```
# compute the sampling weights
# stage1prob contains stage 1 sampling probability;
# Prob has stage 2 sampling probability
ssusample$finalprob<- ssusample$stage1prob*ssusample$Prob
ssusample$finalwt<-1/ssusample$finalprob
sum(ssusample$finalwt) # check sum of weights
## [1] 647
# print selected columns of ssusample
print(ssusample[,c(1,2,3,4,6,8,9)],digits=4)
##      class_size stage1prob studentid class      Prob finalprob finalwt
## 1.3           44      0.3400          4      1 0.09091   0.03091   32.35
## 1.13          44      0.3400         14      1 0.09091   0.03091   32.35
## 1.21          44      0.3400         22      1 0.09091   0.03091   32.35
## 1.26          44      0.3400         27      1 0.09091   0.03091   32.35
## 3.11          26      0.2009         12      3 0.15385   0.03091   32.35
## 3.18          26      0.2009         19      3 0.15385   0.03091   32.35
## 3.19          26      0.2009         20      3 0.15385   0.03091   32.35
## 3.24          26      0.2009         25      3 0.15385   0.03091   32.35
## 7.11          20      0.1546         12      7 0.20000   0.03091   32.35
## 7.13          20      0.1546         14      7 0.20000   0.03091   32.35
## 7.18          20      0.1546         19      7 0.20000   0.03091   32.35
## 7.19          20      0.1546         20      7 0.20000   0.03091   32.35
## 13.16         46      0.3555         17     13 0.08696   0.03091   32.35
## 13.31         46      0.3555         32     13 0.08696   0.03091   32.35
## 13.34         46      0.3555         35     13 0.08696   0.03091   32.35
## 13.42         46      0.3555         43     13 0.08696   0.03091   32.35
## 14.1          100      0.7728          2     14 0.04000   0.03091   32.35
## 14.20         100      0.7728         21     14 0.04000   0.03091   32.35
## 14.35         100      0.7728         36     14 0.04000   0.03091   32.35
## 14.68         100      0.7728         69     14 0.04000   0.03091   32.35
```

6.3 Computing Estimates from an Unequal-Probability Sample

The syntax used to compute estimates from an unequal-probability cluster sample is largely the same as that used in Chapter 5 for equal-probability cluster samples. The *svymean* and *svytotal* functions of the **survey** package (Lumley, 2020) calculate estimates of means, totals, and proportions by using the formulas with survey weights. When the *fpc* argument is omitted from the *svydesign* function call, standard errors are calculated with the formulas for the with-replacement variance in Section 6.4 of SDA.

6.3.1 Estimates from with-Replacement Samples

Example 6.4 of SDA. This example shows how to calculate estimates when the cluster total t_i has already been found for each psu (or when the psus are also the observation units, that is, $M_i = 1$ for all psus). Since the summary statistic has already been calculated for

each psu, the *svydesign* function is called with `id=~1`. We only need to specify the unequal weights using the *weights* argument to calculate the estimates. Class 14 appears twice in the data since it was selected twice for the sample—we call it class 141 for the first appearance and class 142 for the second to distinguish them.

The mean calculated from *svymean* estimates $\bar{t}_{\mathcal{U}}$, the population mean of the cluster totals t_i , which for this example is the total amount of time spent studying by students in class i . The average amount of time spent studying per student is estimated by the ratio $\hat{y}_{\psi} = \hat{t}_{\psi} / \hat{M}_{\psi}$. The *svyratio* function can give the estimate $\hat{y}_{\psi}$. (If the data set consists of the individual values y_{ij} instead of the summary statistics, then the mean $\hat{y}_{\psi}$ can be estimated directly from *svymean*, as seen in the code for Example 6.6 of SDA.)

```
studystat <- data.frame(class = c(12, 141, 142, 5, 1),
                      Mi = c(24, 100, 100, 76, 44),
                      tothours=c(75,203,203,191,168))
studystat$wt<-647/(studystat$Mi*5)
sum(studystat$wt) # check weight sum, which estimates N=15 psus
## [1] 12.62321
# design for with-replacement sample, no fpc argument
d0604 <- svydesign(id = ~1, weights=~wt, data = studystat)
d0604
## Independent Sampling design (with replacement)
## svydesign(id = ~1, weights = ~wt, data = studystat)
# Ratio estimation using Mi as auxiliary variable
ratio0604<-svyratio(~tothours, ~Mi,design = d0604)
ratio0604
## Ratio estimator: svyratio.survey.design2(~tothours, ~Mi, design = d0604)
## Ratios=
##              Mi
## tothours 2.703268
## SEs=
##              Mi
## tothours 0.3437741
confint(ratio0604, level=.95,df=4)
##              2.5 %    97.5 %
## tothours/Mi 1.748798 3.657738
# Can also estimate total hours studied for all students in population
svytotal(~tothours,d0604)
##              total      SE
## tothours 1749 222.42
```

The average amount of time a student spent studying statistics is estimated as 2.70 hours with an estimated standard error of 0.34 hours and a 95% confidence interval of [1.74, 3.66]. Note that 4 degrees of freedom (df; here, 1 less than the number of psus) are used for the confidence interval.

Example 6.6 of SDA. The estimates for a two-stage cluster sample with replacement are calculated exactly the same way as for a one-stage sample. For this example, we have data for the individual students in the psus so we enter those for each student.

Class 14 appears twice in the sample of psus in Example 6.4 of SDA. An independent set of students is selected for each appearance. To enable correct variance calculations, the first occurrence of class 14 is relabeled as class 141, and the second occurrence as class 142. These are counted as two separate psus in the estimation. If you labeled both as 14, then the *id* argument of *svydesign* would treat that as one psu with $m_i = 10$ instead of two psus of size 5.

The weight *studentwt* is calculated as the first-stage weight $M_0/(nM_i)$ times the second-stage weight M_i/m_i . The sample is self-weighting and the weight for each student simplifies to 647/25. For many problems, defining the weights is the trickiest part, and it is also the most important. Always check that the sum of the weights approximately (or exactly, in this case) equals the population size.

```
students <- data.frame(class = rep(studystat$class,each=5),
  popMi = rep(studystat$Mi,each=5),
  sampmi=rep(5,25),
  hours=c(2,3,2.5,3,1.5,2.5,2,3,0,0.5,3,0.5,1.5,2,3,1,2.5,3,5,2.5,4,4.5,3,2,5))
# The 'with' function allows us to calculate using variables from a data frame
# without having to type the data frame name for all of them
students$studentwt <- with(students,(647/(popMi*5)) * (popMi/sampmi))
# check the sum of the weights
sum(students$studentwt)
## [1] 647
# create the design object
d0606 <- svydesign(id = ~class, weights=~studentwt, data = students)
d0606
## 1 - level Cluster Sampling design (with replacement)
## With (5) clusters.
## svydesign(id = ~class, weights = ~studentwt, data = students)
# estimate mean and SE
svymean(~hours,d0606)
##          mean          SE
## hours    2.5 0.3606
degf(d0606)
## [1] 4
confint(svymean(~hours,d0606),level=.95,df=4) #use t-approximation
##          2.5 %    97.5 %
## hours 1.498938 3.501062
# estimate total and SE
svytotal(~hours,d0606)
##          total          SE
## hours 1617.5 233.28
confint(svytotal(~hours,d0606),level=.95,df=4)
##          2.5 %    97.5 %
## hours 969.8132 2265.187
```

In the *svydesign* function, we supply the weights (which, for this example, are the same for all students) but no *fpc* argument. We specify the psu membership in the *id* argument. This means that the variability is calculated at the first stage level using the pps with-replacement formulas, that is, the variability among $\hat{t}_i/\psi_i$. Note that 4 df (1 less than the number of psus) are used for the confidence interval.

6.3.2 Estimates from without-Replacement Samples

Even when an unequal-probability sample was selected without replacement, the with-replacement variance is commonly calculated for simplicity and stability. Use the *weights* argument to provide the sampling weights at the observation-unit level, and use the *id=~psuid* argument to provide the information on psu membership (recall that only the psu membership is needed to calculate the with-replacement variance).

Example 6.11 of SDA. This example analyzes the without-replacement unequal-probability sample the same way as for the sample in Example 6.6. Even though the sample was selected

without replacement, the with-replacement variance provides a good approximation.

```
data(classpps)
nrow(classpps)
## [1] 20
head(classpps)
##   class class_size finalweight hours
## 1     4         22      32.35   5.0
## 2     4         22      32.35   4.5
## 3     4         22      32.35   5.5
## 4     4         22      32.35   5.0
## 5    10         34      32.35   2.0
## 6    10         34      32.35   4.0
d0611 <- svydesign(ids = ~class, weights=~classpps$finalweight, data = classpps)
d0611
## 1 - level Cluster Sampling design (with replacement)
## With (5) clusters.
## svydesign(ids = ~class, weights = ~classpps$finalweight, data = classpps)
# estimate mean and SE
svymean(~hours,d0611)
##      mean      SE
## hours 3.45 0.4819
confint(svymean(~hours,d0611),level=.95,df=4) #use t-approximation
##      2.5 %   97.5 %
## hours 2.112147 4.787853
# estimate total and SE
svytotal(~hours,d0611)
##      total      SE
## hours 2232.2 311.76
confint(svytotal(~hours,d0611),level=.95,df=4)
##      2.5 %   97.5 %
## hours 1366.559 3097.741
```

Calculating the without-replacement variance for a one-stage sample. In general, we recommend calculating the with-replacement variance estimate and omitting the *fpc* argument from *svydesign* when unequal-probability sampling is used. Most of the replication methods for calculating variances in Chapter 9 also calculate with-replacement variances. You can skip the remainder of this section if the with-replacement variances work for your applications.

Functions in the **survey** package will calculate the without-replacement variance for some one-stage designs if you specify the inclusion probabilities in the *fpc* argument. (As of this writing, the package will not yet calculate without-replacement variances for two-stage designs—the situation where unequal-probability sampling is most commonly used.) The formulas for calculating the Horvitz-Thompson (HT), Sen-Yates-Grundy (SYG), and other without-replacement variance estimates require knowledge of the joint inclusion probabilities, so you must also supply those to the *svydesign* function.

Let's calculate some joint inclusion probabilities first and then use them in the *svydesign* function. Functions in the **sampling** package will calculate joint inclusion probabilities for some of the sample-selection methods; for example, function *UPsampfordpi2* will calculate the joint inclusion probabilities for a sample selected using Sampford's method. For some other sample-selection methods, however, the joint inclusion probabilities must be calculated directly.

Example 6.8 of SDA. We use the supermarket example to illustrate the calculation of joint inclusion probabilities when the sample size is 2. We first create a data frame of the supermarket population containing the store identifier, area, and revenue.

```
supermarket<-data.frame(store=c('A','B','C','D'),area=c(100,200,300,1000),
                        ti=c(11,20,24,245))

supermarket
##   store area  ti
## 1     A  100  11
## 2     B  200  20
## 3     C  300  24
## 4     D 1000 245
```

The draw-by-draw method was used to select two supermarkets for the sample, where the selection probability for draw 1 was proportional to store area. We can use that information to calculate π_i , the probability of store i being included in the sample, and π_{ik} , the joint probability that stores i and k are both included in the sample.

Here, we use matrix operations to calculate the probabilities by applying the formulas in Example 6.8 of SDA, noting that if $\mathbf{a}$ and $\mathbf{b}$ are two vectors, the (i, j) entry of $\mathbf{ab}^T$ is $a_i b_j$. The *apply* function sums the entries in each column.

```
supermarket$psi<-supermarket$area/sum(supermarket$area)
psii<-supermarket$area/sum(supermarket$area)
piik<- psii %*% t(psii/(1-psii)) + (psii/(1-psii)) %*% t(psii)
diag(piik)<-rep(0,4) # set the diagonal entries of the matrix equal to zero
piik # joint inclusion probabilities
##           [,1]      [,2]      [,3]      [,4]
## [1,] 0.00000000 0.01726190 0.02692308 0.14583333
## [2,] 0.01726190 0.00000000 0.05563187 0.29761900
## [3,] 0.02692308 0.05563187 0.00000000 0.45673080
## [4,] 0.14583333 0.29761905 0.45673077 0.00000000
pii<-apply(piik,2,sum)
pii # inclusion probabilities
## [1] 0.1900183 0.3705128 0.5392857 0.9001832
```

The results show that $\pi_1 = 0.19$, $\pi_2 = 0.37$, $\pi_3 = 0.539$, and $\pi_4 = 0.90$. The joint inclusion probabilities are given in *piik*: for example, $\pi_{12} = 0.01726$. These are the numbers shown in Table 6.6 of SDA.

Now let's use the values of π_i and π_{ik} to calculate the HT and SYG variance estimates. Of course, since the supermarket sample has only two units, neither estimate will be very accurate, but it will serve to illustrate the methods.

Example 6.9 of SDA. Suppose supermarkets C and D were selected from the population in Example 6.8 of SDA. We will calculate the Horvitz–Thompson (HT) estimate for total revenue and the without-replacement HT and SYG variance estimates.

As always, we specify all the information about the design in the *svydesign* function. We tell the function that this is an unequal-probability sample without replacement through the *fpc* argument. Instead of giving the population sizes in the *fpc* argument, however, for pps sampling without replacement we specify *fpc*=~*pii*, the inclusion probability for each sampled unit.

We also use two other arguments to the *svydesign* function that we have not seen before. The *variance* argument tells whether to calculate the HT or SYG (the function calls this “YG”) formula for the variance. We supply the joint inclusion probabilities in the *pps* argument,

after first placing π_i on the diagonal elements of the joint probabilities matrix and using the function *ppsmat* to get the joint probabilities in the form required by *svydesign*.

```
supermarket2<-supermarket[3:4,]
supermarket2$pii <- pii[3:4] # these are the unit inclusion probs when n=2
jointprob<-piik[3:4,3:4] # joint probability matrix for stores C and D
diag(jointprob)<-supermarket2$pii # set diagonal entries equal to pii
jointprob
##           [,1]      [,2]
## [1,] 0.5392857 0.4567308
## [2,] 0.4567308 0.9001832
# Horvitz-Thompson type
dht<- svydesign(id=~1, fpc=~pii, data=supermarket2,
               pps=ppsmat(jointprob),variance="HT")
dht
## Sparse-matrix design object:
## svydesign(id = ~1, fpc = ~pii, data = supermarket2, pps = ppsmat(jointprob),
## variance = "HT")
svytotal(~ti,dht)
##      total      SE
## ti 316.67 82.358
# Sen-Yates-Grundy type
dsyg<- svydesign(id=~1, fpc=~pii, data=supermarket2,
               pps=ppsmat(jointprob),variance="YG")
dsyg
## Sparse-matrix design object:
## svydesign(id = ~1, fpc = ~pii, data = supermarket2, pps = ppsmat(jointprob),
## variance = "YG")
svytotal(~ti,dsyg)
##      total      SE
## ti 316.67 57.094
```

We can compare the variance estimates from the two methods with the true without-replacement variance $V(\hat{t}_{HT}) = 4383.6$ in SDA (which is known for this small example where the full population is known), with $\hat{V}_{HT}(\hat{t}_{HT}) = (82.358)^2 = 6782.8$, and $\hat{V}_{SYG}(\hat{t}_{HT}) = (57.094)^2 = 3259.8$. In most situations, the SYG variance estimate is preferred because it is more stable.

The *svydesign* function also provides some approximation methods to calculate without-replacement variance estimates for one-stage samples. Option `pps=HR(sum(pii$sq)/n)`, where *pii\$sq* is the vector of squared inclusion probabilities and *n* is the number of psus selected, gives the Hartley and Rao (1962; see Exercise 6.36 in SDA) approximation to the variance.

Example 6.10 of SDA. Let's do one more example, to compare the with-replacement, HT, and SYG variance estimates calculated for the unequal-probability sample in data *agpps*, as well as the Hartley–Rao approximation. We would expect the with-replacement variance estimate to work well here because $n = 15$ is small relative to $N = 3078$.

```
data(agpps)
jtprobag<-as.matrix(agpps[,20:34])
diag(jtprobag)<-agpps$SelectionProb
# Horvitz-Thompson type
dhtag<- svydesign(id=~1, fpc=~SelectionProb, data=agpps,
               pps=ppsmat(jtprobag),variance="HT")
svytotal(~acres92,dhtag)
```

```
##          total      SE
## acres92 936291172 70466858
# Sen-Yates-Grundy type
dsyagag<- svydesign(id=~1, fpc=~SelectionProb, data=agpps,
                  pps=ppsmat(jtprobag),variance="YG")
svytotal(~acres92,dsyagag)
##          total      SE
## acres92 936291172 11715201
# Hartley-Rao approximation
sumsqprob<-sum(agpps$SelectionProb^2)/nrow(agpps)
dHRag<-svydesign(id=~1, fpc=~SelectionProb, data=agpps,
                pps=HR(sumsqprob),variance="YG")
svytotal(~acres92,dHRag)
##          total      SE
## acres92 936291172 12148234
# With-replacement variance
dwrag<-svydesign(id=~1, weights=~SamplingWeight, data=agpps)
svytotal(~acres92,dwrag)
##          total      SE
## acres92 936291172 12293009
```

Note that the with-replacement (12,293,009), SYG (11,715,201), and Hartley–Rao (12,148,234) standard errors are all similar to each other. The HT standard error (70,466,858) is larger and often less stable; in general, we recommend one of the other methods.

6.4 Summary, Tips, and Warnings

Several functions in the **sampling** package will select equal- and unequal-probability cluster samples; some of these are listed in Table 6.1. The *sample* function can be used to select with-replacement unequal-probability samples.

Table 6.2 lists the major R functions used in this chapter.

Tips and Warnings

- When selecting an unequal-probability sample, check the calculation of the selection probabilities to make sure these are roughly proportional to the unit sizes.
- The more complex the sampling plan, the more complicated the weight calculations. Check that the sum of the weights approximately equals the population size.
- For unequal-probability sampling, omitting the *fpc* argument in the *svydesign* function gives the with-replacement variance. In general, this is the approach that we recommend. If the without-replacement variance is desired, use the Sen-Yates-Grundy formula directly.

TABLE 6.2

Functions used for Chapter 6.

Function	Package	Usage
sample	base	Select a with-replacement sample with unequal probabilities
confint	stats	Calculate confidence intervals, add df for t confidence interval
apply	base	Apply a function to the rows or columns of a matrix
cluster	sampling	Select a cluster sample
strata	sampling	Select a stratified random sample (here used to select ssus from the sampled psus)
mstage	sampling	Select a multi-stage cluster sample
UPsampford	sampling	Select an unequal-probability sample of units using Sampford's method
inclusionprobabilities	sampling	Convert a vector of positive size measures to selection probabilities, for use in the UP selection functions
getdata	sampling	Extract the data after selecting a sample
svydesign	survey	Specify the survey design
svymean	survey	Calculate mean and standard error of mean
svyratio	survey	Calculate ratio and standard error of ratio
svytotal	survey	Calculate total and standard error of total

A

Data Set Descriptions

The data sets referenced in SDA and described in this appendix are available from the book website (see the Preface for the website address) and in the contributed R package `SDAResources` (Lu and Lohr, 2021). These data sets are provided for instructional purposes only and without warranty. Anyone wishing to investigate the subject matter further should obtain the original data from the source. In some cases, the data sets referenced in SDA and this book are a subset of the original data; in others, the information has been modified to protect the confidentiality of the respondents.

All data sets ending in `.csv` use commas as a separator between fields.

These data sets have also been stored in SAS format (with the name ending in `.sas7bdat`) and R format with missing values recoded to the symbols used for missing data in the software package (`.` or blank in SAS and `NA` in R).

agpop.csv Data from the 1992 U.S. Census of Agriculture. Source: U.S. Bureau of the Census (1995). In columns 3–14, the value `–99` denotes missing data.

Column	Name	Value
1	county	county name (character variable)
2	state	state abbreviation (character variable)
3	acres92	number of acres devoted to farms, 1992
4	acres87	number of acres devoted to farms, 1987
5	acres82	number of acres devoted to farms, 1982
6	farms92	number of farms, 1992
7	farms87	number of farms, 1987
8	farms82	number of farms, 1982
9	largef92	number of farms with 1,000 acres or more, 1992
10	largef87	number of farms with 1,000 acres or more, 1987
11	largef82	number of farms with 1,000 acres or more, 1982
12	smallf92	number of farms with 9 acres or fewer, 1992
13	smallf87	number of farms with 9 acres or fewer, 1987
14	smallf82	number of farms with 9 acres or fewer, 1982
15	region	S = south; W = west; NC = north central; NE = northeast

agpps.csv Data from a without-replacement probability-proportional-to-size sample from file `agpop.csv`.

Column	Name	Value
1	county	county name
2	state	state abbreviation
3	acres92	number of acres devoted to farms, 1992
4	acres87	number of acres devoted to farms, 1987
5–15	...	same as variables 5–15 in <code>agpop.csv</code>
16	sizemeas	size measure used to select the pps sample
17	SelectionProb	inclusion probability for county i , π_i
18	SamplingWeight	sampling weight for county i , $w_i = 1/\pi_i$
19	Unit	unit number for indexing joint inclusion probabilities
20–34	JtProb_1– JtProb_15	columns of joint inclusion probabilities

agsrs.csv Data from an SRS of size 300 from the 1992 U.S. Census of Agriculture. Variables are the same as in `agpop.csv`. In columns 3–14, the value -99 denotes missing data.

agstrat.csv Data from a stratified random sample of size 300 from the 1992 U.S. Census of Agriculture data in `agpop.csv`. In columns 3–14, the value -99 denotes missing data.

Column	Name	Value
1	county	county name
2	state	state abbreviation
3	acres92	number of acres devoted to farms, 1992
4	acres87	number of acres devoted to farms, 1987
5	acres82	number of acres devoted to farms, 1982
6	farms92	number of farms, 1992
7	farms87	number of farms, 1987
8	farms82	number of farms, 1982
9	largef92	number of farms with 1,000 acres or more, 1992
10	largef87	number of farms with 1,000 acres or more, 1987
11	largef82	number of farms with 1,000 acres or more, 1982
12	smallf92	number of farms with 9 acres or fewer, 1992
13	smallf87	number of farms with 9 acres or fewer, 1987
14	smallf82	number of farms with 9 acres or fewer, 1982
15	region	S = south; W = west; NC = north central; NE = northeast
16	rn	random numbers used to select sample in each stratum
17	strwt	sampling weight for each county in sample

algebra.csv Hypothetical data for an SRS of 12 algebra classes in a city, from a population of 187 classes.

Column	Name	Value
1	class	Class number
2	Mi	Number of students (M_i) in class
3	score	Score of student on test

anthrop.csv Finger length and height for 3,000 criminals. Source: Macdonell (1901). This data set contains information for the entire population.

Column	Name	Value
1	finger	length of left middle finger (cm)
2	height	height (inches)

anthrsrs.csv Finger length and height for an SRS of size 200 from anthrop.csv.

Column	Name	Value
1	finger	length of left middle finger (cm)
2	height	height (inches)
3	wt	sampling weight

anthuneq.csv Finger length and height for a with-replacement unequal-probability sample of size 200 from **anthrop.csv**. The probability of selection, ψ_i , was proportional to 24 for $y < 65$, 12 for $y = 65$, 2 for $y = 66$ or 67, and 1 for $y > 67$.

Column	Name	Value
1	finger	length of left middle finger (cm)
2	height	height (inches)
3	wt	sampling weight

artfratio.csv Values from all possible SRSs for artificial population in Chapter 4 of SDA.

Column	Name	Value
1	sample	sample number
2	i1	first unit in sample
3	i2	second unit in sample
4	i3	third unit in sample
5	i4	fourth unit in sample
6	xbars	$\bar{x}_S$
7	ybars	$\bar{y}_S$
8	bhat	$\hat{B}$
9	tSRS	$\hat{t}_{y,SRS} = N\bar{y}_S$
10	thatr	$\hat{t}_{yr}$

asafellow.csv Information from a stratified random sample of Fellows of the American Statistical Association elected between 2000 and 2018. The list of Fellows serving as the population was downloaded from <https://www.amstat.org/ASA/Your-Career/Awards/ASA-Fellows-list.aspx> on March 18, 2019. All other information was obtained from public sources.

Column	Name	Value
1	awardyr	Year of award
2	gender	Gender of Fellow (character variable, M = male, F = female)

asafellow.csv (continued)

Column	Name	Value
3	popsize	Population size in stratum ($= N_h$)
4	sampsize	Sample size in stratum ($= n_h$)
5	field	Field of employment (character variable) acad = academia, ind = industry, govt = government
6	degreeyr	Year in which Fellow received terminal degree (year of Ph.D. if applicable, otherwise year of Master's or Bachelor's degree)
7	math	= 1 if majored in mathematics as undergraduate, 0 if did not major in math, -99 if missing

auditresult.csv Audit data used in Chapter 6 of SDA.

Column	Name	Value
1	account	audit unit
2	bookvalue	book value of account
3	psi	probability of selection
4	auditvalue	audit value of account

auditselect.csv Selection of accounts for audit data used in Chapter 6 of SDA.

Column	Name	Value
1	account	audit unit
2	bookval	book value of account
3	cumbv	cumulative book value
4	rn1	random number 1 selecting account
5	rn2	random number 2 selecting account
6	rn3	random number 3 selecting account

azcounties.csv Population and housing unit estimates for Arizona counties (excluding Maricopa and Pima counties), from the American Community Survey 2018 5-year estimates. Source: <https://data.census.gov/>, accessed November 27, 2020.

Column	Name	Value
1	number	County number
2	name	County name (character variable, length 15)
3	population	Population estimate for county
4	housing	Housing unit estimate for county
5	ownerocc	Number of owner-occupied housing units for county

baseball.csv Statistics on 797 baseball players, compiled by Jenifer Boshes from the rosters of all major league teams in November 2004. Source: Forman (2004). Missing values (for variables *pball*, *intwalk*, *hbp*, and *sacrflly*; all other variables have complete data) are coded as -9.

Column	Name	Value
1	team	team played for at beginning of the season
2	leagueid	AL or NL
3	player	a unique identifier for each baseball player
4	salary	player salary in 2004
5	pos	primary position coded as P, C, 1B, 2B, 3B, SS, RF, LF, or CF
6	gplay	games played
7	gstart	games started
8	inning	number of innings
9	putout	number of putouts
10	assist	number of assists
11	error	Errors
12	dplay	number of double plays
13	pball	number of passed balls (only applies to catchers)
14	gbat	number of games that player appeared at bat
15	atbat	number of at bats
16	run	number of runs scored
17	hit	number of hits
18	secbase	number of doubles
19	thirdbase	number of triples
20	homerun	number of home runs
21	rbi	number of runs batted in
22	stolenb	number of stolen bases
23	csteal	number of times caught stealing
24	walk	number of times walked
25	strikeout	number of strikeouts
26	intwalk	number of times intentionally walked
27	hbp	number of times hit by pitch
28	sacrhit	number of sacrifice hits
29	sacrfly	number of sacrifice flies
30	gidplay	grounded into double play

books.csv Data from homeowner's survey to estimate total number of books, used in Chapter 5.

Column	Name	Value
1	shelf	shelf number
2	Mi	number of books on shelf
3	booknumber	number of the book selected
4	purchase	purchase cost of book
5	replace	replacement cost of book

census1920.csv Population sizes for each state, from the 1920 U.S. census. The data set contains only the 48 states and excludes Washington D.C., Puerto Rico, and U.S. territories (these areas were not allowed to have voting representatives in Congress). Source: U.S. Bureau of the Census (1921).

Column	Name	Value
1	state	state name
2	population	state population in 1920 census

census2010.csv Population sizes for each state, from the 2010 U.S. census. Source: U.S. Census Bureau (2019). The data set contains only the 50 states and excludes the areas that, as of 2010, were not allowed to have voting representatives in Congress: Washington D.C., Puerto Rico, and U.S. territories.

Column	Name	Value
1	state	state name
2	population	state population in 2010 census

cherry.csv Data for a sample of 31 cherry trees. Source: Hand et al. (1994).

Column	Name	Value
1	diameter	Diameter of tree (inches)
2	height	Height of tree (feet)
3	volume	Timber volume of tree (cubic feet)

classes.csv Population sizes for 15 classes, used in Chapter 6 of SDA to illustrate unequal-probability sampling.

Column	Name	Value
1	class	Class ID number
2	class_size	Number of students in class

classpps.csv Two-stage unequal-probability sample without replacement from the population of classes in `classes.csv`.

Column	Name	Value
1	class	Class ID number
2	class_size	Number of students in class
3	finalweight	Sampling weight for student
4	hours	Number of hours spent studying statistics

classppsjp.csv Joint inclusion probabilities for unequal-probability sample without replacement from the population of classes in `classes.csv`.

Column	Name	Value
1	class	Class ID number
2	class_size	Number of students in class
3	SelectionProb	Probability of being included in sample, π_i
4	SamplingWeight	Sampling weight $w_i = 1/\pi_i$
5–9	JtProb_1–JtProb_5	Columns of joint inclusion probabilities, π_{ik}

college.csv Selected variables from the U.S. Department of Education College Scorecard Data (version updated on June 1, 2020). Source: U.S. Department of Education (2020), downloaded on August 25, 2020. Some of the variables in **college.csv** have been calculated from other variables in the original source; these have been given new variable names that are not found in the data dictionary at <https://collegescorecard.ed.gov/data/documentation/>.

This data set is made available for pedagogical purposes only. Anyone wishing to draw conclusions from College Scorecard data should obtain the full data set from the Department of Education. The original data set has 1,925 variables and includes institutions (such as those that do not grant undergraduate degrees) that are not in the file **college.csv**.

The data set **college.csv** includes institutions in the original data set that: (1) are located in the 50 states plus the District of Columbia, (2) contain information on average net price (NPT4), (3) are predominantly Bachelor's degree-granting, (4) were currently operating as of June 2020, (5) are not private for-profit institutions or "global" campuses, (6) have Carnegie size classification (variable *ccsizset*) between 6 and 17 and Carnegie basic classification (variable *ccbasic*) between 14 and 22 (these offer Bachelor's degrees), (7) enroll first-time students, and (8) are not U.S. Service Academies.

For all variables, missing data are coded as `–99`.

Column	Name	Value
1	unitid	Unit identification number
2	instnm	Institution name (character, length 81)
3	city	City (character, length 24)
4	stabbr	State abbreviation (character, length 2)
5	highdeg	Highest degree awarded 3 = Bachelor's degree, 4 = Graduate degree
6	control	Control (ownership) of institution 1 = Public, 2 = Private nonprofit
7	region	Region where institution is located 1 New England (CT, ME, MA, NH, RI, VT) 2 Mid East (DE, DC, MD, NJ, NY, PA) 3 Great Lakes (IL, IN, MI, OH, WI) 4 Plains (IA, KS, MN, MO, NE, ND, SD) 5 Southeast (AL, AR, FL, GA, KY, LA, MS, NC, SC, TN, VA, WV) 6 Southwest (AZ, NM, OK, TX) 7 Rocky Mountains (CO, ID, MT, UT, WY) 8 Far West (AK, CA, HI, NV, OR, WA)
8	locale	Locale of institution 11 City: Large (population of 250,000 or more) 12 City: Midsize (population of at least 100,000 but less than 250,000) 13 City: Small (population less than 100,000) 21 Suburb: Large (outside principal city, in urbanized area with population of 250,000 or more) 22 Suburb: Midsize (outside principal city, in urbanized area with population of at least 100,000 but less than 250,000)

college.csv (continued)		
Column	Name	Value
		23 Suburb: Small (outside principal city, in urbanized area with population less than 100,000)
		31 Town: Fringe (in urban cluster up to 10 miles from an urbanized area)
		32 Town: Distant (in urban cluster more than 10 miles and up to 35 miles from an urbanized area)
		33 Town: Remote (in urban cluster more than 35 miles from an urbanized area)
		41 Rural: Fringe (rural territory up to 5 miles from an urbanized area or up to 2.5 miles from an urban cluster)
		42 Rural: Distant (rural territory more than 5 miles but up to 25 miles from an urbanized area or more than 2.5 and up to 10 miles from an urban cluster)
		43 Rural: Remote (rural territory more than 25 miles from an urbanized area and more than 10 miles from an urban cluster)
9	ccbasic	Carnegie basic classification
		15 Doctoral Universities: Very High Research Activity
		16 Doctoral Universities: High Research Activity
		17 Doctoral/Professional Universities
		18 Master's Colleges & Universities: Larger Programs
		19 Master's Colleges & Universities: Medium Programs
		20 Master's Colleges & Universities: Small Programs
		21 Baccalaureate Colleges: Arts & Sciences Focus
		22 Baccalaureate Colleges: Diverse Fields
10	ccsizset	Carnegie classification, size and setting
		6 Four-year, very small, primarily nonresidential
		7 Four-year, very small, primarily residential
		8 Four-year, very small, highly residential
		9 Four-year, small, primarily nonresidential
		10 Four-year, small, primarily residential
		11 Four-year, small, highly residential
		12 Four-year, medium, primarily nonresidential
		13 Four-year, medium, primarily residential
		14 Four-year, medium, highly residential
		15 Four-year, large, primarily nonresidential
		16 Four-year, large, primarily residential
		17 Four-year, large, highly residential
11	hbcu	Historically black college or university, 1 = yes, 0 = no
12	openadmp	Does the college have an open admissions policy, that is, does it accept any students that apply or have minimal requirements for admission? 1 = yes, 0 = no
13	adm_rate	Fall admissions rate, defined as the number of admitted undergraduates divided by the number of undergraduates who applied
14	sat_avg	Average SAT score (or equivalent) for admitted students
15	ugds	Number of number of degree-seeking undergraduate students enrolled in the fall term

college.csv (continued)		
Column	Name	Value
16	ugds_men	Proportion of <i>ugds</i> who are men
17	ugds_women	Proportion of <i>ugds</i> who are women
18	ugds_white	Proportion of <i>ugds</i> who are white (based on self-reports)
19	ugds_black	Proportion of <i>ugds</i> who are black/African American (based on self-reports)
20	ugds_hisp	Proportion of <i>ugds</i> who are Hispanic (based on self-reports)
21	ugds_asian	Proportion of <i>ugds</i> who are Asian (based on self-reports)
22	ugds_other	Proportion of <i>ugds</i> who have other race/ethnicity (created from other categories on original data file; race/ethnicity proportions sum to 1)
23	npt4	Average net price of attendance, derived from the full cost of attendance (including tuition and fees, books and supplies, and living expenses) minus federal, state, and institutional grant/scholarship aid, for full-time, first-time undergraduate Title IV-receiving students. NPT4 created from scorecard data variables NPT4_PUB if public institution and NPT4_PRIV if private
24	tuitionfee_in	In-state tuition and fees
25	tuitionfee_out	Out-of-state tuition and fees
26	avgfacsal	Average faculty salary per month
27	pftfac	Proportion of faculty that is full-time
28	c150_4	Proportion of first-year, full-time students who complete their degree within 150% of the expected time to complete; for most institutions, this is the proportion of students who receive a degree within 6 years
29	grads	Number of graduate students

collegerg.csv Five replicate SRSs from the set of public colleges and universities (having *control* = 1) in **college.csv**. Columns 1–29 are as in **college.csv**, with additional columns 30–32 listed below. Note that the selection probabilities and sampling weights are for the separate replicate samples, so that the weights for each replicate sample sum to the population size 500.

Column	Name	Value
30	selectionprob	Selection probability for each replicate sample
31	samplingweight	Sampling weight for each replicate sample
32	repgroup	Replicate group number

collshr.csv Probability-proportional-to-size sample of size 10 from the stratum of small, highly residential colleges (having *ccsiseset* = 11) in **college.csv**. Columns 1–29 are as in **college.csv**, with additional columns 30–34 listed below.

Column	Name	Value
30	mathfac	Number of mathematics faculty
31	psychfac	Number of psychology faculty
32	biolfac	Number of biology faculty
33	psii	Selection probability, = $ugds / (\text{sum of } ugds \text{ for stratum})$
34	wt	Sampling weight = $1 / (10\psi_i)$

coots.csv Selected information on egg size, from a larger study by Arnold (1991). Data provided courtesy of Todd Arnold. Not all observations are used for this data set, so results may not agree with those in Arnold (1991).

Column	Name	Value
1	clutch	clutch number from which eggs were subsampled.
2	csize	number of eggs in clutch (M_i)
3	length	length of egg (mm)
4	breadth	maximum breadth of egg (mm)
5	volume	calculated as $0.000507 * \text{length} * \text{breadth}^2$ (mm ³)
6	tmt	= 1 if received supplemental feeding, 0 otherwise

counties.csv Data (from 1990) from an SRS of 100 of the 3141 counties in the United States. Missing values are coded as −99. Source: U.S. Census Bureau (1994).

Column	Name	Value
1	RN	random number used to select the county
2	state	state abbreviation
3	county	county name
4	landarea	land area, 1990 (square miles)
5	totpop	total number of persons, 1992
6	physician	active non-Federal physicians on Jan. 1, 1990
7	enroll	school enrollment in elementary or high school, 1990
8	percpub	percent of school enrollment in public schools
9	civlabor	civilian labor force, 1991
10	unemp	number unemployed, 1991
11	farmpop	farm population, 1990
12	numfarm	number of farms, 1987
13	farmacre	acreage in farms, 1987
14	fedgrant	total expenditures in federal funds and grants, 1992 (millions of dollars)
15	fedciv	civilians employed by federal government, 1990
16	milit	military personnel, 1990
17	veterans	number of veterans, 1990
18	percviet	percent of veterans from Vietnam era, 1990

crimes.csv Data from selected variables in a simple random sample of 5,000 records from the 7,048,107 records with dates between 2001 and 2019 in the City of Chicago database “Crimes—2001 to Present.” This file was downloaded on August 11, 2020 from <https://data.cityofchicago.org/>. These data are provided for pedagogical purposes only. Anyone

wishing to publish analyses of Chicago crime data should obtain the most recent data from <https://data.cityofchicago.org/>. For a list and map of Community Areas, see https://www.chicago.gov/city/en/depts/dgs/supp_info/citywide_maps.html.

Column	Name	Value
1	year	Year in which crime occurred (between 2001 and 2019)
2	crimetype	Type of crime, determined from detailed crime description in database homicide = homicide, sexualasslt = sexual assault, robbery = robbery, aggasslt = aggravated assault, burglary = burglary, mvtheft = motor vehicle theft, idtheft = identity theft, theft = other type of theft, arson = arson, simpleasslt = simple assault (assaults that are not aggravated), threat = threat or harassment, fraud = fraud, weapon = weapons violation, trespass = trespassing, narcotics = narcotics or liquor law violation, vandalism = vandalism, other = other
3	violent	= 1 if violent crime, 0 otherwise
4	arrest	= 1 if an arrest was made, 0 otherwise
5	domestic	= 1 if crime was domestic-related as defined by the Illinois Domestic Violence Act, 0 otherwise
6	commarea	Number of the Community Area in Chicago where the crime occurred
7	location	Type of location where crime occurred (e.g., street, apartment)

deadtrees.csv Number of dead trees recorded by photograph and field count for a (fictional) SRS of 25 plots taken from a population of 100 plots.

Column	Name	Value
1	photo	Number of dead trees in plot from photograph
2	field	Number of dead trees in plot from field observation

divorce.csv Data from a sample of divorce records for states in the Divorce Registration Area. Source: National Center for Health Statistics (1987).

Column	Name	Value
1	state	state name (character variable)
2	abbrev	state abbreviation (character variable)
3	samprate	sampling rate for state
4	numrecs	number of records sampled in state
5	hsblt20	number of records in sample with husband's age < 20
6	hsb20to24	number of records with $20 \leq$ husband's age ≤ 24
7	hsb25to29	number of records with $25 \leq$ husband's age ≤ 29
8	hsb30to34	number of records with $30 \leq$ husband's age ≤ 34
9	hsb35to39	number of records with $35 \leq$ husband's age ≤ 39
10	hsb40to44	number of records with $40 \leq$ husband's age ≤ 44

divorce.csv (continued)		
Column	Name	Value
11	hsb45to49	number of records with $45 \leq \text{husband's age} \leq 49$
12	hsbge50	number of records with husband's age ≥ 50
13	wflt20	number of records with wife's age < 20
14	wf20to24	number of records with $20 \leq \text{wife's age} \leq 24$
15	wf25to29	number of records with $25 \leq \text{wife's age} \leq 29$
16	wf30to34	number of records with $30 \leq \text{wife's age} \leq 34$
17	wf35to39	number of records with $35 \leq \text{wife's age} \leq 39$
18	wf40to44	number of records with $40 \leq \text{wife's age} \leq 44$
19	wf45to49	number of records with $45 \leq \text{wife's age} \leq 49$
20	wfge50	number of records with wife's age ≥ 50

gini.csv Data from the population of districts for the 1921 Italian general census. Source: Gini and Galvani (1929, pp. 73–78).

Column	Name	Value
1	id	ID number
2	district	District name
3	birth_rate	Births per 1,000 population
4	death_rate	Deaths per 1,000 population
5	marriage_rate	Marriages per 1,000 population
6	agricultural_pop	Percentage of males over 10 years old who work in agriculture
7	urban_population	Percentage of population in urban areas
8	income	Average income
9	altitude	Average altitude above sea level (meters)
10	pop_density	Number of inhabitants per square kilometer
11	natural_growth	Rate of average increase of the population
12	population	Population of area
13	area	Land area (square kilometers)
14	in_GG_sample	= 1 if in the purposive sample selected by Gini and Galvani; 0 otherwise

golfsrs.csv A simple random sample of 120 golf courses, taken from the population on the website ww2.golfcourse.com on August 5, 1998. Missing data in the .csv file are denoted by blanks.

Column	Name	Value
1	RN	random number used to select golf course for sample
2	state	state name
3	holes	number of holes
4	type	type of course: priv = private, semi = semi-private, pub = public, mili = military, resort
5	yearblt	year course was built
6	wkday18	greens fee for 18 holes during week
7	wkday9	greens fee for 9 holes during week
8	wkend18	greens fee for 18 holes on weekend

golfsrs.csv (continued)		
Column	Name	Value
9	wkend9	greens fee for 9 holes on weekend
10	backtee	back tee yardage
11	rating	course rating
12	par	par for course
13	cart18	golf cart rental fee for 18 holes
14	cart9	golf cart rental fee for 9 holes
15	caddy	Are caddies available? (y or n)
16	pro	Is a golf pro available? (y or n)

gpa.csv GPA data from Chapter 5 of SDA.

Column	Name	Value
1	suite	Suite (psu) identifier
2	gpa	Grade point average of person in suite
3	wt	Sampling weight, = 20 for every observation

healthjournals.csv Randomization and statistical inference practices in a stratified random sample of 196 public health articles. The data, provided courtesy of Dr. Matt Hayat, are discussed in Hayat and Knapp (2017). The variables provided in **healthjournals.csv** are a subset of the variables collected by the authors.

Column	Name	Value
1	journal	Journal that published the article AJPH = <i>American Journal of Public Health</i> AJPM = <i>American Journal of Preventive Medicine</i> PM = <i>Preventive Medicine</i>
2	NumAuthors	Number of authors
3	RandomSel	= “Yes” if data in the article were from a randomly selected (probability) sample; “No” otherwise
4	RandomAssn	= “Yes” if study subjects for the article were randomly assigned to treatment groups; “No” otherwise
5	ConfInt	= “Yes” if a confidence interval appeared in the article’s main text, tables, or figures; “No” otherwise
6	HypTest	= “Yes” if a p-value or significance test appeared in the article’s main text, tables, or figures; “No” otherwise
7	Asterisks	= “Yes” if asterisks were used to represent p-value ranges; “No” otherwise

htcdf.csv Empirical distribution function and empirical probability mass function of data in **htpop.csv**.

Column	Name	Value
1	height	height value, cm
2	frequency	number of times height value in column 1 occurs in population
3	epmf	empirical probability mass function
4	ecdf	empirical distribution function

htpop.csv Height and gender of 2,000 persons in an artificial population.

Column	Name	Value
1	height	height of person, cm
2	gender	M=male, F=female

htsrs.csv Height and gender for a SRS of 200 persons, taken from **htpop.csv**.

Column	Name	Value
1	rn	random number used to select unit
2	height	height of person, cm
3	gender	M=male, F=female

htstrat.csv Height and gender for a stratified random sample of 160 women and 40 men, taken from **htpop.csv**. The columns and names are as in **htsrs.csv**.

hunting.csv Population and sample sizes for the poststrata used for the Sunday hunting survey. Source: Virginia Polytechnic and State University/Responsive Management (2006).

Column	Name	Value
1	region	Region of state (East, Central, West)
2	gender	Gender (female, male)
3	age	Age group (16-24, 25-34, 35-44, 45-54, 55-64, 65+)
4	popsize	Population size in poststratum from the 2000 U.S. census
5	sampsize	Sample size in poststratum

impute.csv Small artificial data set used to illustrate imputation methods. Missing values are denoted by -99.

Column	Name	Value
1	person	identification number for person
2	age	age in years
3	gender	M=male, F=female
4	education	number of years of education
5	crime	= 1 if victim of any crime, 0 otherwise
6	violcrime	= 1 if victim of violent crime, 0 otherwise

integerwt.csv Artificial population of 2,000 observations.

Column	Name	Value
1	stratum	Stratum number
2	y	y value of observation

intellonline.csv Data from the online (Mechanical Turk) survey. Source: Heck et al. (2018). The data were downloaded from <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0200103> on February 8, 2020; the variables extracted from the full data set are provided here for educational purposes only.

Column	Name	Value
1	int	Response to question about agreement with the statement “I am more intelligent than the average person.” 1 = Strongly Agree; 2 = Mostly Agree; 3 = Mostly Disagree; 4 = Strongly Disagree; 5 = Don’t Know or Not Sure
2	region	Census region of respondent (character variable, length 10): Northeast, South, Midwest, West
3	sex	Sex (character variable, length 8): Male, Female
4	race	Race (character variable, length 18): White, African American, Asian American, Hispanic American, Another origin
5	age	Age, years
6	income	Household income level (character variable, length 8): <\$40k, \$40–80k, or >\$80k
7	education	Highest education level attained (character variable, length 12): No College, Some College, College Grad, Grad School
8	postwt	Relative weight, obtained by poststratifying to demographic proportions in the 2010 U.S. Census. The weights are normed so that they sum to 750.

intelltel.csv Data from the telephone survey studied by Heck et al. (2018). The data were downloaded from <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0200103> and are provided here for educational purposes only. The variables are the same as in **intellonline.csv**.

intellwts.csv Relative weights for demographic groups in **intellonline.csv** and **intelltel.csv** (Heck et al., 2018). Each sample was weighted using the 2010 U.S. Census demographics for sex (male, female), age (< 44, ≥ 44), and race/ethnicity (white, nonwhite). The table entries give the weights for each of these eight demographic groups.

Column	Name	Value
1	sex	Sex
2	agegroup	Age group: Young = (age less than 44), Old = (age greater than or equal to 44)
3	race	Race: White or Nonwhite
4	tel_n	Number of telephone survey respondents in the sex/age-group/race class

intellwts.csv (continued)		
Column	Name	Value
5	online_n	Number of online survey respondents in the sex/agegroup/race class
6	tel_wgt	Relative weight for each respondent to the telephone survey in this sex/agegroup/race class
7	online_wgt	Relative weight for each respondent to the telephone survey in this sex/agegroup/race class

ipums.csv Data extracted from the 1980 Census Integrated Public Use Microdata Series, using the “Small Sample Density” option in the data extract tool, on September 17, 2008. The stratum and psu variables were constructed for use in the book exercises. Data analyses on this file do NOT give valid results for inference to the 1980 U.S. population. Source: Ruggles et al. (2004).

Column	Name	Value
1	stratum	stratum number (1–9)
2	psu	psu number (1–90)
3	inctot	total personal income (dollars), topcoded at \$75,000
4	age	age, with range 15–90
5	sex	1 = Male, 2 = Female
6	race	1 = White, 2 = Black, 3 = American Indian or Alaska Native, 4 = Asian or Pacific Islander, 5 = Other Race
7	hispanic	0 = Not Hispanic, 1 = Hispanic
8	marstat	Marital Status: 1 = Married, 2 = Separated, 3 = Divorced, 4 = Widowed, 5 = Never married/single
9	ownershg	Ownership of housing unit: 0 = Not Applicable (N/A), 1 = Owned or being bought, 2 = Rents
10	yrsusa	Number of years a foreign-born person has lived in the U.S.: 0= N/A, 1= 0–5 years, 2= 6–10 years, 3= 11–15 years, 4= 16–20 years, 5= 21+ years
11	school	Is person in school? 0 = N/A, 1 = No, not in school, 2 = Yes, in school
12	educrec	Educational Attainment: 1= None or preschool, 2= Grade 1, 2, 3, or 4, 3= Grade 5, 6, 7, or 8, 4= Grade 9, 5= Grade 10, 6= Grade 11, 7= Grade 12, 8= 1 to 3 years of college, 9= 4+ years of college
13	labforce	In labor force? 0 = Not Applicable, 1 = No, 2 = Yes
14	classwk	class of worker: 0=Not applicable, 13= Self-employed, not incorporated, 14= Self-employed, incorporated, 22= Wage/salary, private, 25= Federal government employee, 27= State government employee, 28= Local government employee, 29= Unpaid family worker
15	vetstat	Veteran Status 0 = Not Applicable, 1 = No Service, 2 = Yes

journal.csv Types of sampling used for articles in a sample of journals. Source: Jacoby and Handlin (1991).

Note that columns 2 and 3 do not always sum to column 1; for some articles, the investigators could not determine which type of sampling was used. When working with these data, you may wish to create a fourth column, “indeterminate,” which equals $\text{column1} - (\text{column2} + \text{column3})$.

Column	Name	Value
1	numemp	number of articles in 1988 that used sampling
2	prob	number of articles that used probability sampling
3	nonprob	number of articles that used non-probability sampling

measles.csv Roberts et al. (1995) reported on the results of a survey of parents whose children had not been immunized against measles during a recent campaign to immunize all children in the first five years of secondary school. The original data were unavailable; univariate and multivariate summary statistics from these artificial data, however, are consistent with those in the paper. All variables are coded as 1 for yes, 0 for no, and 9 for no answer. A parent who refused consent (variable 4) was asked why, with responses in variables 5 through 10. If a response in variables 5 through 10 was checked, it was assigned value 1; otherwise, it was assigned value 0. A parent could give more than one reason for not having the child immunized.

Column	Name	Value
1	school	school attended by child
2	form	Parent received consent form
3	returnf	Parent returned consent form
4	consent	Parent gave consent for measles immunization
5	hadmeas	Child had already had measles
6	previmm	Child had been immunized against measles
7	sideeff	Parent concerned about side effects
8	gp	Parent wanted general practitioner (GP) to give vaccine
9	noshot	Child did not want injection
10	notser	Parent thought measles was not a serious illness
11	gpadv	GP advised that vaccine was not needed
12	Mitotal	Population size in school
13	mi	Sample size in school

mysteries.csv Data from a stratified random sample of books nominated for the Edgar[®] awards for Best Novel and Best First Novel. The sample was drawn from the population listing of 655 books at <http://theedgars.com/awards/> on August 14, 2020.

Column	Name	Value
1	stratum	Stratum number, from 1 to 12, computed from the stratification variables in columns 2–4
2	time	Time period in which award was given: 1 = 1946–1980, 2 = 1981–2000, 3 = 2001–2020
3	category	Award category (character variable, length 16): Best Novel, or Best First Novel

mysteries.csv (continued)		
Column	Name	Value
4	winner	= 1 if book won the award that year, = 0 if book was nominated but did not win award
5	popsize	Number of population books in stratum ($= N_h$)
6	sampsize	Number of sampled books in stratum ($= n_h$)
7	obtained	= 1 if book was obtained (responded) in original sample, = 2 if book was obtained in phase II subsample of nonrespondents, = 0 if not obtained
8	p1weight	Weight for phase I sample, calculated as N_h/n_h ; use for exercises in Chapters 1–11 of SDA
9	p2weight	Final weight for phase II sample; use for exercises in Chapter 12 of SDA and analyses involving variables <i>victims</i> and <i>firearm</i>
10	genre	Genre of book (character variable, length 11). Values “private eye” (protagonist is a private detective), “procedural” (a detailed, step-by-step analysis of how the crime is solved, using the skills of the detective), or “suspense” (the protagonist is at the center of action or is involved in espionage, but is not a professional detective)
11	historical	= 1 if the main action in the book takes place at least 20 years before the book’s publication date, = 0 if book action is within 20 years of the publication date
12	urban	= 1 if the main action in the book takes place primarily in urban areas, = 0 otherwise
13	authorgender	Gender of author (character variable, length 1) = “F” if author is female, “M” if author is male
14	fdetect	Number of female detectives (or protagonists, if book has no detective) in book
15	mdetect	Number of male detectives (or protagonists, if book has no detective) in book
16	victims	Number of murder victims in book (missing value set to -9 if <i>obtained</i> = 0)
17	firearm	Number of murders committed with firearms in book (missing value set to -9 if <i>obtained</i> = 0)

nhanes.csv Selected variables from the 2015–2016 National Health and Nutrition Examination Survey (NHANES). Source: Centers for Disease Control and Prevention (2017). This data set is provided for educational purposes only. Anyone wishing to publish or use results from analyses of NHANES data should obtain the data files directly from the source.

The data files merged to create **nhanes.csv** can be read directly from the SAS transport files DEMO_I.XPT, BMX_I.XPT, TCHOL_I.XPT, and BPX_I.XPT from the NHANES website. Variables 1–23 have the same names as in the SAS transport files.

The blood pressure variables *sbp* and *dbp* were created as follows. In the medical examination, three consecutive blood pressure readings were obtained after participants sat quietly for 5 minutes, and the maximum inflation level was determined. A fourth measurement was conducted for some persons who had an incomplete or interrupted blood pressure reading.

The variables *sbp* and *dbp* were calculated by discarding the first blood pressure reading and calculating the average of the remaining valid readings. Note that some of the diastolic blood pressure readings are 0.

In the comma-delimited file `nhanes.csv`, missing values are denoted by `-9`. In the SAS data file, missing values are denoted by a period. In the R data file, missing values are denoted by `NA`. Note that some of the codes for variables in the table below also denote missing values; for example, the value 7 for *dmdeduc2* indicates “Refused,” and these codes for special types of missing values remain in the SAS and R data files.

Column	Name	Value
1	<i>sdmvstra</i>	Pseudo-stratum. These are groups of secondary sampling units used for variance estimation on the publicly available data. Pseudo-strata and pseudo-psus are released instead of the actual strata and psus to protect the confidentiality of respondents’ information. Use <i>sdmvstra</i> as the variable defining the strata.
2	<i>sdmvpsu</i>	Pseudo-psu. Use <i>sdmvpsu</i> as the primary sampling unit (psu). There are two pseudo-psus per pseudo-stratum, numbered 1 and 2.
3	<i>wtint2yr</i>	Interview weight (use as weight for variables 5–12)
4	<i>wtmec2yr</i>	Mobile Examination Center weight (use as weight for any analysis involving variables 13–25)
5	<i>ridstatr</i>	Interview/examination status, = 1 if interviewed only, = 2 if interviewed and had medical examination
6	<i>ridageyr</i>	Age in years at screening, from 0 to 80. Anyone with age > 80 years is recorded (topcoded) as 80. No values are missing for this variable.
7	<i>ridagemn</i>	Age in months at screening (reported only for persons with age 24 months or younger at the time of exam, otherwise missing)
8	<i>riagendr</i>	= 1 if male, 2 if female (no missing values)
9	<i>ridreth3</i>	Race/ethnicity code (no missing values) 1 = Mexican American 2 = Other Hispanic 3 = Non-Hispanic White 4 = Non-Hispanic Black 6 = Non-Hispanic Asian 7 = Other Race, Including Multi-Racial
10	<i>dmdeduc2</i>	Education level of person interviewed (given for adults age 20+ only) 1 = Less than 9th grade 2 = 9th to 11th grade (including 12th grade with no diploma) 3 = High school graduate (including GED) 4 = Some college or associate’s degree 5 = College graduate or above 7 = Refused 9 = Don’t know
11	<i>dmdfmsiz</i>	Total number of people in the family. Values 1–6 indicate the number of people is that number; value 7 indicates 7 or more people in family. No missing values.

nhanes.csv (continued)		
Column	Name	Value
12	indfmpir	Ratio of family income to poverty guideline. A value less than 1 indicates the family is below the poverty threshold. Variable <i>indfmpir</i> is a continuous variable where values between 0 and 4.99 indicate the actual poverty ratio. A value of 5 indicates that the ratio of family income to the poverty guideline for that family is 5 or more.
13	bmxtwt	Weight (kg)
14	bmxtht	Standing height (cm)
15	bmxbmi	Body mass index (kg/m ²), calculated as $\text{bmxtwt}/(\text{bmxtht}/100)^2$
16	bmxtwaist	Waist circumference (cm)
17	bmxtleg	Upper leg length (cm)
18	bmxtarm1	Upper arm length (cm)
19	bmxtarmc	Upper arm circumference (cm)
20	bmdavasad	Average sagittal abdominal diameter (SAD, the distance from the small of the back to the upper abdomen), in cm. Calculated by averaging the SAD readings on the person (up to four).
21	lbxtc	Serum total cholesterol (mg/dL)
22	bpxpls	60-second pulse
23	sbp	Average systolic blood pressure (mm Hg)
24	dbp	Average diastolic blood pressure (mm Hg)
25	bpread	Number of blood pressure readings

nybight.csv Data collected in the New York Bight for June 1974 and June 1975. Two of the original strata were combined because of insufficient sample sizes. For variable *catchwt*, weights less than 0.5 were recorded as 0.5 kg. Source: Wilk et al. (1977).

Column	Name	Value
1	year	year of data collection, 1974 or 1975
2	stratum	stratum membership, based on depth
3	catchnum	number of fish caught during trawl
4	catchwt	total weight (kg) of fish caught during trawl
5	numsp	number of species of fish caught during trawl
6	depth	depth of station (m)
7	temp	surface temperature (degrees C), missing = -99

otters.csv Data on number of holts (dens) in Shetland, U.K., used in Kruuk et al. (1989). Data courtesy of Hans Kruuk.

Column	Name	Value
1	section	section of coastline
2	habitat	type of habitat (stratum)
3	holts	number of holts (dens)

ozone.csv Hourly ozone readings (parts per billion, ppb) from a site in Monterey County, California, for 2018 and 2019. Source: https://aqs.epa.gov/aqsweb/airdata/download_files.html#Raw, accessed November 19, 2020. Missing values are denoted by -9.

Column	Name	Value
1	year	year of reading (2018 or 2019)
2	month	month of reading (1–12)
3	day	day of reading (1–31)
4	hr0	ozone reading (ppb) at 0:00 local time
5	hr1	ozone reading (ppb) at 1:00 local time
⋮	⋮	⋮
27	hr23	ozone reading (ppb) at 23:00 local time

pitcount.csv Fictional data from a fictional point-in-time (PIT) survey taken to estimate the number of persons experiencing homelessness.

Column	Name	Value
1	strat	Stratum number (from 1 to 8)
2	division	Geographic division, used to form strata
3	density	Expected density of persons experiencing homelessness (character variable, with values High or Low)
4	popsize	$= N_h$, the number of areas in the population for stratum h
5	sampsize	$= n_h$, the number of areas in the sample for stratum h
6	areawt	$= N_h/n_h$, the sampling weight for the area
7	y	Number of persons experiencing unsheltered homelessness found in the area during the PIT count

profresp.csv The data described in Zhang et al. (2020) were downloaded from <http://doi.org/10.3886/E109021V1> on January 22, 2020, from file **survey4.rds**. The data set **profresp.csv** contains selected variables from the set of 2,407 respondents who completed the survey and provided information on the demographic variables and the information needed to calculate “professional respondent” status. The full data set **survey4.rds** contains numerous additional questions about behavior that are not included here, as well as the data from the partially completed surveys. The website also contains data for three other online panel surveys. Because **profresp.csv** is a subset of the full data, statistics calculated from it may differ from those in Zhang et al. (2020).

Missing values are denoted by –9.

Column	Name	Value
1	prof_cat	Level of professionalism 1 = novice, 2 = average, 3 = professional
2	panelnum	Number of panels respondent has belonged to. A response between 1 and 6 means that the person has belonged to that number of panels; 7 means 7 or more.
3	survnum_cat	How many Internet surveys have you completed before this one? 1 = This is my first one, 2 = 1–5, 3 = 6–10, 4 = 11–15, 5 = 16–20, 6 = 21–30, 7 = More than 30
4	panelq1	Are you a member of any online survey panels besides this one? 1 = yes, 2 = no
5	panelq2	To how many other online panels do you belong?

profresp.csv (continued)		
Column	Name	Value
		1 = None, 2 = 1 other panel, 3 = 2 others, 4 = 3 others, 5 = 4 others, 6 = 5 others, 7 = 6 others or more. This question has a missing value if <i>panelq1</i> = 2. If you want to estimate how many panels a respondent belongs to, create a new variable <i>numpanel</i> that equals <i>panelq2</i> if <i>panelq2</i> is not missing and equals 1 if <i>panelq1</i> = 2.
6	age4cat	Age category. 1 = 18 to 34, 2 = 35 to 49, 3 = 50 to 64, 4 = 65 and over
7	edu3cat	Education category. 1 = high school or less, 2 = some college or associates' degree, 3 = college graduate or higher
8	gender	Gender: 1 = male, 2 = female
9	non_white	1 = race is non-white, 0 = race is white
10	motive	Which best describes your main reason for joining on-line survey panels? 1 = I want my voice to be heard, 2 = Completing surveys is fun, 3 = To earn money, 4 = Other (Please specify)
11	freq_q1	During the PAST 12 MONTHS, how many times have you seen a doctor or other health care professional about your own health? Response is number between 0 and 999.
12	freq_q2	During the PAST MONTH, how many days have you felt you did not get enough rest or sleep?
13	freq_q3	During the PAST MONTH, how many times have you eaten in restaurants? Please include both full-service and fast food restaurants.
14	freq_q4	During the PAST MONTH, how many times have you shopped in a grocery store? If you shopped at more than one grocery store on a single trip, please count them separately.
15	freq_q5	During the PAST 2 YEARS, how many overnight trips have you taken?

profrespacs.csv Population estimates from the 2011 American Community Survey (ACS) for age/gender/education categories measured in **profresp.csv** (Zhang et al., 2020). Note that *age3cat* has 3 categories, while the age variable in **profresp.csv** has 4 categories.

Column	Name	Value
1	gender	Gender: 1 = male, 2 = female
2	age3cat	Age category. 1 = 18 to 34, 2 = 35 to 64, 3 = 65 and over
3	edu3cat	Education category. 1 = high school or less, 2 = some college or associates' degree, 3 = college graduate or higher
4	count	Population size from ACS for the gender/age/education level combination

radon.csv Radon readings for a stratified sample of 1003 homes in Minnesota. Source: Nolan and Speed (2000). The data were downloaded in April 2008 from an earlier version of the website now located at www.stat.berkeley.edu/users/statlabs/labs.html.

Column	Name	Value
1	countyname	County Name
2	countynum	County Number
3	sampsize	Sample size in county
4	popsize	Population size in county
5	radon	Radon concentration (picocuries per liter)

rectlength.csv Lengths of rectangles.

Column	Name	Value
1	rectangle	Rectangle number
2	length	Rectangle length

rnt.csv Page from a random number table. Open the .csv file in a text editor instead of a spreadsheet, because a spreadsheet strips off the leading zeroes. The columns have format z5.0 in the SAS file, and are character variables in the R file, so that leading zeroes are displayed in those formats.

Column	Name	Value
1	col1	Column of 5-digit random numbers
2	col2	Column of 5-digit random numbers
3	col3	Column of 5-digit random numbers
4	col4	Column of 5-digit random numbers
5	col5	Column of 5-digit random numbers
6	col6	Column of 5-digit random numbers

sample70.csv All possible simple random samples that can be generated from the population in Example 2.2 of SDA.

Column	Name	Value
1	sampnum	Sample number
2–5	u1–u4	Sampled units in $\mathcal{S}$
6–9	y1–y4	Values of y_i in sample $\mathcal{S}$
10	total	Estimated population total

santacruz.csv The number of seedlings in the sampled psus on Santa Cruz Island, California, in 1992 and 1994. Source: Peart (1994).

Column	Name	Value
1	tree	Tree number
2	seed92	Number of seedlings in 1992
3	seed94	Number of seedlings in 1994

schools.csv Math and reading test results from a two-stage cluster sample of tenth-grade students. An SRS of 10 schools was selected from the 75 schools in the population, and then 20 students were sampled from each school. These data are fictional but the summary statistics are consistent with those seen in educational studies.

Column	Name	Value
1	schoolid	School number (use as cluster variable)
2	gender	Gender of student (character variable, F = female, M = male)
3	math	Score on math test
4	reading	Score on reading test
5	mathlevel	Category level for math test score: 1 if $1 \leq \text{math} \leq 40$ 2 if $41 \leq \text{math}$
6	readlevel	Category level for reading test score: 1 if $1 \leq \text{read} \leq 32$ 2 if $33 \leq \text{read} \leq 50$
7	Mi	Number of students in school, M_i
8	finalwt	Weight for student in sample

seals.csv Data on number of breathing holes found in sampled areas of Svalbard fjords, reconstructed from summary statistics given in Lydersen and Ryg (1991).

Column	Name	Value
1	zone	zone number for sampled area
2	holes	number of breathing holes Imjak found in area

shapespop.csv Population of black and gray squares and circles.

Column	Name	Value
1	ID	identification number for object
2	shape	shape of object (square or circle)
3	color	color of object (gray or black)
4	area	area of object (cm^2)
5	conv	= 1 if object can be reached through convenience sample, 0 otherwise

shorebirds.csv Two-phase sample of shorebird nests. These are artificial data constructed from summary statistics given in Bart and Earnst (2002).

Column	Name	Value
1	plot	Plot number
2	rapid	Rapid-method count of number of birds in plot
3	intense	Intensive-method count of number of nests in plot = -9 if the plot is not in the phase II sample

sp500.csv Companies in the S&P 500[®] Stock Market Index as of September 15, 2020. Source: Downloaded from <https://fknol.com/list/eps-sp-500-index-companies.php> on September 19, 2020.

Column	Name	Value
1	Company	Company name (character variable, length 37)
2	Symbol	Stock symbol (character variable, length 5)
3	MarketCap	Market capitalization, in billions of U.S. dollars
4	StockPrice	Price per share of stock
5	PE_Ratio	Price-to-earnings ratio
6	EPS	Earnings per share

spanish.csv Fictional cluster sample of introductory Spanish students.

Column	Name	Value
1	class	Class number
2	score	Score on vocabulary test (out of 100)
3	trip	= 1 if plan a trip to a Spanish-speaking country, 0 otherwise

srs30.csv An SRS of size 30 taken from an artificial population of size 100.

Column	Name	Value
1	y	Value of observation

ssc.csv SRS of 150 members of the Statistical Society of Canada, downloaded from `ssc.ca` in August, 2006.

Column	Name	Value
1	gender	m = male, f=female
2	occupation	a = academic, g = government, i = industry, n = not determined
3	ASA	= 1 if person is member of American Statistical Association, 0 otherwise

statepop.csv Data from an unequal-probability sample of 100 counties from the 1994 *County and City Data Book* (U.S. Census Bureau, 1994). The sample was selected with probability proportional to population.

Column	Name	Value
1	county	county name (character variable, length 14)
2	state	state name (character variable)
3	landarea	land area of county, 1990 (square miles)
4	popn	population of county, 1992
5	phys	number of physicians, 1990
6	farmpop	farm population, 1990
7	numfarm	number of farms, 1987
8	farmacre	number of acres devoted to farming, 1987
9	veterans	number of veterans, 1990
10	percviet	percent of veterans from Vietnam era, 1990
11	psii	ψ_i , probability of selection
12	wt	sampling weight, $= 1/(100\psi_i)$

statepps.csv Number of counties (or county equivalents; Alaska has boroughs, Louisiana has parishes, and some states have independent cities), population estimates for 2019, land area, and water area for the 50 states plus the District of Columbia. Total area for a state can be calculated by summing land area and water area.

Source: Population estimates are from U.S. Census Bureau (2019). Land and water areas are from U.S. Census Bureau (2012).

Column	Name	Value
1	state	state name (character variable, length 20)
2	counties	number of counties or county equivalents
3	pop2019	population of state, 2019
4	landarea	land area of state (square kilometers)
5	waterarea	water area of state (square kilometers)

swedishlcs.csv Data on call attempts from the Swedish Survey of Living Conditions. Source: Lundquist and Särndal (2013).

Column	Name	Value
1	attempt	call attempt number
2	resprate	response rate at call attempt (percent)
3	benefits	relative bias for variable <i>benefits</i>
4	income	relative bias for variable <i>income</i>
5	employed	relative bias for variable <i>employed</i>
6	note	Character variable, length 25: notes about data collection

The variable *attempt* takes on values 1–25 for the initial fieldwork period. Values 31–40 denote the follow-up period, and value 45 gives the final estimates. The gaps in the attempt variable allow one to see the separation of the periods on the graph.

syc.csv Selected variables from the Survey of Youth in Custody (Beck et al., 1988). Source: U.S. Department of Justice (1989). Strata 6–16 each contain one facility; the psus in those strata are residents. In strata 1–5, the psus are facilities. The number of facilities in the population (N_h) for those five strata are: $N_1 = 99$, $N_2 = 39$, $N_3 = 30$, $N_4 = 13$, $N_5 = 14$. Eleven facilities are sampled from stratum 1, and seven facilities are sampled from each of strata 2–5.

The table gives missing value codes for individual variables in the .csv file (these codes are the same as in the original data source, but have been changed to the appropriate missing value codes for the respective software packages in the SAS and R data files).

Column	Name	Value
1	stratum	stratum number
2	psu	psu number, = facility number for residents in strata 1–5 and person number for residents in strata 6–16
3	facility	facility number
4	facsize	number of eligible residents in psu
5	finalwt	final weight
6	randgrp	random group number
7	age	age of resident (99=missing)

syc.csv (continued)		
Column	Name	Value
8	race	race of resident 1 = white; 2 = Black; 3 = Asian/Pacific Islander; 4 = American Indian, Alaska Native; 5 = Other; 9 = Missing
9	ethnicity	1 = Hispanic, 2 = not Hispanic, 9=missing
10	educ	highest grade attended before sent to correctional institution 0 = Never attended school; 1–12 = highest grade attended; 13 = GED; 14 = Other
11	gender	1 = male, 2 = female
12	livewith	Who did you live with most of the time you were growing up? 1 = Mother only, 2 = Father only 3 = Both mother and father, 4 = Grandparents, 5 = Other relatives, 6 = Friends, 7 = Foster home, 8 = Agency or institution, 9 = Someone else, 99 = Blank
13	famtime	Has anyone in your family, such as your mother, father, brother, sister, ever served time in jail or prison? 1 = Yes, 2 = No, 7 = Don't know, 9 = Blank
14	crimtype	most serious crime in current offense 1 = violent (e.g., murder, rape, robbery, assault) 2 = property (e.g., burglary, larceny, arson, fraud, motor vehicle theft) 3 = drug (drug possession or trafficking) 4 = public order (weapons violation, perjury, failure to appear in court) 5 = juvenile status offense (truancy, running away, incorrigible behavior) 9 = missing
15	everviol	ever put on probation or sent to correctional inst for violent offense: 1 = yes, 0 = no
16	numarr	number of times arrested (99=missing)
17	probtn	number of times on probation (99=missing)
18	corrinst	number of times previously committed to correctional institution (99=missing)
19	vertime	Prior to being sent here did you ever serve time in a correctional institution? 1 = yes, 2 = no, 9 = missing
20	prviol	=1 if previously arrested for violent offense, 0 otherwise
21	prprop	=1 if previously arrested for property offense, 0 otherwise
22	prdrug	=1 if previously arrested for drug offense, 0 otherwise
23	prpub	=1 if previously arrested for public order offense, 0 otherwise
24	prjuv	=1 if previously arrested for juvenile status offense, 0 otherwise
25	agefirst	age first arrested (99=missing)
26	usewepn	Did you use a weapon ... for this incident? 1 = Yes, 2 = No, 9 = Blank
27	alcuse	Did you drink alcohol at all during the year before being sent here this time? 1 = Yes; 2 = No, didn't drink during year before; 3 = No, don't drink at all, 9 = missing
28	everdrug	Ever used illegal drugs; 0=no, 1=yes, 9=missing

teachers.csv Selected variables from a study on elementary school teacher workload in Maricopa County, Arizona. Data courtesy of Rita Gnap (Gnap, 1995). The psu sizes are given in file **teachmi.csv**. The large stratum had 245 schools; the small/medium stratum had 66 schools. Missing values are coded as -9 .

Column	Name	Value
1	dist	school district size. Character variable: large or med/small
2	school	school identifier
3	hrwork	number of hours required to work at school per week
4	size	class size
5	preprmin	minutes spent per week in school on preparation
6	assist	minutes per week that a teacher's aide works with the teacher in the classroom

teachmi.csv Cluster sizes for data in **teachers.csv**.

Column	Name	Value
1	dist	School district size: large or med/small
2	school	school identifier
3	popteach	number of teachers in that school
4	ssteach	number of surveys returned from that school

teachnr.csv Data from a follow-up study of nonrespondents from Gnap (1995).

Column	Name	Value
1	hrwork	number of hours required to work at school per week
2	size	class size
3	preprmin	minutes spent per week in school on preparation
4	assist	minutes per week that a teacher's aide works with the teacher in the classroom

uneqvar.csv Artificial data used in exercises of Chapter 11.

Column	Name	Value
1	x	x
2	y	y

vietnam.csv Vietnam-service data from Stockford and Page (1984).

Column	Name	Value
1	apc	APC stratum. Character variable with options "Yes," "No," "NotAvail"
2	p2sample	Indicator variable for phase II sample, = 1 if in phase II sample, 0 otherwise
3	vietnam	= 1 if service in Vietnam, = 0 if service not in Vietnam, = -9 if not in phase II sample

vietnam.csv (continued)		
Column	Name	Value
4	phase1wt	weight for phase I sample
5	phase2wt	conditional weight for phase II sample, calculated as (phase I sample size in stratum) / (phase II sample size in stratum). phase2wt = -9 for observations not in phase 2 sample.
6	finalwt	final weight for phase II sample, calculated as phase1wt*phase2wt (= -9 for observations not in phase II sample)
7	p1apcsize	number of observations in the observation's APC stratum that are in the phase I sample (n_h)
8	p2apcsize	number of observations in the observation's APC stratum that are in the phase II sample (m_h)

vius.csv Selected variables from the 2002 U.S. Vehicle Inventory and Use Survey (VIUS). Source: U.S. Census Bureau (2006). The data were downloaded from www.census.gov/svsd/www/vius in May, 2006. The website from which the data were downloaded no longer exists, and online information about VIUS may now be found at <https://www.bts.gov/vius>, which provides a link to the archived 2002 data. The missing value of *state* for records with *adm_state* = 42 was recoded to “PA,” the state that has code 42. This data set has 98,682 records, which may be too large for some software packages to handle; the file **viusca.csv** is a smaller data set, with the same columns described below, containing only vehicles from California. The variable descriptions below are taken from the VIUS Data Dictionary.

Missing values are coded as -99 . For some variables, the value is missing because the question is not applicable or the vehicle is not in use; see the individual variable descriptions.

Note that a new VIUS is planned for 2022, with data to be released in 2023; see <https://www.bts.gov/vius>.

Column	Name	Value
1	stratum	stratum number (contains all 255 strata)
2	adm_state	state number
3	state	state name
4	trucktype	type of truck, used in stratification 1. pickups 2. minivans, other light vans, and sport utility vehicles 3. light single-unit trucks with gross vehicle weight less than 26,000 pounds 4. heavy single-unit trucks with gross vehicle weight greater than or equal to 26,000 pounds 5. truck-tractors
5	tabtrucks	column of sampling weights
6	bodytype	body type of vehicle 01. Pickup 02. Minivan 03. Light van other than minivan 04. Sport utility

vius.csv (continued)		
Column	Name	Value
7	adm_modelyear	05. Armored
		06. Beverage
		07. Concrete mixer
		08. Concrete pumper
		09. Crane
		10. Curtainside
		11. Dump
		12. Flatbed, stake, platform, etc.
		13. Low boy
		14. Pole, logging, pulpwood, or pipe
		15. Service, utility
		16. Service, other
		17. Street sweeper
		18. Tank, dry bulk
		19. Tank, liquids or gases
		20. Tow/Wrecker
		21. Trash, garbage, or recycling
		22. Vacuum
		23. Van, basic enclosed
		24. Van, insulated non-refrigerated
		25. Van, insulated refrigerated
		26. Van, open top
		27. Van, step, walk-in, or multistop
		28. Van, other
		99. Other not elsewhere classified
8	vius_gvw	model year
		01. 2003, 2002
		02. 2001
		03. 2000
		04. 1999
		05. 1998
		06. 1997
		07. 1996
		08. 1995
		09. 1994
		10. 1993
		11. 1992
		12. 1991
		13. 1990
		14. 1989
		15. 1988
		16. 1987
		17. Pre-1987
		Gross vehicle weight based on average reported weight
		01. Less than 6,001 lbs.
		02. 6,001 to 8,500 lbs.
		03. 8,501 to 10,000 lbs.
		04. 10,001 to 14,000 lbs.

vius.csv (continued)		
Column	Name	Value
		05. 14,001 to 16,000 lbs.
		06. 16,001 to 19,500 lbs.
		07. 19,501 to 26,000 lbs.
		08. 26,001 to 33,000 lbs.
		09. 33,001 to 40,000 lbs.
		10. 40,001 to 50,000 lbs.
		11. 50,001 to 60,000 lbs.
		12. 60,001 to 80,000 lbs.
		13. 80,001 to 100,000 lbs.
		14. 100,001 to 130,000 lbs.
		15. 130,001 lbs. or more
9	miles_annl	Number of Miles Driven During 2002
10	miles_life	Number of Miles Driven Since Manufactured
11	mpg	Miles Per Gallon averaged during 2002. Range from 0.3 to 35. -99 denotes not reported or not applicable.
12	opclass	Operator Classification With Highest Percent
		1. Private
		2. Motor carrier
		3. Owner operator
		4. Rental
		5. Personal transportation
		6. Not applicable (Vehicle not in use)
13	opclass_mtr	Percent of Miles Driven as a Motor Carrier. -99 denotes vehicle not in use
14	opclass_own	Percent of Miles Driven as an Owner Operator. -99 denotes vehicle not in use
15	opclass_psl	Percent of Miles Driven for Personal Transportation. -99 denotes vehicle not in use
16	opclass_pvt	Percent of Miles Driven as Private (Carry Own Goods or Internal Company Business Only). -99 denotes vehicle not in use
17	opclass_rnt	Percent of Miles Driven as Rental. -99 denotes vehicle not in use
18	transmssn	Type of Transmission
		1. Automatic
		2. Manual
		3. Semi-Automated Manual
		4. Automated Manual
19	trip_primary	Primary Range of Operation
		1. Off-the-road
		2. Less than 50 miles
		3. 51 to 100 miles
		4. 101 to 200 miles
		5. 201 to 500 miles
		6. 501 miles or more
		7. Not reported
		8. Not applicable (Vehicle not in use)
20	trip0_50	Percent of Annual Miles Accounted for with Trips

vius.csv (continued)		
Column	Name	Value
21	trip051_100	50 Miles or Less from the Home Base Percent of Annual Miles Accounted for with Trips 51 to 100 Miles from the Home Base
22	trip101_200	Percent of Annual Miles Accounted for with Trips 101 to 200 Miles from the Home Base
23	trip201_500	Percent of Annual Miles Accounted for with Trips 201 to 500 Miles from the Home Base
24	trip500more	Percent of Annual Miles Accounted for with Trips 501 or More Miles from Home Base
25	adm_make	Make of vehicle 01. Chevrolet 02. Chrysler 03. Dodge 04. Ford 05. Freightliner 06. GMC 07. Honda 08. International 09. Isuzu 10. Jeep 11. Kenworth 12. Mack 13. Mazda 14. Mitsubishi 15. Nissan 16. Peterbilt 17. Plymouth 18. Toyota 19. Volvo 20. White 21. Western Star 22. White GMC 23. Other (domestic) 24. Other (foreign)
26	business	Business in which vehicle was most often used during 2002 01. For-hire transportation or warehousing 02. Vehicle leasing or rental 03. Agriculture, forestry, fishing, or hunting 04. Mining 05. Utilities 06. Construction 07. Manufacturing 08. Wholesale trade 09. Retail trade 10. Information services 11. Waste management, landscaping, or administra- tive/support services

vius.csv (continued)		
Column	Name	Value
		12. Arts, entertainment, or recreation services
		13. Accommodation or food services
		14. Other services
		–99. Not reported or not applicable

winter.csv Selected variables from the Arizona State University Winter Closure Survey, taken in January 1995 (provided courtesy of the ASU Office of University Evaluation). This survey was taken to investigate the attitudes and opinions of university employees towards the closing of the university (for budgetary reasons) between December 25 and January 1. For the yes/no questions, the responses are coded as 1 = No, 2 = Yes. The variables *treatsta* and *treatme* are coded as 1=strongly agree, 2=agree, 3=undecided, 4=disagree, 5=strongly disagree. The variables *process* and *satbreak* are coded as 1=very satisfied, 2=satisfied, 3=undecided, 4=dissatisfied, 5=very dissatisfied. Variables *ownsupp* through *offclose* are coded 1 if the person checked that the statement applied to him/her, and 2 if the statement was not checked.

Missing values are coded as 9.

Column	Name	Value
1	class	stratum number 1 = faculty ; 2 = classified staff; 3 = administrative staff; 4 = academic professional
2	yearasu	number of years worked at ASU 1 = 1–2 years; 2 = 3–4 years; 3 = 5–9 years; 4 = 10–14 years; 5 = 15 or more years
3	vacation	In the past, have you <i>usually</i> taken vacation days the entire period between December 25 and January 1?
4	work	Did you work on campus during Winter Break Closure?
5	havediff	Did the Winter Break Closure cause you any difficulty/concerns?
6	negaeffe	Did the Winter Break Closure <i>negatively</i> affect your work productivity?
7	ownsupp	I was unable to obtain staff support in my department/office
8	othersup	I was unable to obtain staff support in other departments/offices
9	utility	I was unable to access computers, copy machine, etc. in my department/office
10	environ	I was unable to endure environmental conditions, e.g., not properly climatized
11	uniserve	I was unable to access university services necessary to my work
12	workelse	I was unable to work on my assignments because I work in another department/office
13	offclose	I was unable to work on my assignments because my office was closed
14	treatsta	Compared to other departments/offices, I feel staff in my department/office were treated fairly

winter.csv (continued)		
Column	Name	Value
15	treatme	Compared to other people working in my department/office, I feel I was treated fairly
16	process	How satisfied are you with the process used to inform staff about Winter Break Closure?
17	satbreak	How satisfied are you with the fact that ASU had a Winter Break Closure this year?
18	breakaga	Would you want to have Winter Break Closure again?

wtshare.csv Hypothetical sample of size 100, with indirect sampling. The data set has multiple records for adults with more than one child; if adult 254 has 3 children, adult 254 is listed 3 times in the data set. Note that to obtain L_k , you need to take $numadult + 1$.

Column	Name	Value
1	id	Identification number of adult in sample
2	child	= 1 if record is for a child, 0 if adult has no children
3	preschool	= 1 if child is in preschool, 0 otherwise
4	numadult	number of other adults in population who link to that child

Bibliography

- Allison, P. D. (2012). *Logistic Regression using SAS[®]: Theory and Application*. Cary, NC: SAS Institute, Inc.
- Arnold, T. W. (1991). Intraclutch variation in egg size of American coots. *The Condor* 93, 19–27.
- Asian Development Bank (2020). *Introduction to Small Area Estimation Techniques: A Practical Guide for National Statistics Offices*. Manila: Asian Development Bank.
- Azur, M. J., E. A. Stuart, C. Frangakis, and P. J. Leaf (2011). Multiple imputation by chained equations: What is it and how does it work? *International Journal of Methods in Psychiatric Research* 20(1), 40–49.
- Baillargeon, S. and L.-P. Rivest (2007). Rcapture: Loglinear models for capture-recapture in R. *Journal of Statistical Software* 19(5), 1–31.
- Baillargeon, S. and L.-P. Rivest (2011). The construction of stratified designs in R with the package stratification. *Survey Methodology* 37(1), 53–65.
- Barcaroli, G. (2014). SamplingStrata: An R package for the optimization of stratified sampling. *Journal of Statistical Software* 61(4), 1–24.
- Barcaroli, G., M. Ballin, H. Odendaal, D. Pagliuca, E. Willighagen, and D. Zardetto (2020). *SamplingStrata: Optimal Stratification of Sampling Frames for Multipurpose Sampling Surveys*. R package version 1.5-1, <https://CRAN.R-project.org/package=SamplingStrata> (accessed March 11, 2021).
- Bart, J. and S. Earnst (2002). Double-sampling to estimate density and population trends in birds. *The Auk* 119, 36–45.
- Bates, D., M. Mächler, B. Bolker, and S. Walker (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software* 67(1), 1–48.
- Bates, D., M. Mächler, B. Bolker, S. Walker, R. H. B. Christensen, H. Singmann, B. Dai, F. Scheipl, G. Grothendieck, P. Green, J. Fox, A. Bauer, and P. N. Krivitsky (2020). *lme4: Linear Mixed-Effects Models using ‘Eigen’ and S4*. R package version 1.1-26, <https://CRAN.R-project.org/package=lme4> (accessed March 20, 2021).
- Beck, A. J., S. A. Kline, and L. A. Greenfeld (1988). Survey of Youth in Custody. Technical Report NCJ-113365, Bureau of Justice Statistics, Washington, DC.
- Bretz, F., T. Hothorn, and P. Westfall (2016). *Multiple Comparisons using R*. Boca Raton, FL: CRC Press.
- Brewer, K. R. W. (1963). Ratio estimation and finite populations: Some results deducible from the assumption of an underlying stochastic process. *The Australian Journal of Statistics* 5(3), 93–105.

- Brewer, K. R. W. (1975). A simple procedure for sampling π pswor. *The Australian Journal of Statistics* 17, 166–172.
- Brick, J. M., D. Morganstein, and R. Valliant (2000). *Analysis of Complex Sample Data Using Replication*. Rockville, MD: Westat.
- Bueno, E. (2020). *optimStrat: Choosing the Sample Strategy*. R package version 2.3. <https://CRAN.R-project.org/package=optimStrat> (accessed March 12, 2021).
- Canty, A. J. and A. C. Davison (1999). Resampling-based variance estimation for labour force surveys. *The Statistician* 48, 379–391.
- Centers for Disease Control and Prevention (2017). NHANES Questionnaires, Datasets, and Related Documentation. <https://wwwn.cdc.gov/nchs/nhanes/> (accessed August 15, 2020).
- Chao, A., K. H. Ma, T. C. Hsieh, and C.-H. Chiu (2016). *SpadeR: Species-Richness Prediction and Diversity Estimation with R*. R package version 0.1.1, <https://CRAN.R-project.org/package=SpadeR> (accessed March 15, 2021).
- Chapman, D. G. (1951). Some properties of the hypergeometric distribution with applications to zoological sample censuses. *University of California Publications in Statistics* 1, 131–160.
- Chauvet, G. and Y. Tillé (2006). A fast algorithm for balanced sampling. *Computational Statistics* 21(1), 53–62.
- Cho, I., J.-K. Kim, J. Im, and Y. Yang (2020). *FHDI: Fractional Hot Deck and Fully Efficient Fractional Imputation*. R package version 1.4.1, <https://CRAN.R-project.org/package=FHDI> (accessed April 15, 2021).
- Cormack, R. M. (1992). Interval estimation for mark-recapture studies of closed populations. *Biometrics* 48, 567–576.
- Dippo, C. S., R. E. Fay, and D. H. Morganstein (1984). Computing variances from complex samples with replicate weights. In *Proceedings of the Survey Research Methods Section*, pp. 489–494. Alexandria, VA: American Statistical Association.
- Domingo-Salvany, A., R. L. Hartnoll, A. Maquire, J. M. Suelves, and J. M. Anto (1995). Use of capture-recapture to estimate the prevalence of opiate addiction in Barcelona, Spain, 1989. *American Journal of Epidemiology* 141, 567–574.
- Fienberg, S. E. and A. Rinaldo (2007). Three centuries of categorical data analysis: Log-linear models and maximum likelihood estimation. *Journal of Statistical Planning and Inference* 137(11), 3430–3445.
- Forman, S. L. (2004). Baseball-reference.com—Major league statistics and information. www.baseball-reference.com (accessed November 2004).
- Gambino, J. G. (2021). *pps: PPS Sampling*. R package version 1.0. <https://CRAN.R-project.org/package=pps> (accessed March 12, 2021).
- Gini, C. and L. Galvani (1929). Di una applicazione del metodo rappresentativo all’ultimo censimento italiano della popolazione. *Annali di Statistica* 6(4), 1–105.
- Gnap, R. (1995). *Teacher Load in Arizona Elementary School Districts in Maricopa County*. Ph.D. dissertation. Tempe, AZ: Arizona State University.

- Goga, C. (2018). Brief overview of survey sampling techniques with R. *Romanian Statistical Review* 2018(1), 83–94.
- Grafström, A. and J. Lisic (2019). *BalancedSampling: Balanced and Spatially Balanced Sampling*. R package version 1.5.5, <https://CRAN.R-project.org/package=BalancedSampling> (accessed March 12, 2021).
- Hand, D. J., F. Daly, A. D. Lunn, K. J. McConway, and E. Ostrowski (1994). *A Handbook of Small Data Sets*. London: Chapman and Hall.
- Hanurav, T. V. (1967). Optimum utilization of auxiliary information: π ps sampling of two units from a stratum. *Journal of the Royal Statistical Society, Series B* 29, 374–391.
- Harmening, S., A.-K. Kreutzmann, S. Pannier, N. Rojas-Perilla, N. Salvati, T. Schmid, M. Templ, N. Tzavidis, and N. Würz (2021). *emdi: Estimating and Mapping Disaggregated Indicators*. R package version 2.0.2, <https://CRAN.R-project.org/package=emdi> (accessed April 30, 2021).
- Harrell, F. E. (2021). *Hmisc: Harrell miscellaneous*. R package version 4.5-0, <https://CRAN.R-project.org/package=Hmisc> (accessed April 7, 2021).
- Hartley, H. O. and J. N. K. Rao (1962). Sampling with unequal probabilities and without replacement. *The Annals of Mathematical Statistics* 33, 350–374.
- Hayat, M. and T. Knapp (2017). Randomness and inference in medical and public health research. *Journal of the Georgia Public Health Association* 7(1), 7–11.
- Haziza, D. (2009). Imputation and inference in the presence of missing data. In D. Pfeffermann and C. R. Rao (Eds.), *Sample Surveys: Design, Methods, and Applications. Handbook of Statistics, Volume 29A*, pp. 215–246. Amsterdam: North-Holland.
- Heck, P. R., D. J. Simons, and C. F. Chabris (2018). 65% of Americans believe they are above average in intelligence: Results of two nationally representative surveys. *PloS One* 13(7), 1–11.
- Hidiroglou, M. A., J.-F. Beaumont, and W. Yung (2019). Development of a small area estimation system at Statistics Canada. *Survey Methodology* 45(1), 101–126.
- Horton, N. J. and K. Kleinman (2015). *Using R and RStudio for Data Management, Statistical Analysis, and Graphics, 2nd ed.* Boca Raton, FL: CRC Press.
- Horton, N. J., R. Pruim, and D. T. Kaplan (2018). *A Student's Guide to R*. Amherst, MA: Project MOSAIC.
- Hyndman, R. J. and Y. Fan (1996). Sample quantiles in statistical packages. *The American Statistician* 50(4), 361–365.
- Im, J., I. H. Cho, and J.-K. Kim (2018). FHDI: An R package for fractional hot deck imputation. *R Journal* 10(1), 140–154.
- Ismay, C. and P. C. Kennedy (2019). *Getting Used to R, RStudio, and R Markdown*. <https://rbasics.netlify.app/> (accessed March 2, 2021).
- Jacoby, J. and A. H. Handlin (1991). Non-probability sampling designs for litigation surveys. *Trademark Reporter* 81, 169–179.

- Judkins, D. (1990). Fay's method for variance estimation. *Journal of Official Statistics* 6, 223–240.
- Kabacoff, R. I. (2021). *R in Action, 3rd ed.* Shelter Island, NY: Manning Publications.
- Kim, J. K. (2011). Parametric fractional imputation for missing data analysis. *Biometrika* 98(1), 119–132.
- Kim, J. K. and W. Fuller (2004). Fractional hot deck imputation. *Biometrika* 91(3), 559–578.
- Koch, G. G., D. H. Freeman, and J. L. Freeman (1975). Strategies in the multivariate analysis of data from complex surveys. *International Statistical Review* 43, 59–78.
- Koenker, R. (2005). *Quantile Regression*. Cambridge: Cambridge University Press.
- Koenker, R., S. Portnoy, P. T. Ng, B. Melly, A. Zeileis, P. Grosjean, C. Moler, Y. Saad, V. Chernozhukov, I. Fernandez-Val, and B. D. Ripley (2021). *quantreg: Quantile Regression*. R package version 2.23-18, <https://CRAN.R-project.org/package=quantreg> (accessed April 7, 2021).
- Korn, E. L. and B. I. Graubard (1998). Confidence intervals for proportions with small expected number of positive counts estimated from survey data. *Survey Methodology* 24, 193–201.
- Kott, P. S. (2012). Why one should incorporate the design weights when adjusting for unit nonresponse using response homogeneity groups. *Survey Methodology* 38(1), 95–99.
- Kowarik, A. and M. Templ (2016). Imputation with the R package VIM. *Journal of Statistical Software* 74(7), 1–16.
- Kreutzmann, A.-K., S. Pannier, N. Rojas-Perilla, T. Schmid, M. Templ, and N. Tzavidis (2019). The R package emdi for estimating and mapping regionally disaggregated indicators. *Journal of Statistical Software* 91, 1–33.
- Kruuk, H., A. Moorhouse, J. W. H. Conroy, L. Durbin, and S. Frears (1989). An estimate of numbers and habitat preferences of otters *Lutra lutra* in Shetland, UK. *Biological Conservation* 49, 241–254.
- Little, R. J. and S. Vartivarian (2003). On weighting the rates in non-response weights. *Statistics in Medicine* 22(9), 1589–1599.
- Lohr, S. L. (2022). *SAS® Software Companion for Sampling: Design and Analysis*, 3rd ed. Boca Raton, FL: CRC Press.
- Lopez-Vizcaino, E., M. Lombardia, and D. Morales (2019). *mme: Multinomial Mixed Effects Models*. R package version 0.1-6. <https://CRAN.R-project.org/package=mme> (accessed January 16, 2021).
- Lu, Y. and S. L. Lohr (2021). *SDAResources: Datasets and Functions for “Sampling: Design and Analysis.”* R package version 0.1.0. <https://CRAN.R-project.org/package=SDAResources> (accessed May 17, 2021).
- Lumley, T. (2004). Analysis of complex survey samples. *Journal of Statistical Software* 9(1), 1–19.
- Lumley, T. (2010). *Complex Surveys: A Guide to Analysis using R*. Hoboken, NJ: Wiley.

- Lumley, T. (2020). *survey: Analysis of Complex Survey Samples*. R package version 4.0. <https://CRAN.R-project.org/package=survey> (accessed September 20, 2020).
- Lundquist, P. and C.-E. Särndal (2013). Aspects of responsive design with applications to the Swedish Living Conditions Survey. *Journal of Official Statistics* 29(4), 557–582.
- Luraschi, J. (2021). Importing data with RStudio. <https://support.rstudio.com/hc/en-us/articles/218611977-Importing-Data-with-RStudio> (accessed March 12, 2021).
- Lydersen, C. and M. Ryg (1991). Evaluating breeding habitat and populations of ringed seals *Phoca hispida* in Svalbard fjords. *Polar Record* 27, 223–228.
- Macdonell, W. R. (1901). On criminal anthropometry and the identification of criminals. *Biometrika* 1, 177–227.
- Molina, I. and Y. Marhuenda (2015). sae: An R package for small area estimation. *The R Journal* 7(1), 1–98.
- Molina, I. and Y. Marhuenda (2020). *sae: Small Area Estimation*. R package version 1.3. <https://CRAN.R-project.org/package=sae> (accessed December 20, 2020).
- National Center for Health Statistics (1987). *Vital Statistics of the United States, Volume 3: Marriage and Divorce*. Washington, DC: U.S. Government Printing Office.
- Nolan, D. and T. Speed (2000). *Stat Labs: Mathematical Statistics Through Applications*. New York: Springer.
- Oetiker, T., H. Partl, I. Hyna, and E. Schlegl (2021). *The Not So Short Introduction to L^AT_EX 2_ε, Version 6.4*. Olten, Switzerland: Tobias Oetiker. <https://tobi.oetiker.ch/lshort/lshort.pdf> (accessed March 18, 2021).
- Peart, D. (1994). *Impacts of Feral Pig Activity on Vegetation Patterns Associated with Quercus agrifolia on Santa Cruz Island, California*. Ph.D. dissertation. Tempe, AZ: Arizona State University.
- Pinheiro, J., D. Bates, S. DebRoy, D. Sarkar, and R Core Team (2021). *nlme: Linear and Nonlinear Mixed Effects Models*. R package version 3.1-152, <https://CRAN.R-project.org/package=nlme> (accessed March 20, 2021).
- Pratesi, M. (Ed.) (2016). *Analysis of Poverty Data by Small Area Estimation*. Hoboken, NJ: Wiley.
- Preston, J. (2009). Rescaled bootstrap for stratified multistage sampling. *Survey Methodology* 35(2), 227–234.
- R Core Team (2021). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. Version 4.0.4, <https://www.R-project.org>.
- Rao, J. N. K. and I. Molina (2015). *Small Area Estimation, 2nd ed.* Hoboken, NJ: Wiley.
- Rao, J. N. K. and A. J. Scott (1981). The analysis of categorical data from complex sample surveys: Chi-squared tests for goodness of fit and independence in two-way tables. *Journal of the American Statistical Association* 76, 221–230.
- Rao, J. N. K. and A. J. Scott (1984). On chi-squared tests for multiway contingency tables with cell proportions estimated from survey data. *The Annals of Statistics* 12, 46–60.

- Rao, J. N. K., C. F. J. Wu, and K. Yue (1992). Some recent work on resampling methods for complex surveys. *Survey Methodology* 18, 209–217.
- Reiter, J. P., T. E. Raghunathan, and S. K. Kinney (2006). The importance of modeling the sampling design in multiple imputation for missing data. *Survey Methodology* 32(2), 143–149.
- Rivest, L.-P. and S. Baillargeon (2017). *stratification: Univariate Stratification of Survey Populations*. R package version 2.2-6, <https://CRAN.R-project.org/package=stratification> (accessed March 10, 2021).
- Rivest, L.-P. and S. Baillargeon (2019). *Rcapture: Loglinear Models for Capture-Recapture Experiments*. R package version 1.4-3, <https://CRAN.R-project.org/package=Rcapture> (accessed March 15, 2021).
- Roberts, R. J., Q. D. Sandifer, M. R. Evans, M. Z. Nolan-Ferrell, and P. M. Davis (1995). Reasons for non-uptake of measles, mumps, and rubella catch up immunisation in a measles epidemic and side effects of the vaccine. *British Medical Journal* 310, 1629–1632.
- Ruggles, S., M. Sobek, T. Alexander, C. A. Fitch, R. Goeken, P. K. Hall, M. King, and C. Ronnander (2004). Integrated Public Use Microdata Series: Version 3.0 [machine-readable database]. www.ipums.org/usa (accessed September 17, 2008).
- Sampford, M. R. (1967). On sampling without replacement with unequal probabilities of selection. *Biometrika* 54, 499–513.
- SAS Institute Inc. (2021). *SAS/STAT[®] User's Guide*. Cary, NC: SAS Institute Inc. <https://documentation.sas.com/> (accessed April 27, 2021).
- Stockford, D. D. and W. F. Page (1984). Double sampling and the misclassification of Vietnam service. In *Proceedings of the Social Statistics Section*, pp. 261–264. Alexandria, VA: American Statistical Association.
- Talbot, N. L. C. (2012). *LaTeX for Complete Novices*. Saxlingham Nethergate, UK: Dickimaw Books. <https://www.dickimaw-books.com/latex/novices/> (accessed March 10, 2021).
- Templ, M., A. Kowarik, A. Alfons, G. de Cillia, B. Prantner, and W. Rannetbauer (2021). *VIM: Visualization and Imputation of Missing Values*. R package version 6.1.0, <https://CRAN.R-project.org/package=VIM> (accessed April 21, 2021).
- Thomas, D. R. and J. N. K. Rao (1987). Small-sample comparisons of level and power for simple goodness-of-fit statistics under cluster sampling. *Journal of the American Statistical Association* 82, 630–636.
- Tillé, Y. (2006). *Sampling Algorithms*. New York: Springer.
- Tillé, Y. and A. Matei (2010). Teaching survey sampling with the ‘sampling’ R package. In *Proceedings of the 8th International Conference on Teaching Statistics (ICOTS)*, pp. 1–6. Auckland, NZ: International Association of Statistical Education.
- Tillé, Y. and A. Matei (2021). *sampling: Survey Sampling*. R package version 2.9, <https://CRAN.R-project.org/package=sampling> (accessed March 12, 2021).
- Tillé, Y. and M. Wilhelm (2017). Probability sampling designs: Principles for choice of design and balancing. *Statistical Science* 32(2), 176–189.

- Tzavidis, N., L.-C. Zhang, A. Luna, T. Schmid, and N. Rojas-Perilla (2018). From start to finish: A framework for the production of small area official statistics. *Journal of the Royal Statistical Society: Series A* 181(4), 927–979.
- U.S. Bureau of the Census (1921). *Fourteenth Census of the United States Taken in the Year 1920*. Washington, DC: U.S. Government Printing Office. <https://www.census.gov/library/publications/1921/dec/vol-01-population.html> (accessed August 4, 2020).
- U.S. Bureau of the Census (1995). *1992 Census of Agriculture, Volume 1: Geographic Area Series*. Washington, DC: U.S. Bureau of the Census.
- U.S. Census Bureau (1994). *County and City Data Book: 1994*. Washington, DC: U.S. Census Bureau.
- U.S. Census Bureau (2006). *Vehicle Inventory and Use Survey—Methods*. Washington, DC: U.S. Census Bureau.
- U.S. Census Bureau (2012). *United States Summary, 2010*. Washington, DC: U.S. Census Bureau. <https://www.census.gov/prod/cen2010/cph-2-1.pdf> (accessed October 3, 2020).
- U.S. Census Bureau (2019). State population totals: 2010-2019. Table 1. Annual estimates of the resident population for the United States, regions, states, and Puerto Rico: April 1, 2010 to July 1, 2019 (NST-EST2019-01). <https://www.census.gov/data/datasets/time-series/demo/popest/2010s-state-total.html> (accessed August 3, 2020).
- U.S. Department of Education (2020). College scorecard data. <https://collegescorecard.ed.gov/data/> (accessed August 25, 2020).
- U.S. Department of Justice (1989). *Survey of Youth in Custody, 1987, United States computer file, Conducted by Department of Commerce, Bureau of the Census, 2nd ICPSR ed.* Ann Arbor, MI: Inter-University Consortium for Political and Social Research.
- Valliant, R., J. A. Dever, and F. Kreuter (2018). *Practical Tools for Designing and Weighting Survey Samples*. New York: Springer.
- Valliant, R., J. A. Dever, and F. Kreuter (2020). *PracTools: Tools for Designing and Weighting Survey Samples*. R package version 1.2.2. <https://CRAN.R-project.org/package=PracTools> (accessed March 12, 2021).
- Valliant, R. and K. F. Rust (2010). Degrees of freedom approximations and rules-of-thumb. *Journal of Official Statistics* 26(4), 585–602.
- van Buuren, S. (2018). *Flexible Imputation of Missing Data*. Boca Raton, FL: CRC Press.
- van Buuren, S., K. Groothuis-Oudshoorn, G. Vink, R. Schouten, A. Robitzsch, P. Rocken-schaub, L. Doove, S. Jolani, M. Moreno-Betancur, I. White, P. Gaffert, F. Meinfelder, B. Gray, and V. Arel-Bundock (2021). *mice: Multivariate Imputation by Chained Equations*. R package version 3.13.0, <https://CRAN.R-project.org/package=mice> (accessed April 15, 2021).
- Vijayan, K. (1968). An exact π ps sampling scheme: Generalization of a method of Hanurav. *Journal of the Royal Statistical Society, Series B* 30, 556–566.

- Virginia Polytechnic and State University/Responsive Management (2006). *An Assessment of Public and Hunter Opinions and the Costs and Benefits to North Carolina of Sunday Hunting*. Blacksburg, VA: Virginia Polytechnic and State University.
- Wand, M., C. Moler, and B. Ripley (2020). *KernSmooth: Functions for Kernel Smoothing Supporting Wand & Jones (1995)*. R package version 2.23-18, <https://CRAN.R-project.org/package=KernSmooth> (accessed April 7, 2021).
- Wand, M. P. and M. C. Jones (1995). *Kernel Smoothing*. London: Chapman & Hall.
- Wang, J. (2021). The pseudo maximum likelihood estimator for quantiles of survey variables. *Journal of Survey Statistics and Methodology* 9(1), 185–201.
- Wickham, H. (2015). *R Packages: Organize, Test, Document, and Share Your Code*. Sebastopol, CA: O'Reilly Media.
- Wickham, H. (2019). *Advanced R, 2nd ed.* Boca Raton, FL: CRC Press.
- Wickham, H., W. Chang, L. Henry, T. L. Pedersen, K. Takahashi, C. Wilke, K. Woo, H. Yutani, and D. Dunnington (2020). *ggplot2: Create Elegant Data Visualisations Using the Grammar of Graphics*. R package version 3.3.3, <https://CRAN.R-project.org/package=ggplot2> (accessed March 30, 2021).
- Wikibooks contributors (2021). *LaTeX*. Wikibooks, The Free Textbook Project. <https://en.wikibooks.org/wiki/LaTeX> (accessed March 18, 2021).
- Wilk, S. J., W. W. Morse, D. E. Ralph, and T. R. Azarovitz (1977). *Fishes and Associated Environmental Data Collected in New York Bight, June 1974–June 1975*. NOAA Tech. Rep. No. NMFS SSRF-716. Washington, DC: U.S. Government Printing Office.
- Woodruff, R. S. (1952). Confidence intervals for medians and other position measures. *Journal of the American Statistical Association* 47, 636–646.
- Xie, Y. (2015). *Dynamic Documents with R and knitr, 2nd ed.* Boca Raton, FL: CRC Press.
- Yadav, M. L. and B. Roychoudhury (2018). Handling missing values: A study of popular imputation packages in R. *Knowledge-Based Systems* 160, 104–118.
- Zhang, C., C. Antoun, H. Y. Yan, and F. G. Conrad (2020). Professional respondents in opt-in online panels: What do we really know? *Social Science Computer Review* 38(6), 703–719.
- Zhang, G., F. Christensen, and W. Zheng (2015). Nonparametric regression estimators in complex surveys. *Journal of Statistical Computation and Simulation* 85(5), 1026–1034.

Index

- Allocation, stratified sampling, 27–29
- Analysis of variance (ANOVA), 37, 52
- Balanced repeated replication (BRR), 113, 116–124, 144
- Balanced sampling, 151
- Bootstrap, 113, 116–124
 - capture-recapture estimation, 158–159
- Boxplot
 - complex survey, 92–94, 100
 - domains, 148
 - stratified random sample, 32
- Capture-recapture estimation, 157–159
- Categorical variable, 23
- Chi-square test, 129–132
- Clopper-Pearson confidence interval, 24, 37
- Cluster sample
 - equal probabilities, 57–68
 - estimating means, 57–63
 - estimating proportions, 61
 - finite population correction, 62–63
 - selection, 69–77
 - unequal probabilities, 69–83
- Combined ratio estimator, 49
- Contingency table, 129–131
- Conventions used in book, 1
- Cumulative distribution function, empirical, 103–105
- Degrees of freedom (df)
 - bootstrap SRS, 123
 - bootstrap stratified random sample, 124
 - chi-square test, 129
 - domain estimation, 47, 119
 - jackknife cluster sampling, 122
 - jackknife SRS, 121
 - one-stage cluster sample, 57–59, 78
 - random group, 116
 - regression, 45, 141, 143
 - replicated samples, 115
 - simple random sample, 21
 - stratified random sample, 35
 - two-stage cluster sample, 60, 79
- Density estimate, 93–94
- Design effect (deff), 92, 132–133
- Domain, 46
 - comparison, 145–148
 - estimation, 46–48, 90, 148
- Empirical cumulative distribution function, 88, 103–105
- Empirical probability mass function, 103–105
- Examples in SDA
 - Example 02.05, 15–17
 - Example 02.06, 17, 19–23
 - Example 02.07, 19–23
 - Example 02.11, 19–23
 - Example 03.02, 30–37
 - Example 03.06, 32–37
 - Example 04.02, 41–42
 - Example 04.03, 41–42
 - Example 04.05, 42–44
 - Example 04.07, 44–45
 - Example 04.08, 46–47
 - Example 04.09, 48, 124–125
 - Example 04.11, 50–52
 - Example 04.12, 52–54
 - Example 05.02, 57–58
 - Example 05.06, 58–59
 - Example 05.07, 61–63
 - Example 05.08, 59–60
 - Example 05.12, 63
 - Example 05.14, 64
 - Example 06.02, 69–70
 - Example 06.04, 77–78
 - Example 06.06, 78–79
 - Example 06.08, 81
 - Example 06.09, 81–82
 - Example 06.10, 82–83
 - Example 06.11, 72–77, 79–80
 - Example 07.05, 88–89, 103
 - Example 07.06, 89
 - Example 07.09, 90–92

- Example 07.10, 93–95
- Example 07.11, 93–95
- Example 07.12, 93–95
- Example 08.09, 111
- Example 09.03, 113–115
- Example 09.04, 115–116
- Example 09.05, 117–119, 126
- Example 09.06, 119–120
- Example 09.07, 120–122
- Example 09.08, 122
- Example 09.09, 122–123
- Example 09.10, 123–124
- Example 09.12, 89
- Example 10.01, 129–130
- Example 10.05, 130–131
- Example 10.06, 131–133
- Example 10.08, 133–135
- Example 10.09, 135
- Example 11.02, 139–140
- Example 11.04, 140
- Example 11.06, 141–142
- Example 11.07, 142–145
- Example 11.08, 146–147
- Example 11.09, 147–148
- Example 11.12, 149–150
- Example 12.01, 155–157
- Example 12.04, 155–157
- Example 13.01, 157–159
- Example 13.03, 159–161
- Exercises in SDA
 - Exercise 02.27, 24
 - Exercise 02.32, 24
 - Exercise 02.34, 24
 - Exercise 05.40, 65–67
 - Exercise 06.36, 82
 - Exercise 06.45, 70
 - Exercise 07.19, 88
 - Exercise 11.37, 151
- Finite population correction (fpc)
 - chi-square test, 131
 - one-stage cluster sampling, 58
 - regression coefficients, 141
 - simple random sampling, 19, 21
 - stratified random sampling, 34–35
 - two-stage cluster sampling, 60–63
 - unequal-probability sampling, 77
- Functions in R
 - abline, 42
 - anova, 135–137
 - apply, 81, 158–159
 - as.data.frame, 4
 - as.svrepdesign, 116–125, 141
 - axis, 150
 - binom.test, 24
 - box, 150
 - boxplot, 9, 33, 94
 - calibrate, 110
 - captureci, 158, 160–161
 - chisq.test, 129, 133–134
 - cluster, 13, 70–71
 - coef, 22, 42, 114–115
 - complete.cases, 107–108
 - confint, 21–23, 44–45, 57–59, 117–125
 - cor, 41–44
 - cut, 100
 - data, 7
 - data.frame, 78, 79, 117–119, 124–125, 133–134
 - degf, 33–35, 120–122, 140–141
 - density, 92–94
 - ecdf, 104–105
 - emppmf, 103–104
 - factor, 23, 37, 47, 62–64, 91, 110, 147
 - getdata, 17–18, 32, 75–76
 - glm, 44, 50–159
 - help, 4
 - help.search, 4
 - hist, 19, 158–159
 - inclusionprobabilities, 72, 75–76
 - install.packages, 3
 - intervals_ex40, 65–67
 - is.na, 12, 107–108
 - jpeg, 9
 - length, 144–145
 - library, 3
 - lines, 92–95, 101–102, 144–145, 150
 - lm, 52, 54, 63, 111, 139–140
 - lme, 64
 - lmer, 64
 - locpoly, 101
 - matrix, 4, 129
 - mean, 107–108
 - mstage, 71–75, 85–88
 - nobs, 143
 - order, 85
 - par, 92–94
 - pdf, 9
 - plot, 42, 50–54, 96–97, 150
 - postStratify, 48, 109, 124–125
 - predict, 41–42, 44–45, 50–54, 111, 144–145

- qt, 20
 - rake, 109–110
 - read.csv, 6–7
 - read.table, 6–7
 - regTermTest, 149–150
 - round, 100
 - rq, 102
 - sample, 15, 97
 - bootstrap for capture-recapture, 158–159
 - SRS, 15–16
 - stratified random sample, 30
 - unequal-probability sample, 69–70
 - seq, 144–145, 150
 - set.seed, 16, 25, 30, 65–66, 69–70, 114–115, 123
 - sink (saving output), 7–8
 - source, 8
 - srswor, 15–18, 72, 114–115
 - srswor1, 24
 - srswr, 15–18, 72
 - stargazer, 8–9
 - strata, 13, 31–32, 76
 - subset, 46–48, 91–92, 94, 99, 108, 150
 - summary, 130–131, 140–143, 146–150
 - svrepdesign, 126–127
 - svyboxplot, 92, 94–95, 100
 - svyby, 35–36, 47, 115–116, 145, 146
 - svycdf, 95, 104–105
 - svychisq, 130–134
 - svyciprop, 24, 37
 - svydesign, 15, 20, 100, 124–125
 - complex sample, 89–132, 142–143
 - one-stage cluster sample, 57–59
 - one-stage unequal probability sample, 78
 - pps sampling without replacement, 82
 - random group method, 115–116
 - replicated sample design, 114–115
 - SRS, 21–22, 140–141, 157–158
 - two-stage cluster sample, 59
 - use in calculating replicate weight variances, 141
 - use in creating replicate weights, 117–120, 122–124
 - svyglm, 44–45, 139–143, 146–150
 - svyhist, 92–95
 - svyloglin, 133–137
 - svymean, 15, 20, 92–95, 107–108, 114–125, 131–133, 139, 155–157
 - complex sample, 89
 - one-stage cluster sample, 57–59
 - proportion, 22–23
 - proportions with stratified random sample, 36–37
 - SRS, 21–22
 - stratified random sample, 33–35
 - two-stage cluster sample, 59
 - with replicate weights, 126
 - svyplot, 95, 98–99, 144–145
 - svyquantile, 88–92, 104, 119–120, 123
 - svyratio, 120–122, 157–158
 - one-stage unequal probability sample, 78
 - SRS, 41–44
 - stratified sampling, 49–50
 - svysmooth, 92–95, 101
 - svytable, 107–108, 130–132, 135–137
 - svytotal, 15, 20, 124–125
 - one-stage cluster sample, 57–59
 - SRS, 21–22
 - stratified random sample, 33–37
 - t.test, 114–115
 - table, 27–28, 32, 107–108, 145–146
 - tapply, 30, 35–36, 57
 - trimWeights, 110
 - twophase, 155–157
 - unique, 5, 27, 32
 - UPbrewer, 72
 - UPpoisson, 72
 - UPsampford, 72, 75–76
 - UPsampfordpi2, 80
 - UPsystematic, 72
 - write.csv, 9
 - write.table, 9
- Histogram
- complex survey, 92–94, 150
 - SRS, 19
- Imputation, 110–112
- Intraclass correlation coefficient, 65
- Jackknife, 113, 116–124
- Lists, 73
- Logistic regression, 149–150
- Loglinear model, 133–137
- Masking of R functions, 13
- Missing data, 12–13, 90
 - regression analyses, 143–144

- Mixed model, 64–65, 151, 162
- Multivariate imputation by chained equations, 111
- National Health and Nutrition Examination Survey (NHANES), 85, 89–92
 - graphs, 95–102
 - regression models, 142–150
- Nonresponse, 107–112
- Numeric variable, 23
- Odds ratio, 129–130, 150
- Packages in R, 2
 - BalancedSampling, 151
 - emdi, 162
 - FHDI, 111
 - ggplot2, 102
 - Hmisc, 111
 - installr, 2
 - kernSmooth, 101
 - knitr, 10
 - lme4, 64, 151
 - mice, 111
 - mme, 162
 - nlme, 64, 151
 - pps, 71
 - PracTools, 29
 - quantreg, 102
 - Rcapture, 161
 - sae, 161
 - sampling, xi, 3, 15, 27, 69, 85, 113, 151
 - SamplingStrata, 29
 - SDAResources, 3, 7, 15, 27, 65, 103, 158
 - SpadeR, 161
 - stargazer, 8
 - stratification, 29
 - survey, xi, 3, 15, 27, 41, 57, 77, 80, 88, 109, 113, 129, 139, 155, 162
 - VIM, 111
- Poststratification, 48, 108–110
- Probability proportional to size (pps) sampling, 69
- Quantiles, 88–89, 119, 123
- R chunk, 10–11
- R software, obtaining, 2
- R-squared, 143
- Raking, 108–110
- Random effects model, 64–65
- Random group method for estimating variances, 113–116
- Rao–Scott test, 131–132, 143
- Ratio estimation, 41–44
 - combined, 49
 - separate, 49–50
 - stratified sampling, 49–50
- Regression
 - design-based, 140–145
 - model-based, 139–140, 151
- Regression estimation, 44–45
- Replicate weights
 - constructing, 116–124
 - from survey data file, 126–127
 - nonresponse adjustments, 124–125
- Replication variance estimation, 80, 113–127
- Reproducible research, 10
- RStudio, 2, 10
- Scatterplot
 - complex survey, 95–102, 144
 - SRS, 42–43
- SDA, acronym for Sampling Design and Analysis, xi, 1
- SDAResources
 - agpop, 7, 17–18, 27–32
 - agpps, 82
 - agsrs, 19–23, 41–48, 50–52, 124–125
 - agstrat, 32–37, 49–50
 - algebra, 58–59
 - anthrs, 139–141
 - classes, 9, 69–75, 85–88
 - classeslong, 9, 85–88
 - classpps, 80
 - college, 114–115
 - collegerg, 120–122
 - coots, 59–60, 122
 - deadtrees, 44–45, 52–54
 - gpa, 57–58
 - htpop, 104
 - htsrs, 88–89, 123
 - htstrat, 89, 92–94, 103–105, 123–124
 - impute, 107–108, 110–111
 - intervals_ex40, 65–67
 - nhanes, 89–92, 95, 119–120, 142–150
 - santacruz, 42–44
 - schools, 61–65
 - syc, 115–116, 130–133, 135–137
 - vietnam, 155–157
- Separate ratio estimator, 49–50

Simple random sample
 computing estimates, 18–23
 selection, 15–18
 weights, 18

Small area estimation, 161–162

Stratified sample
 allocation, 27–29
 computing estimates, 32–37
 selection, 30–32

Sweave, 10–11

Trend line
 quantiles, 102
 smoothed for mean, 101

Wald test, 131, 136, 143, 150

Website for book, xi

Weight
 cluster sample, 71, 75
 SRS, 18
 stratified random sample, 32

Weighted least squares, 50

Weighting class adjustments for
 nonresponse, 108

